

CSCI-564 Advanced Computer Architecture

Pipelining Hazards

Credit to Dr.Bo Wu
Colorado School of Mines

Optional Readings (After class)

Virtual memory III:

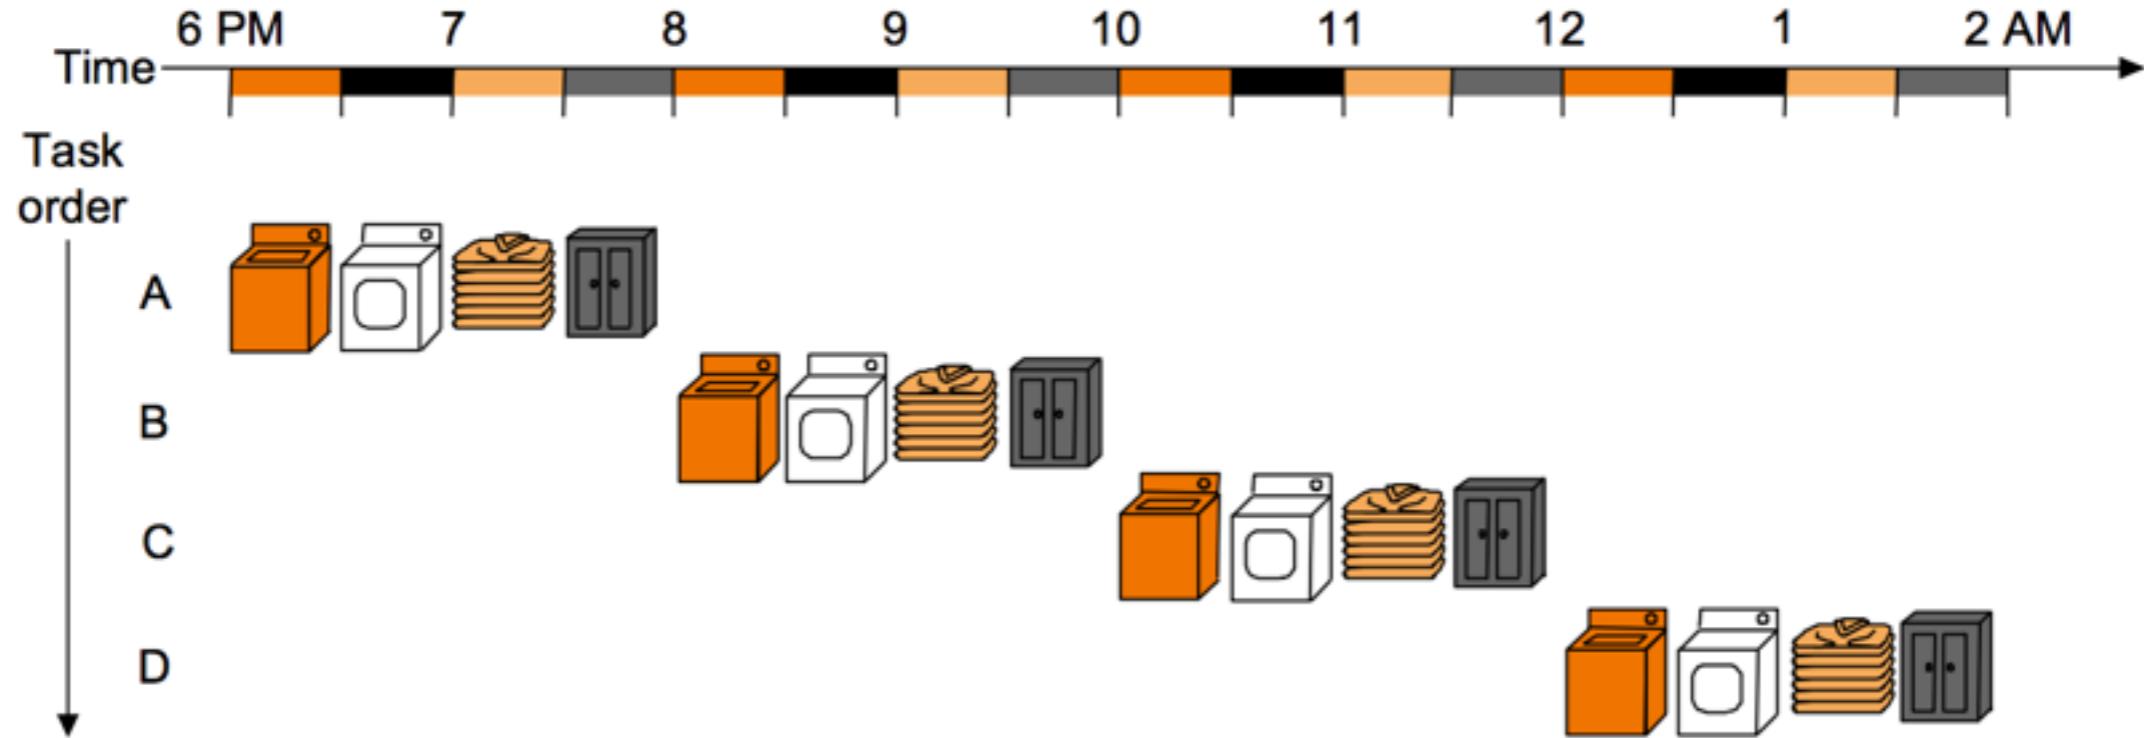
<https://blog.httrack.com/blog/2014/04/05/a-story-of-realloc-and-laziness/>

Virtual memory IV:

<http://thebeardsage.com/virtual-memory-translation-lookaside-buffer-tlb/>

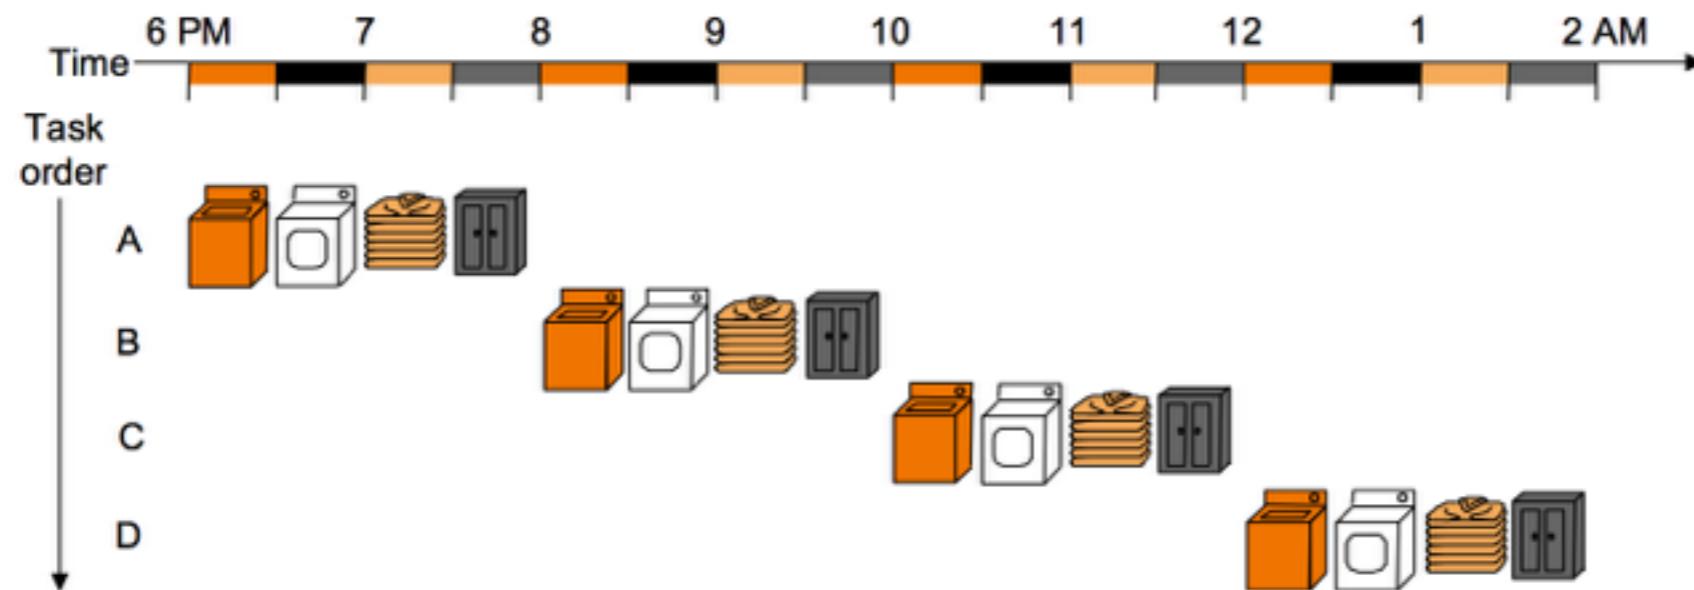
Wake up! Time to do laundry!

The Laundry Analogy

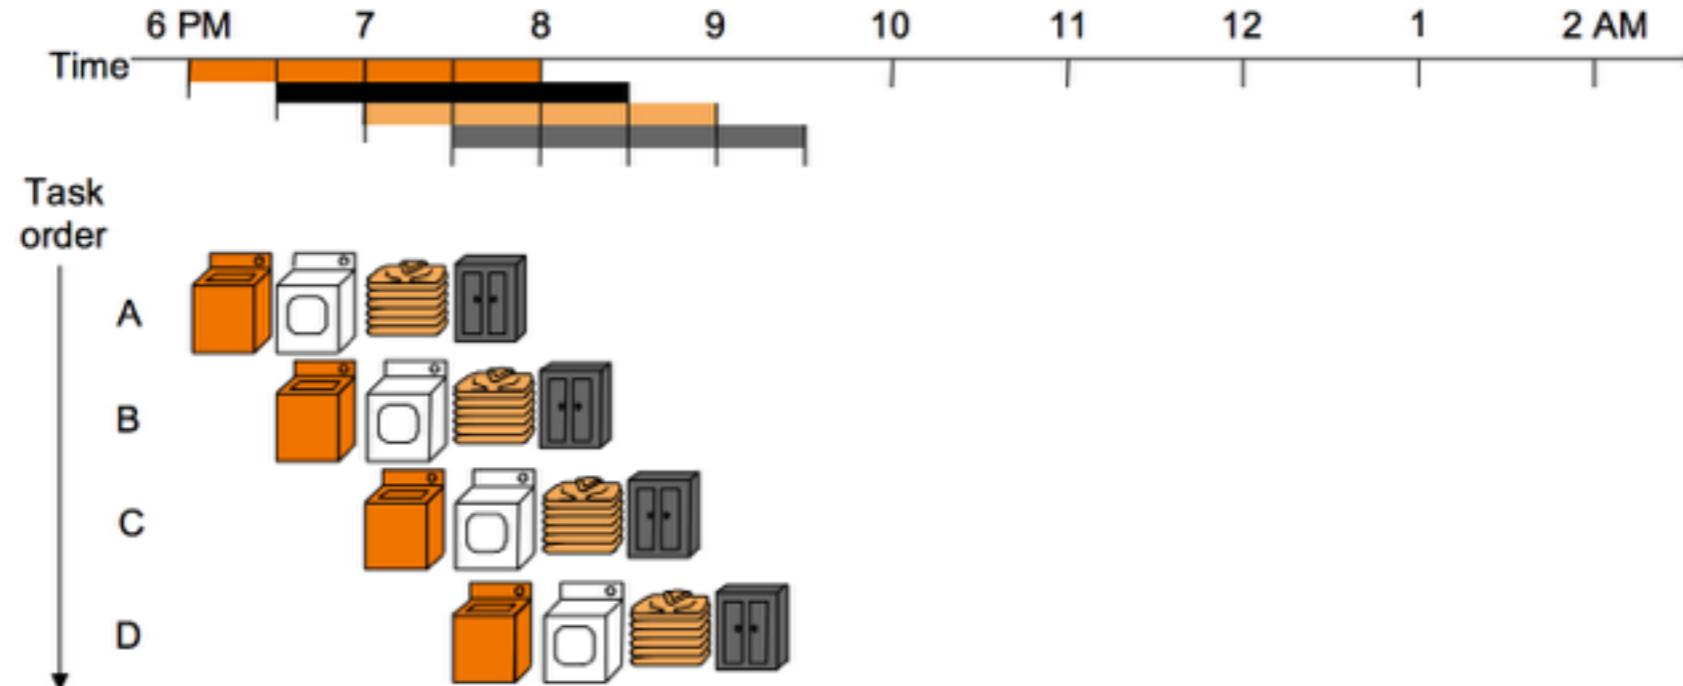
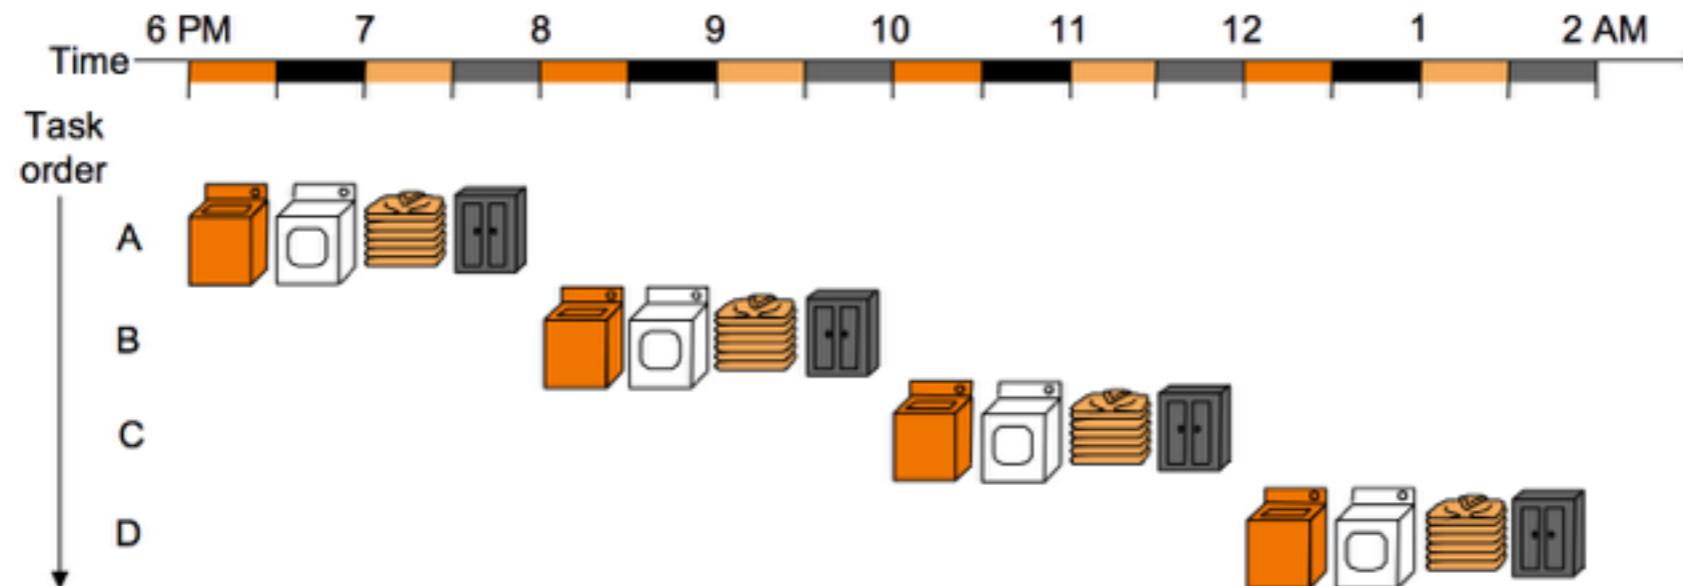


- Place one dirty load of clothes in the washer
- When the washer is finished, place the wet load in the dryer
- When the dryer is finished, take out the dry clothes and fold
- When folding is finished, ask your roommate (?) to put the clothes away

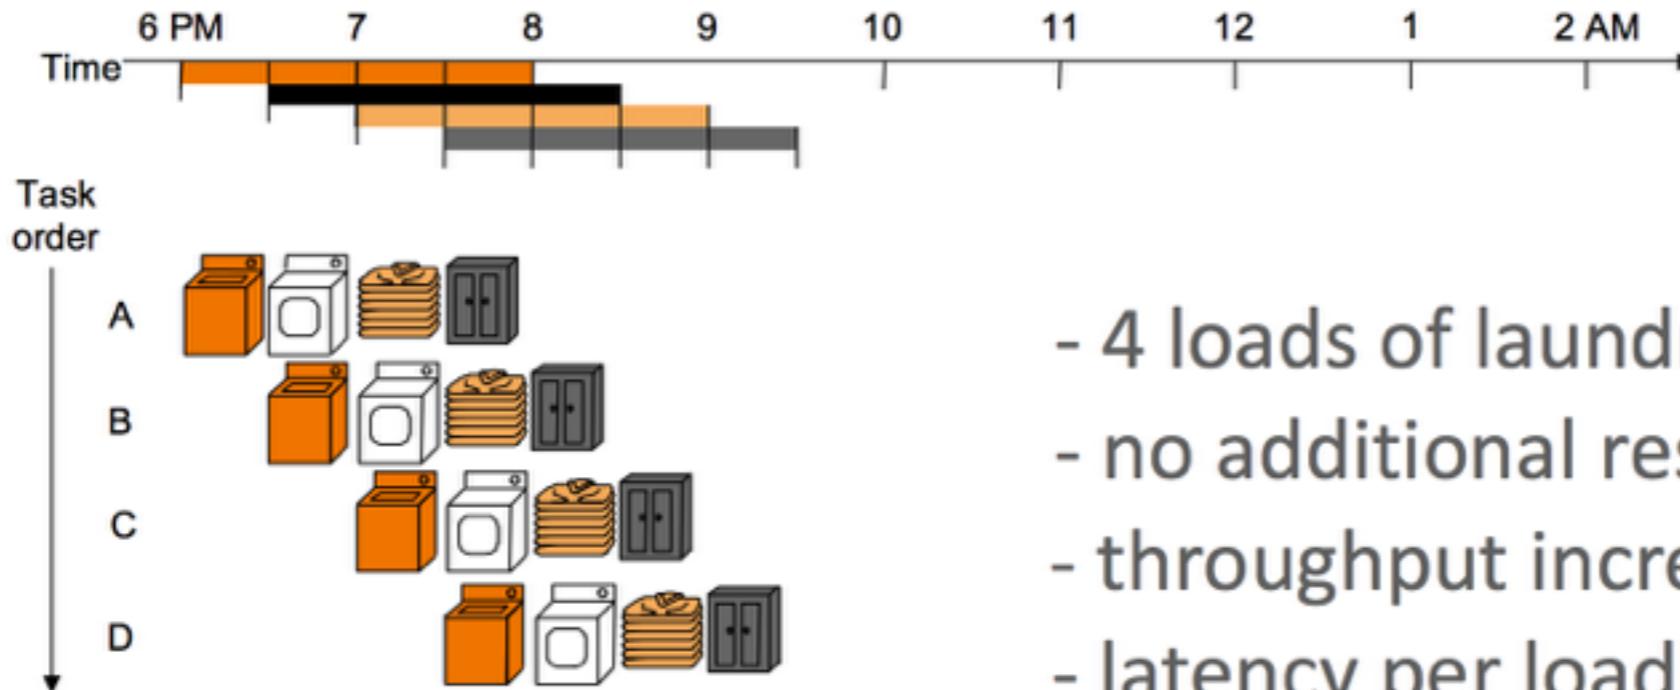
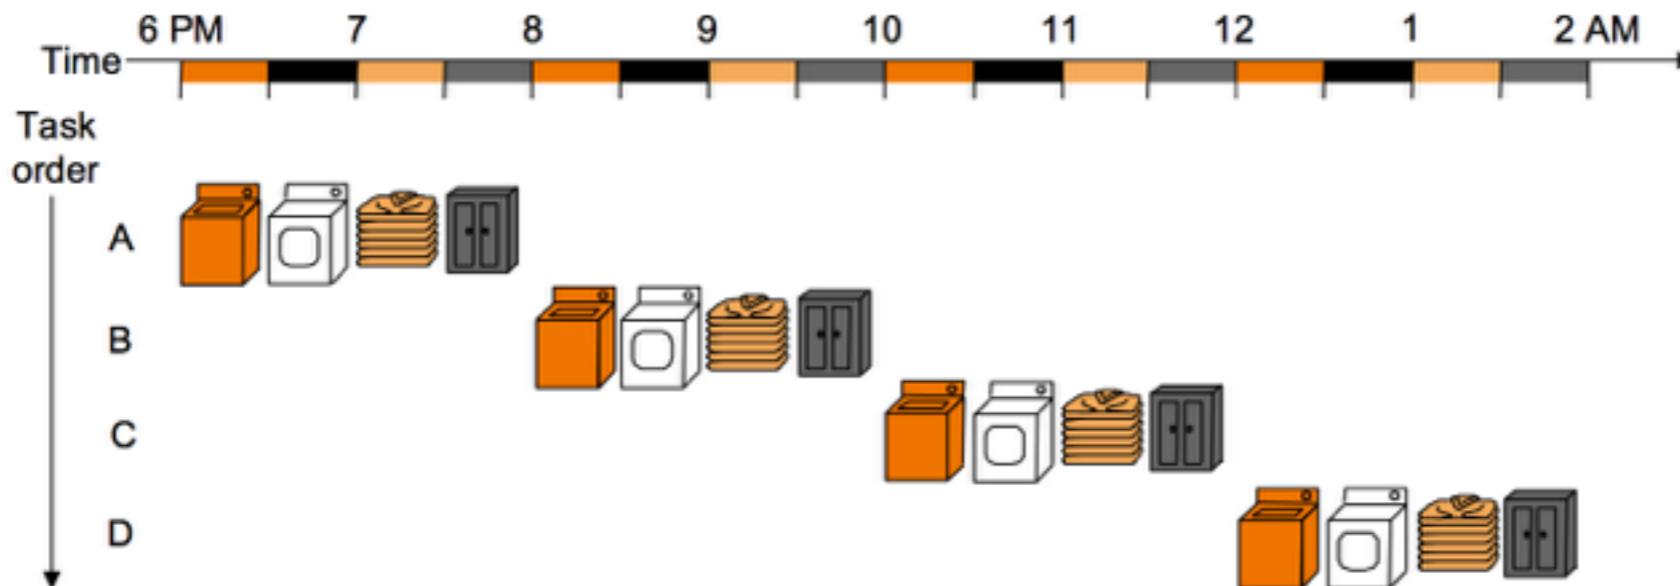
Pipelining Multiple Loads of Laundry



Pipelining Multiple Loads of Laundry



Pipelining Multiple Loads of Laundry



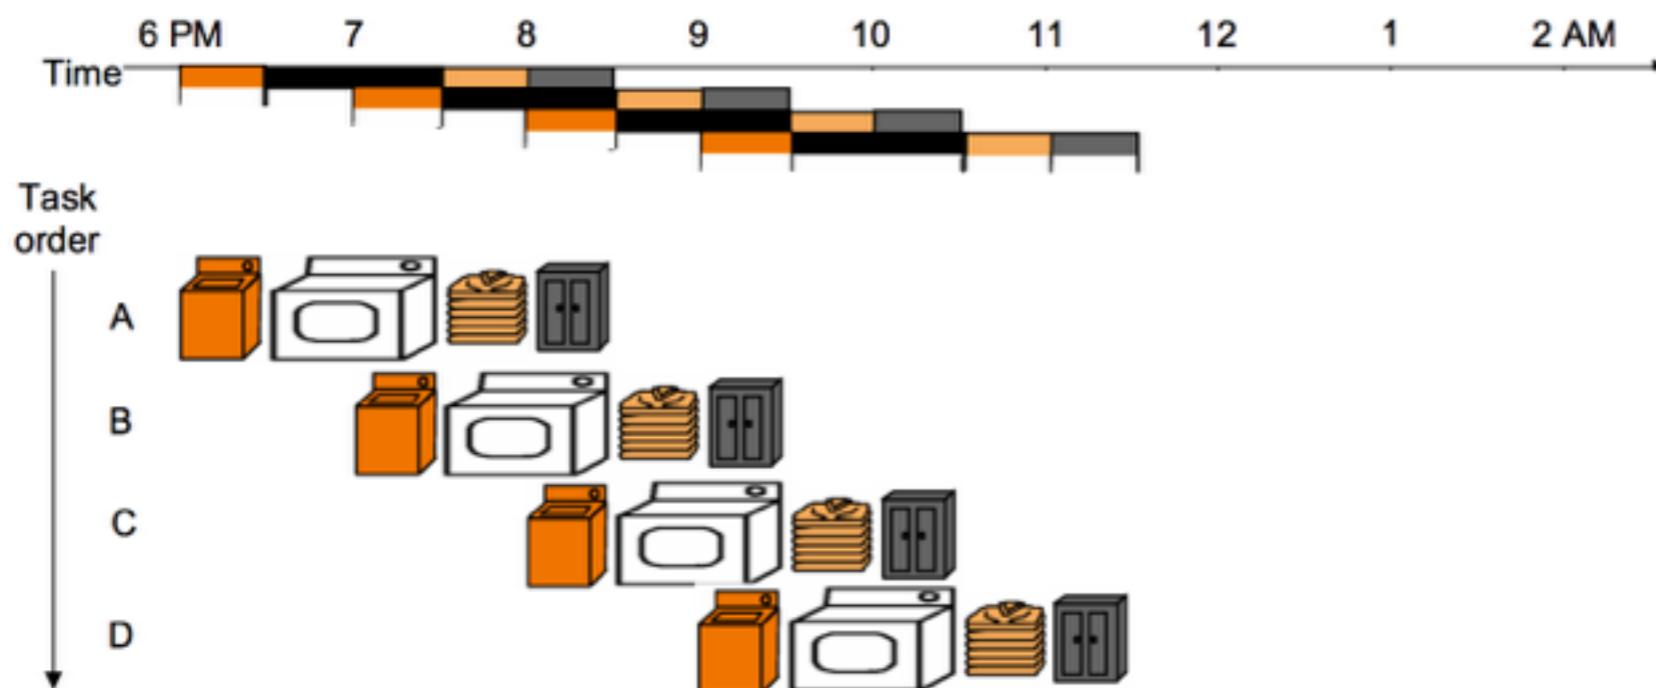
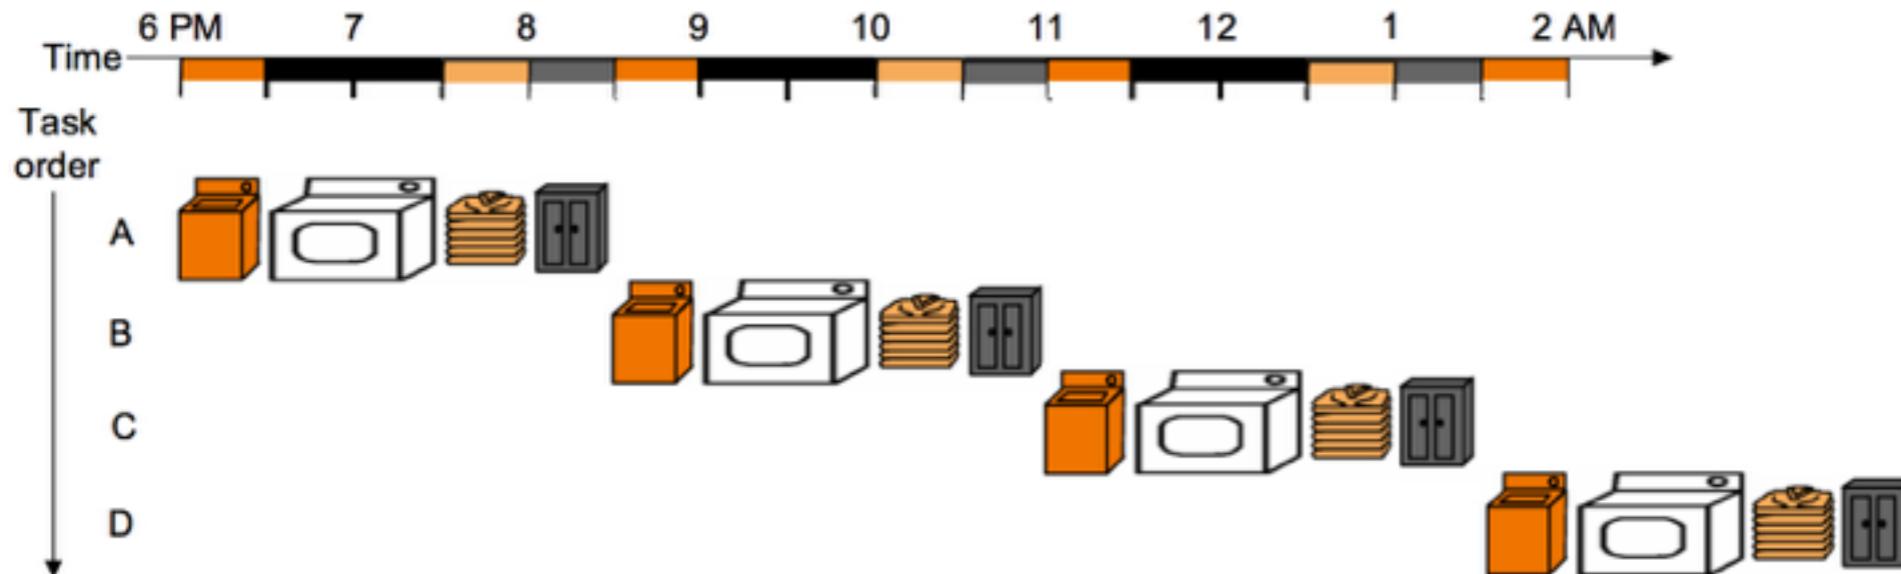
- 4 loads of laundry in parallel
- no additional resources
- throughput increased by 4
- latency per load is the same

Question: We have a four-stage pipeline. Every stage takes 1 hour. How long does it take to finish 100 loads?

Question: We have a four-stage pipeline. Every stage takes 1 hour. How long does it take to finish 100 loads?

Question: We have a four-stage pipeline. Every stage takes 1 hour. How long does it take to finish **N** loads?

In Reality, maybe...



the slowest step decides throughput

Pipelining is Everywhere

Build
frame

Install
parts

Paint

Build
frame

Install
parts

Paint

Build
frame

Install
parts

Paint



Pipelining is Everywhere

Build
frame

Install
parts

Paint

Build
frame

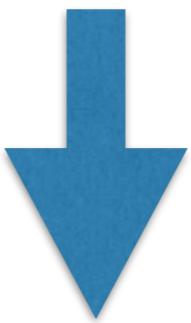
Install
parts

Paint

Build
frame

Install
parts

Paint



Build
frame

Install
parts

Paint

Build
frame

Install
parts

Paint

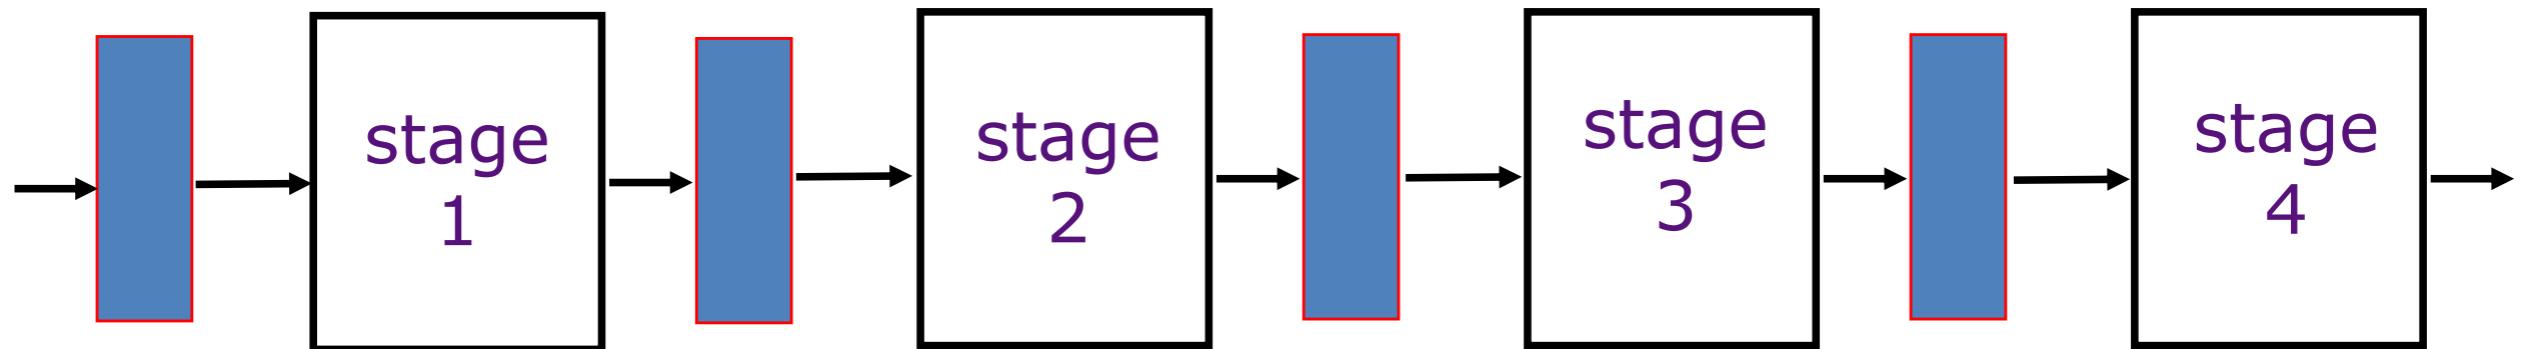
Build
frame

Install
parts

Paint

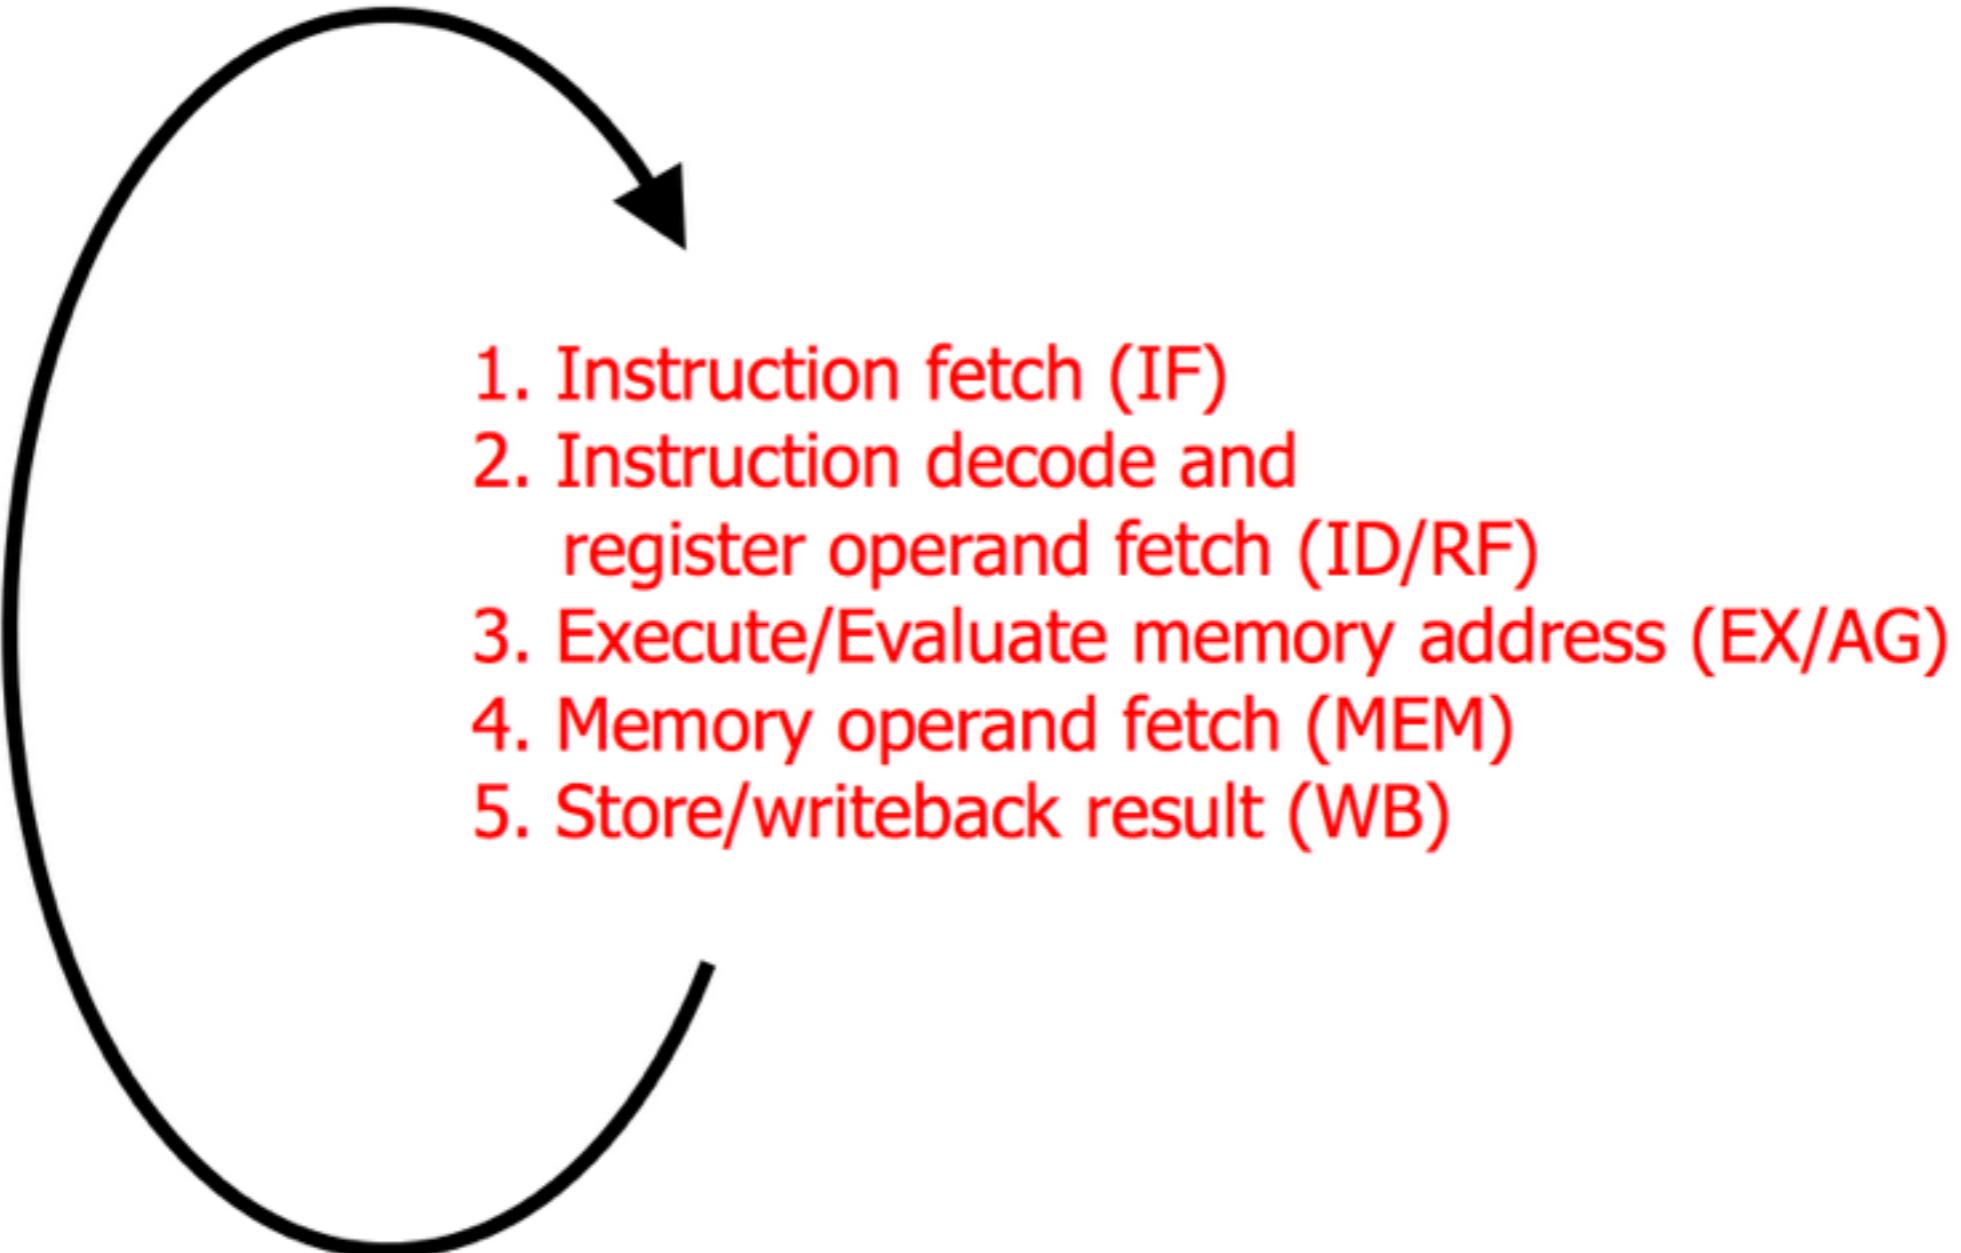


An Ideal Pipeline

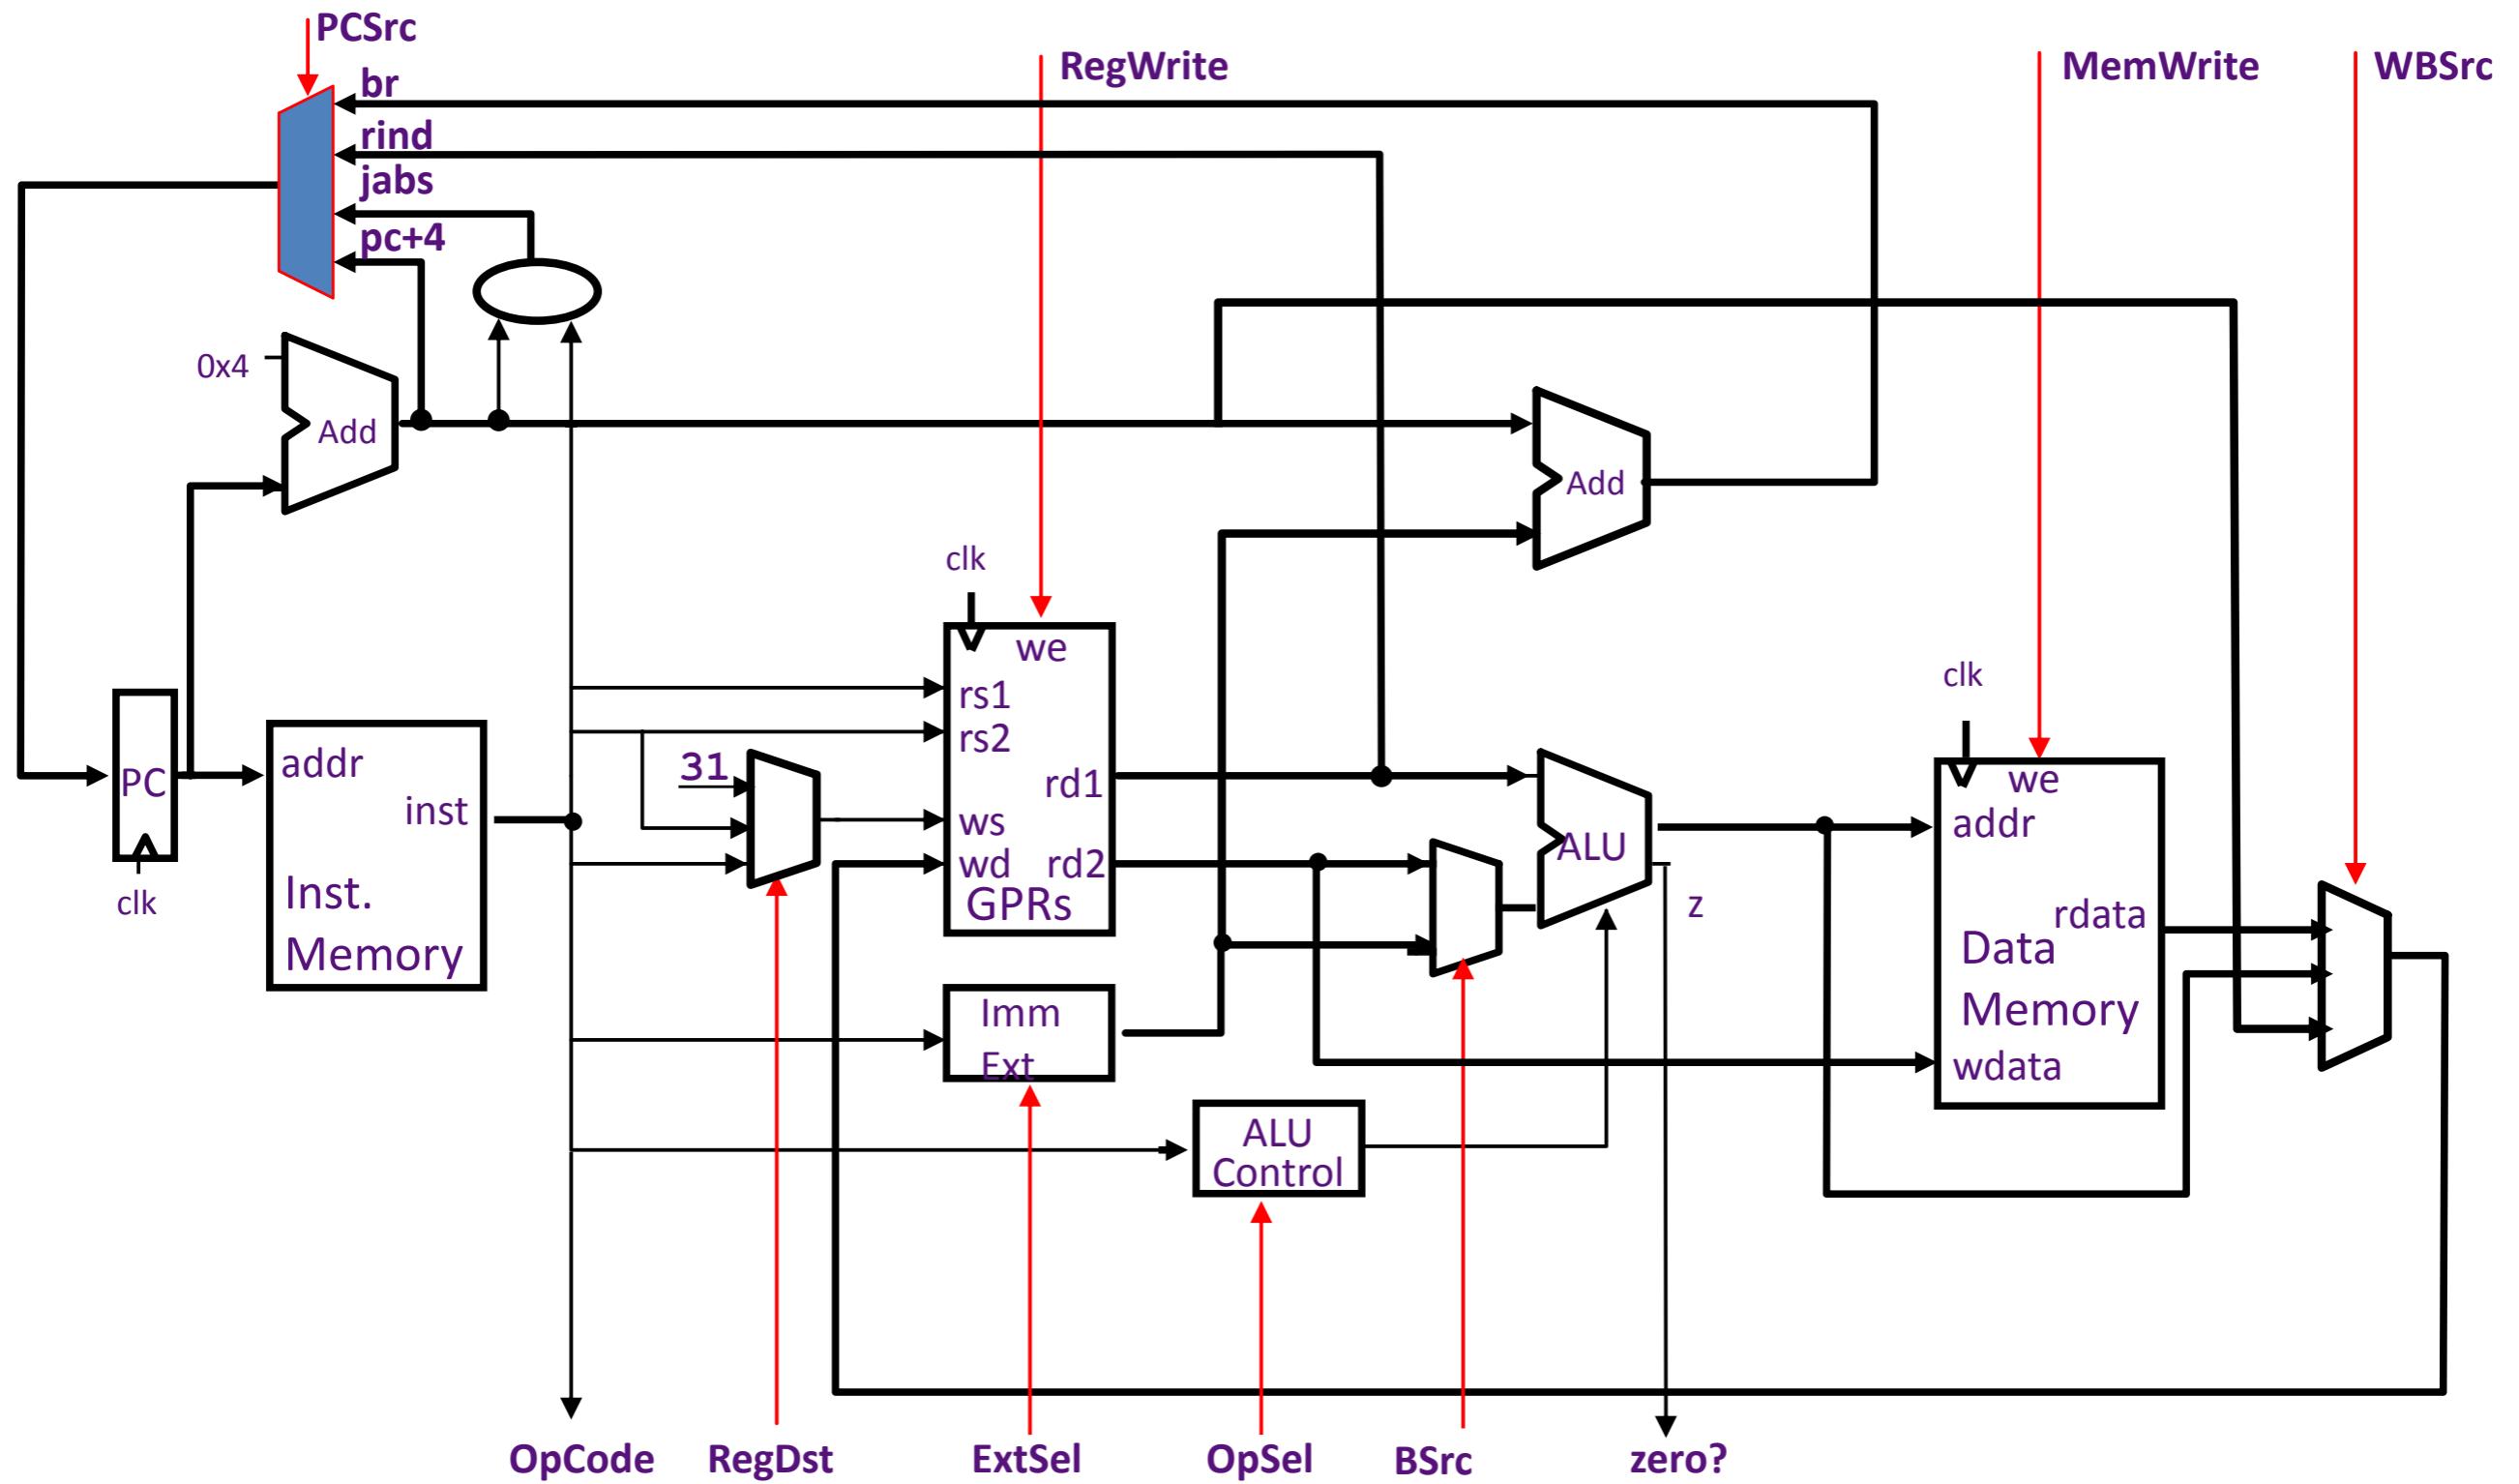


- All objects go through the same stages
- No sharing of resources between any two stages
- Propagation delay through all pipeline stages is equal
- Scheduling of a transaction entering the pipeline is not affected by the transactions in other stages

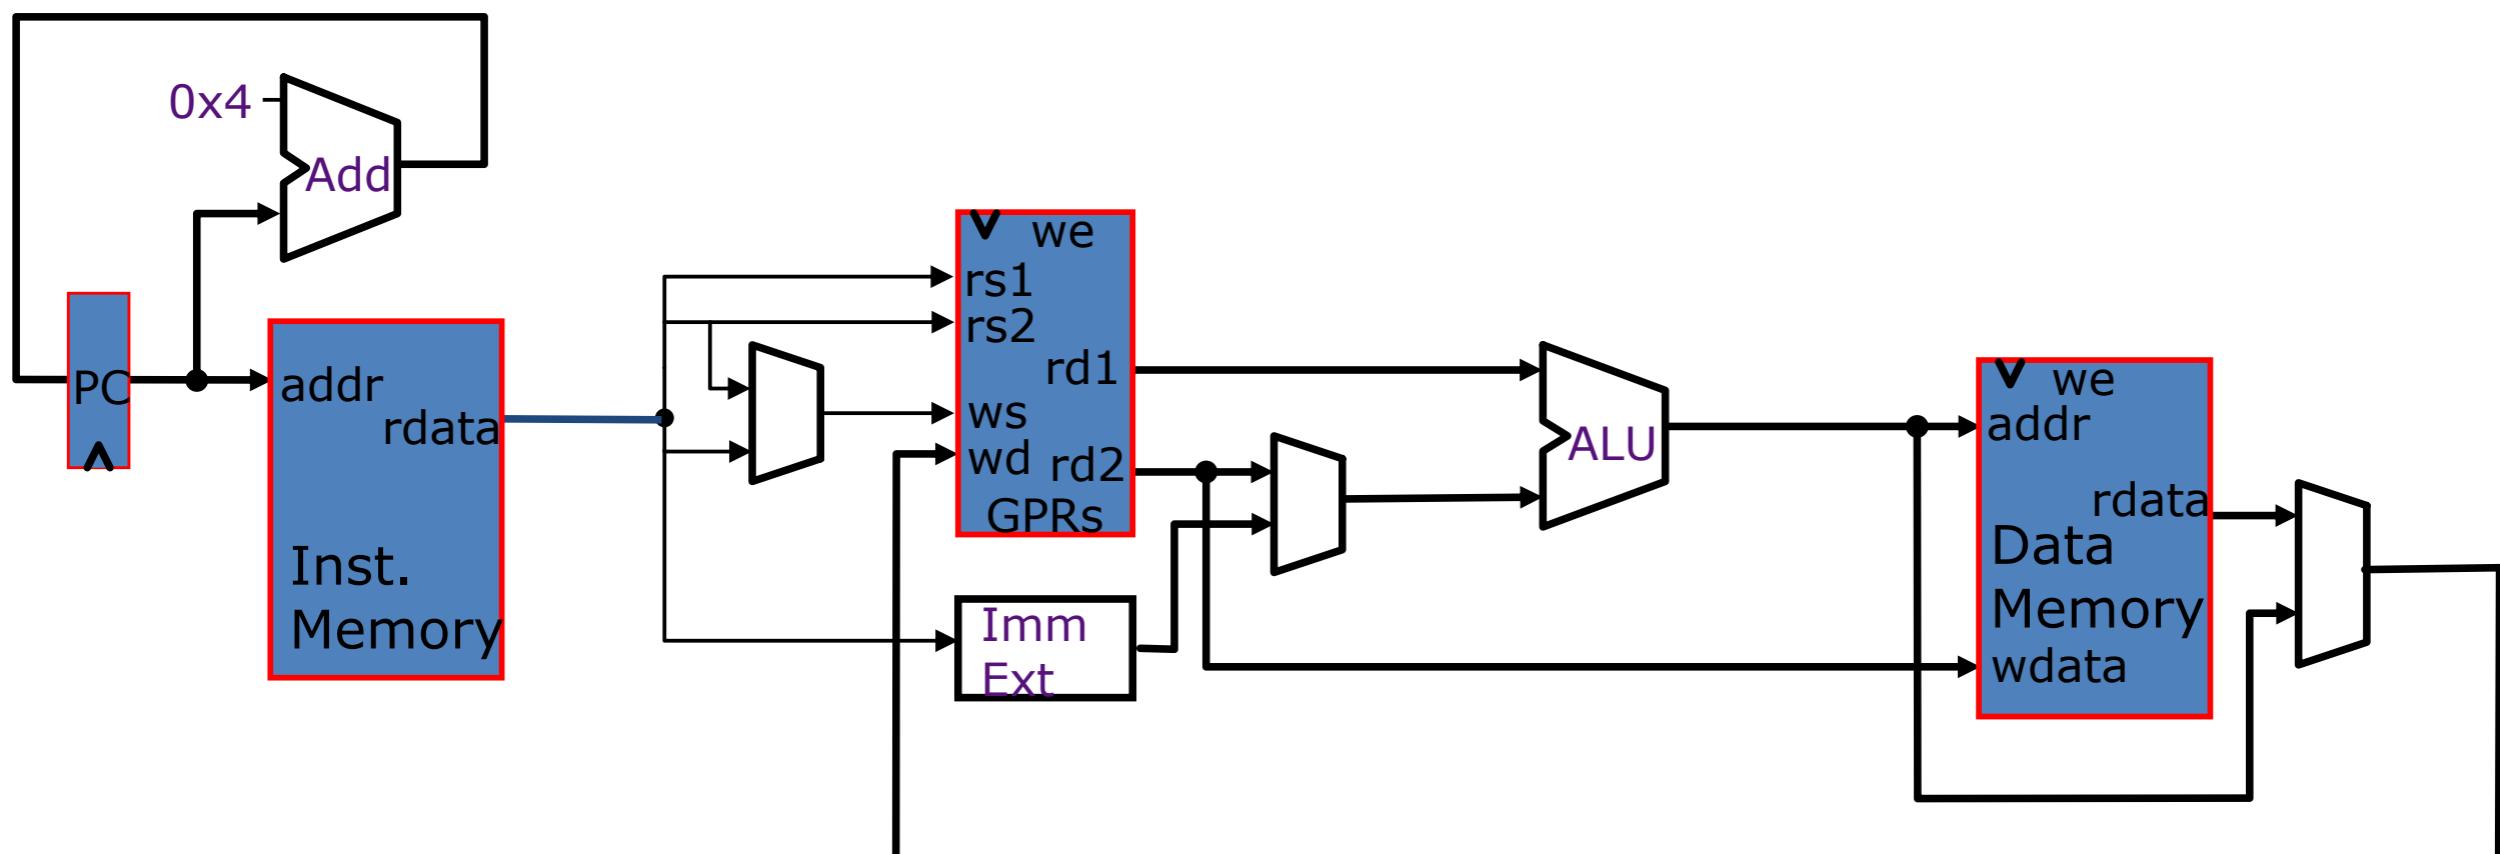
The Instruction Execution Cycle



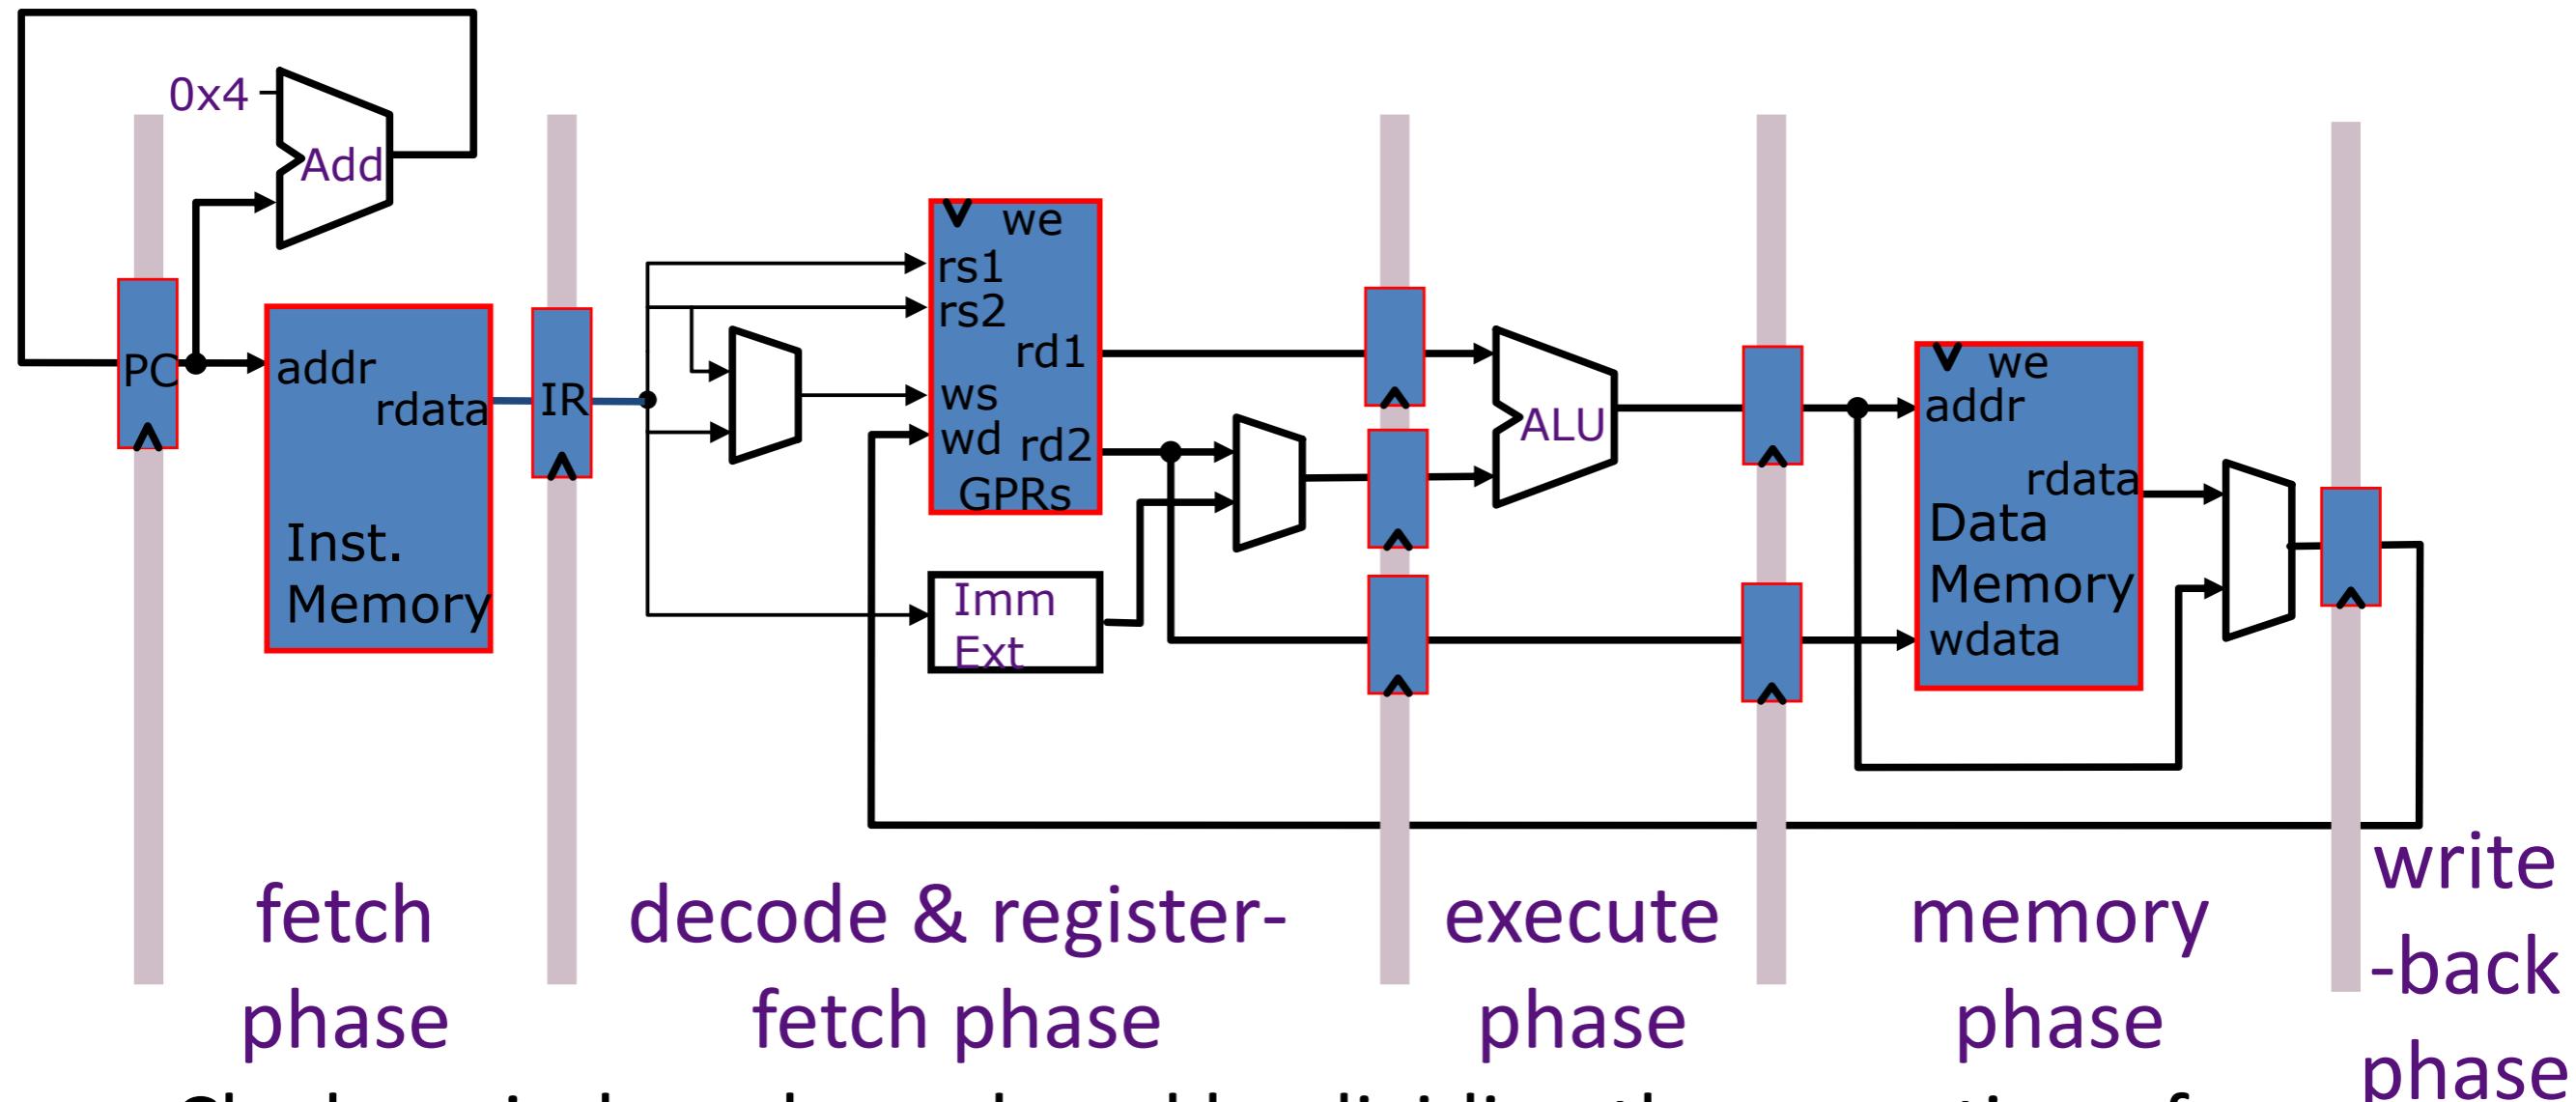
Unpipelined Datapath for MIPS



Simplified Unpipelined Datapath



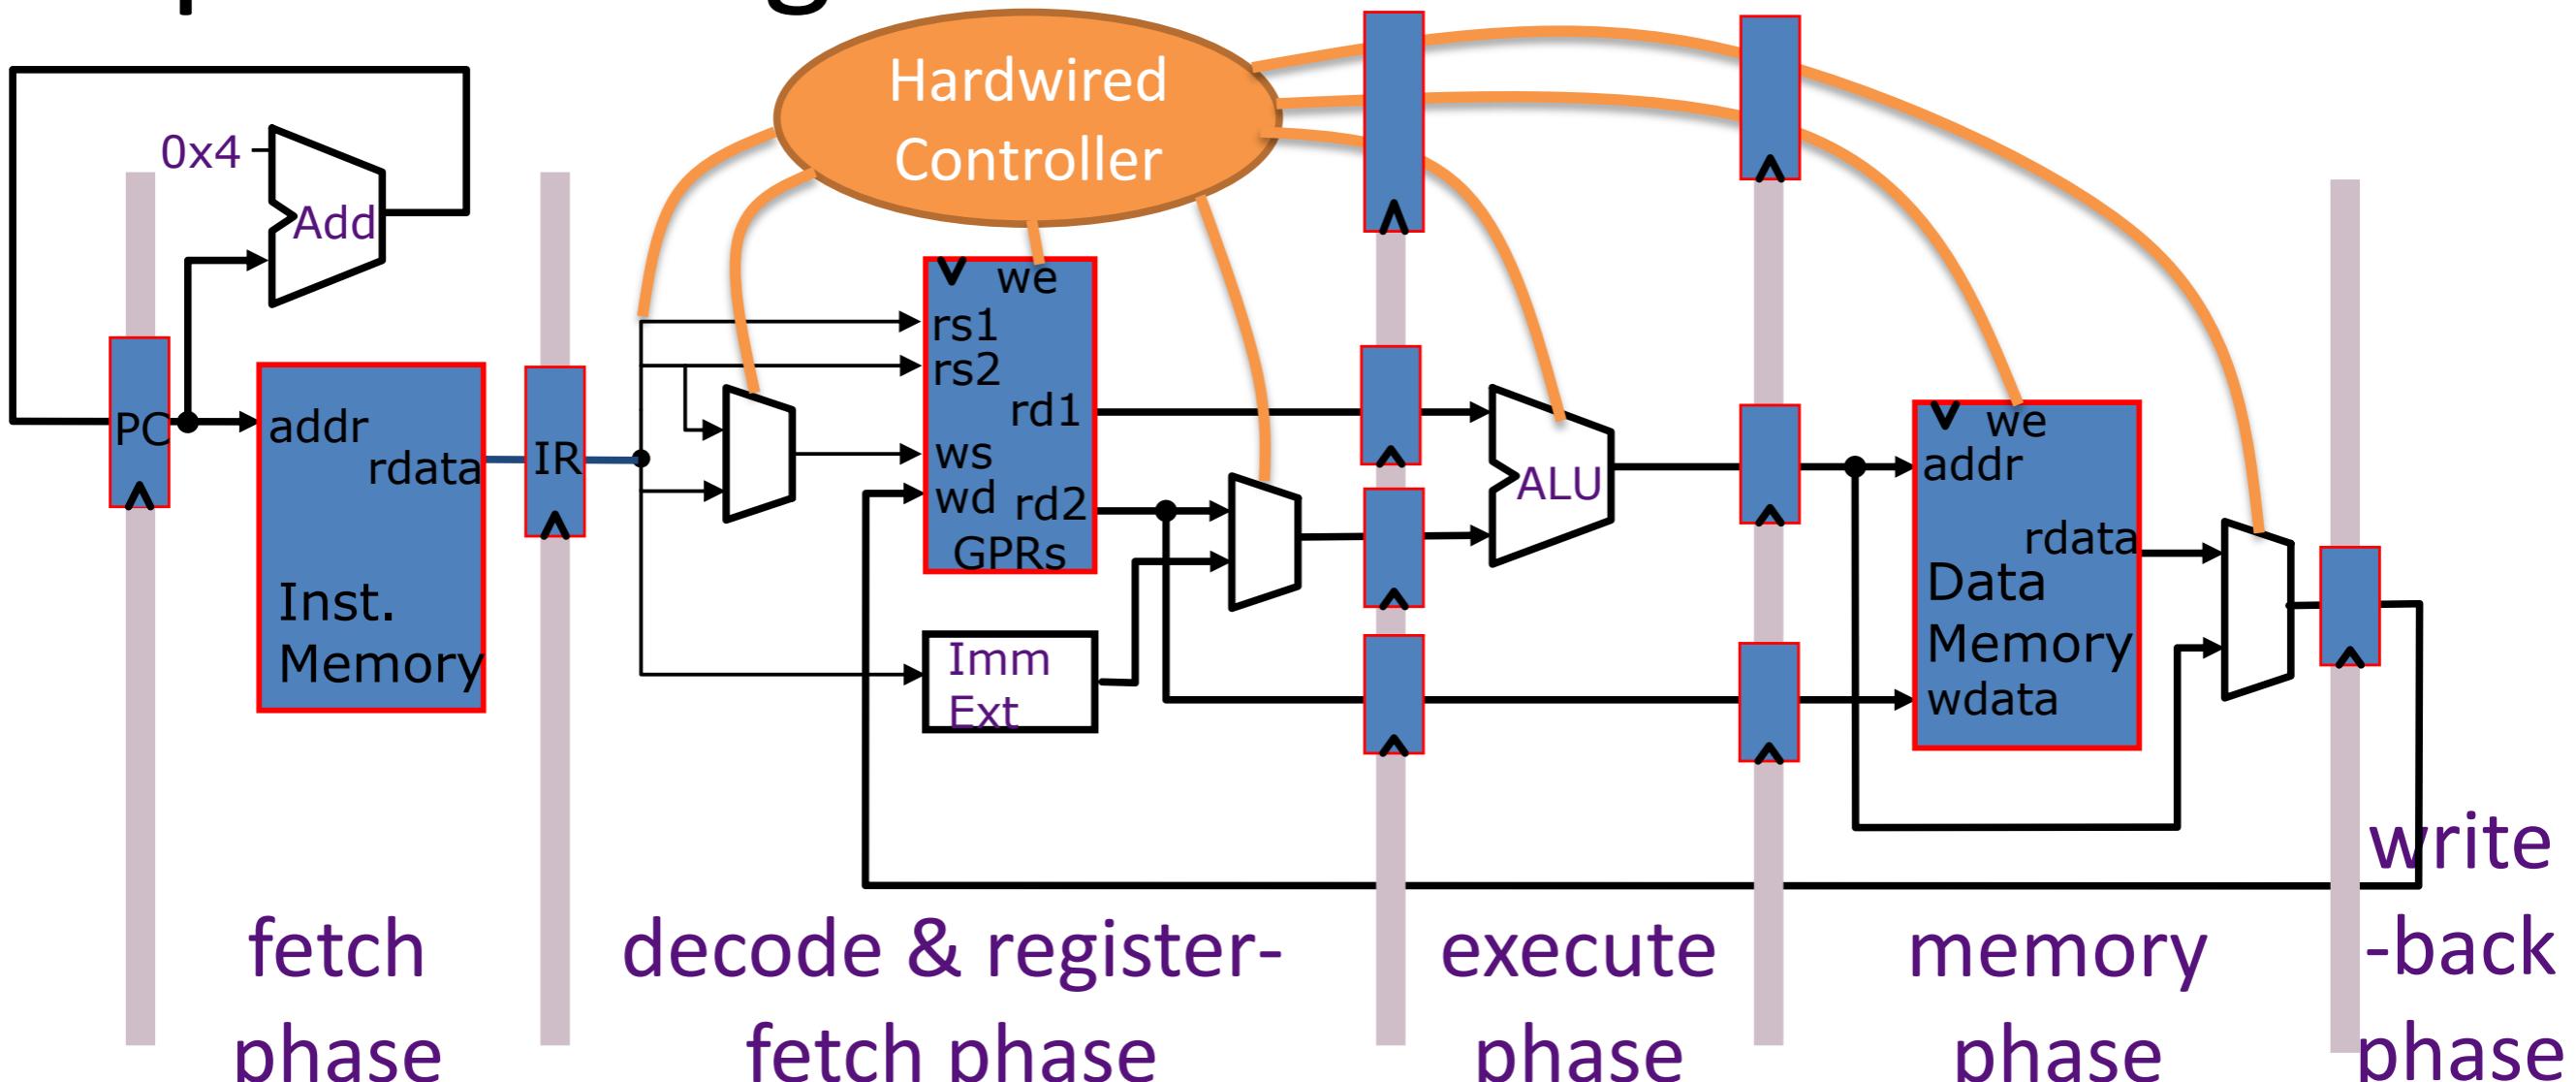
Pipelined Datapath



Clock period can be reduced by dividing the execution of an instruction into multiple cycles

$$t_C > \max \{t_{IM}, t_{RF}, t_{ALU}, t_{DM}, t_{RW}\} \quad (= t_{DM} \text{ probably})$$

Pipeline Diagrams: Transactions vs. Time



time

instruction1

instruction2

instruction3

instruction4

instruction5

t0
 IF_1

t1
 ID_1

t2
 EX_1

t3
 MA_1

t4
 WB_1

t5
 MA_2

t6
 WB_2

t7
 MA_3

WB3

MA4

WB4

...

IF_2

ID_2

EX_2

MA_2

WB_2

MA_3

WB_3

MA_4

WB_4

ID_4

WB_4

ID_5

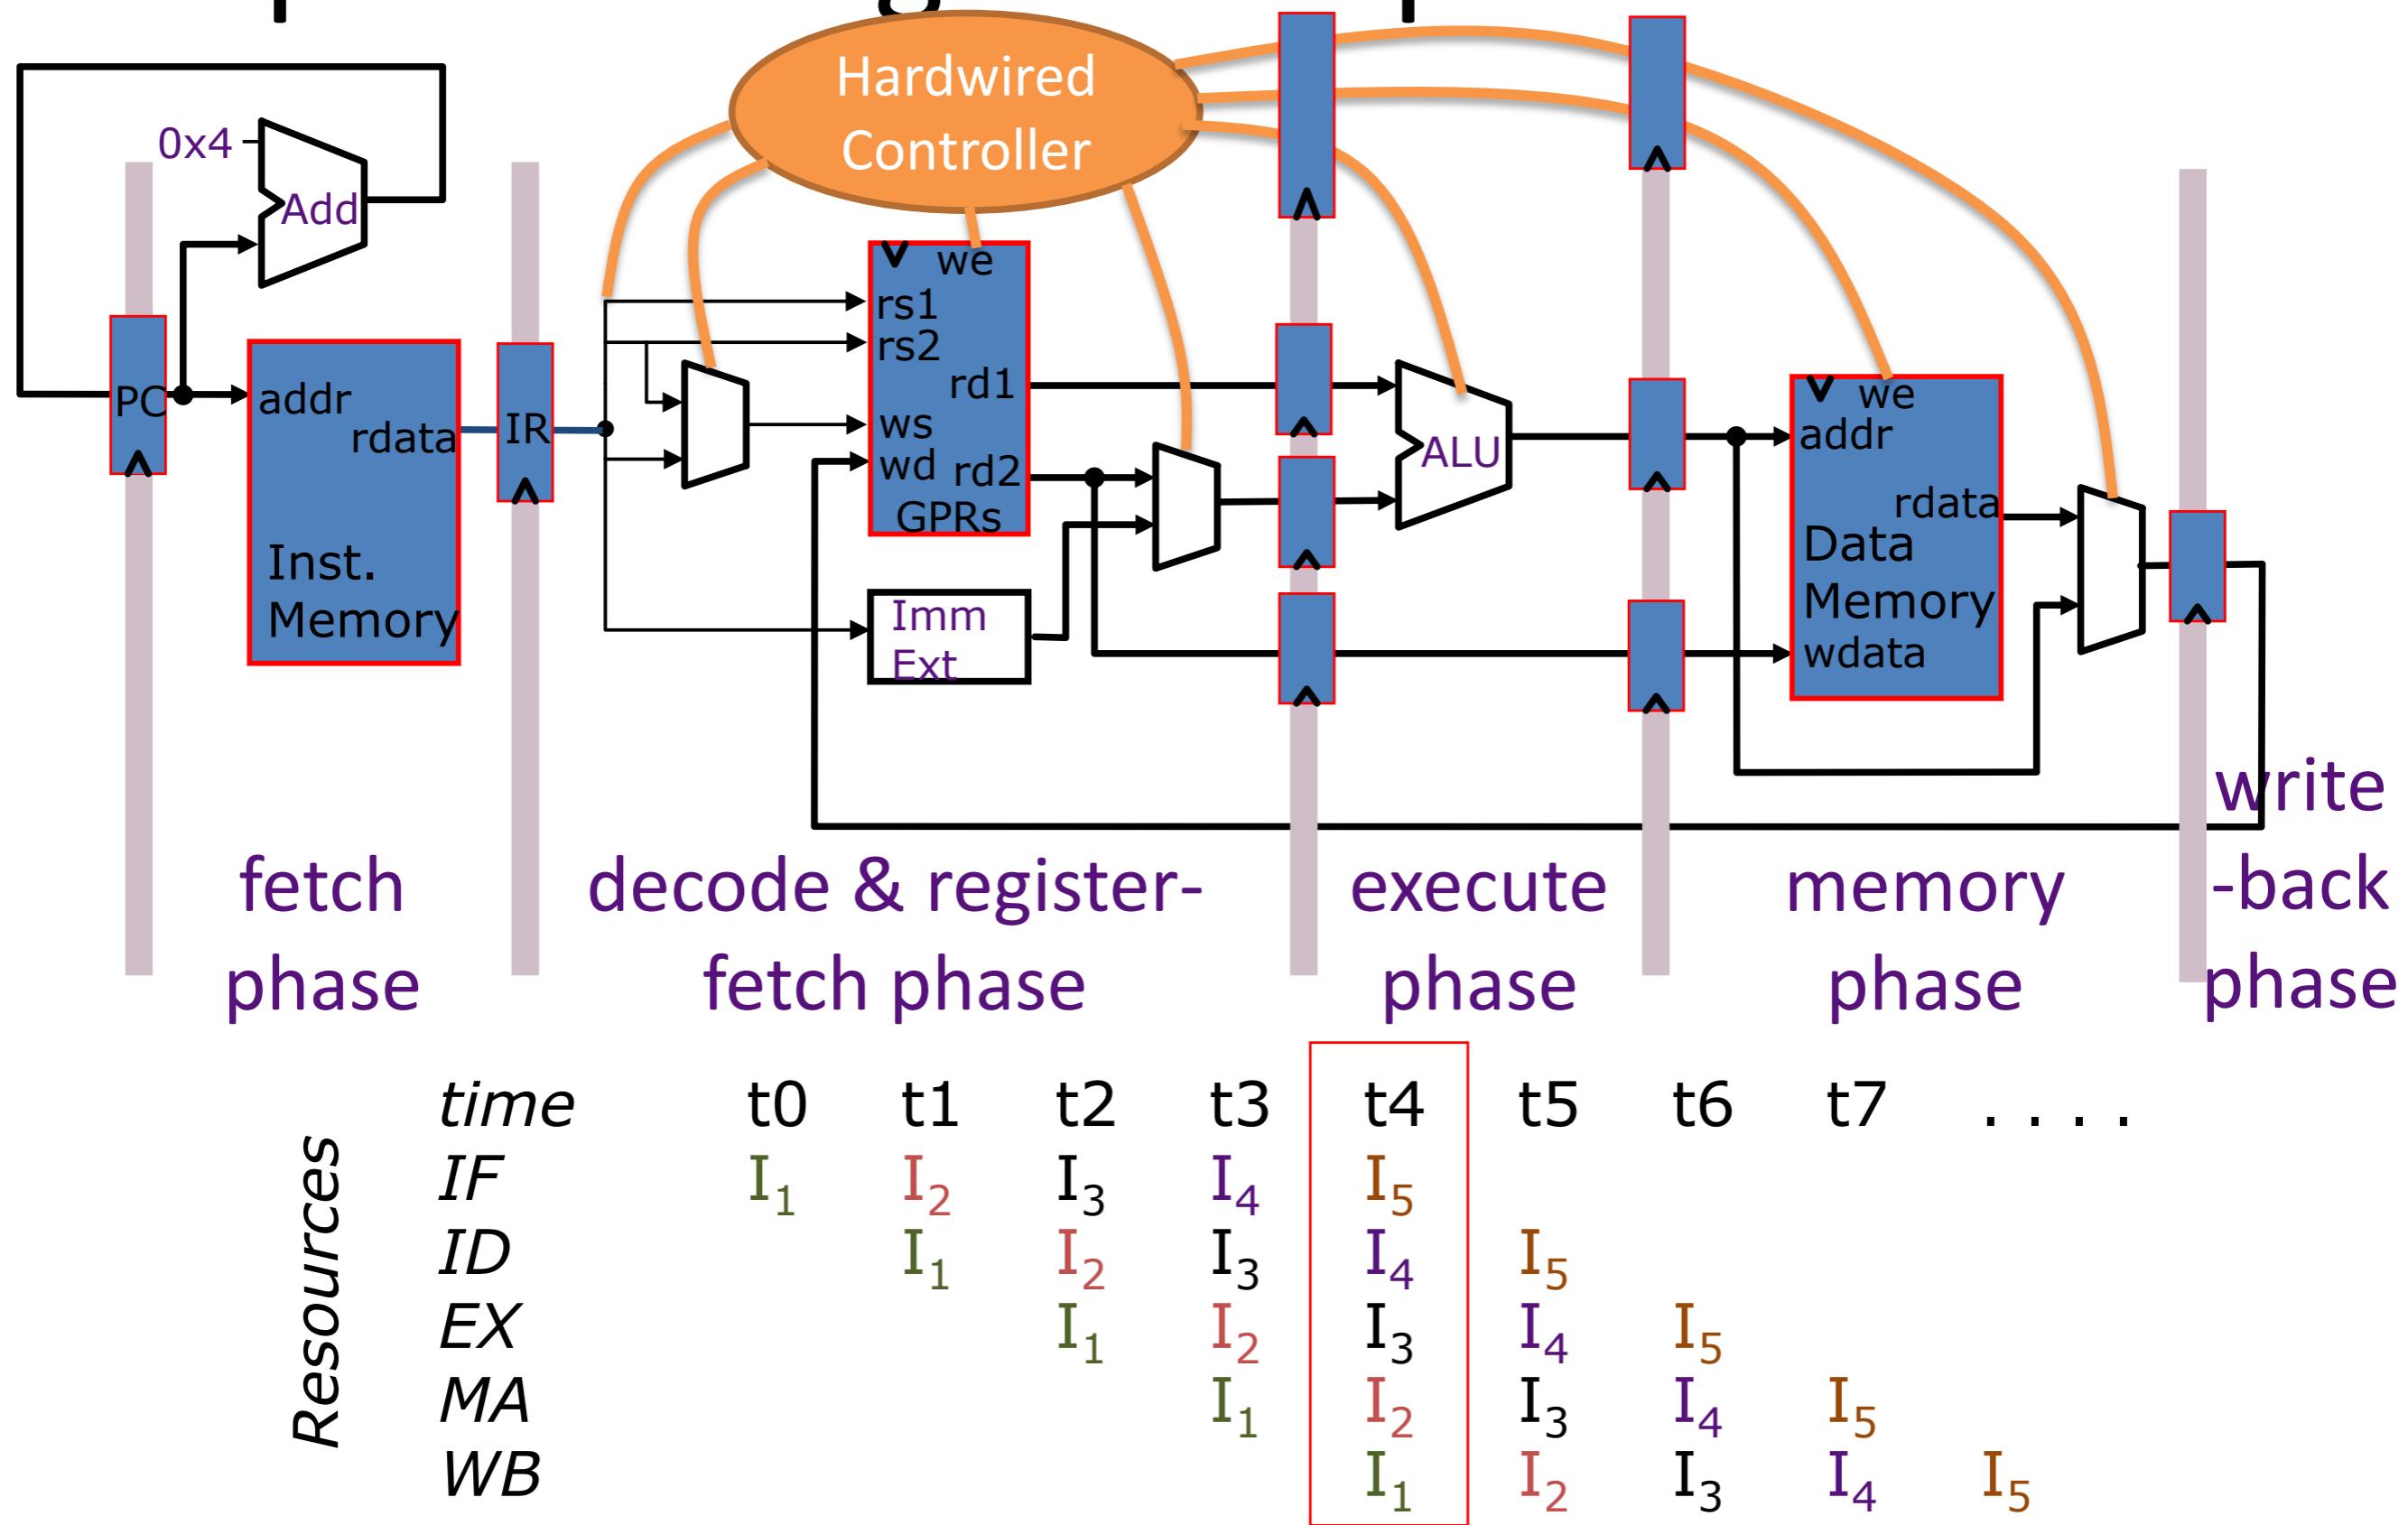
EX_5

MA_5

WB_5

IF_5

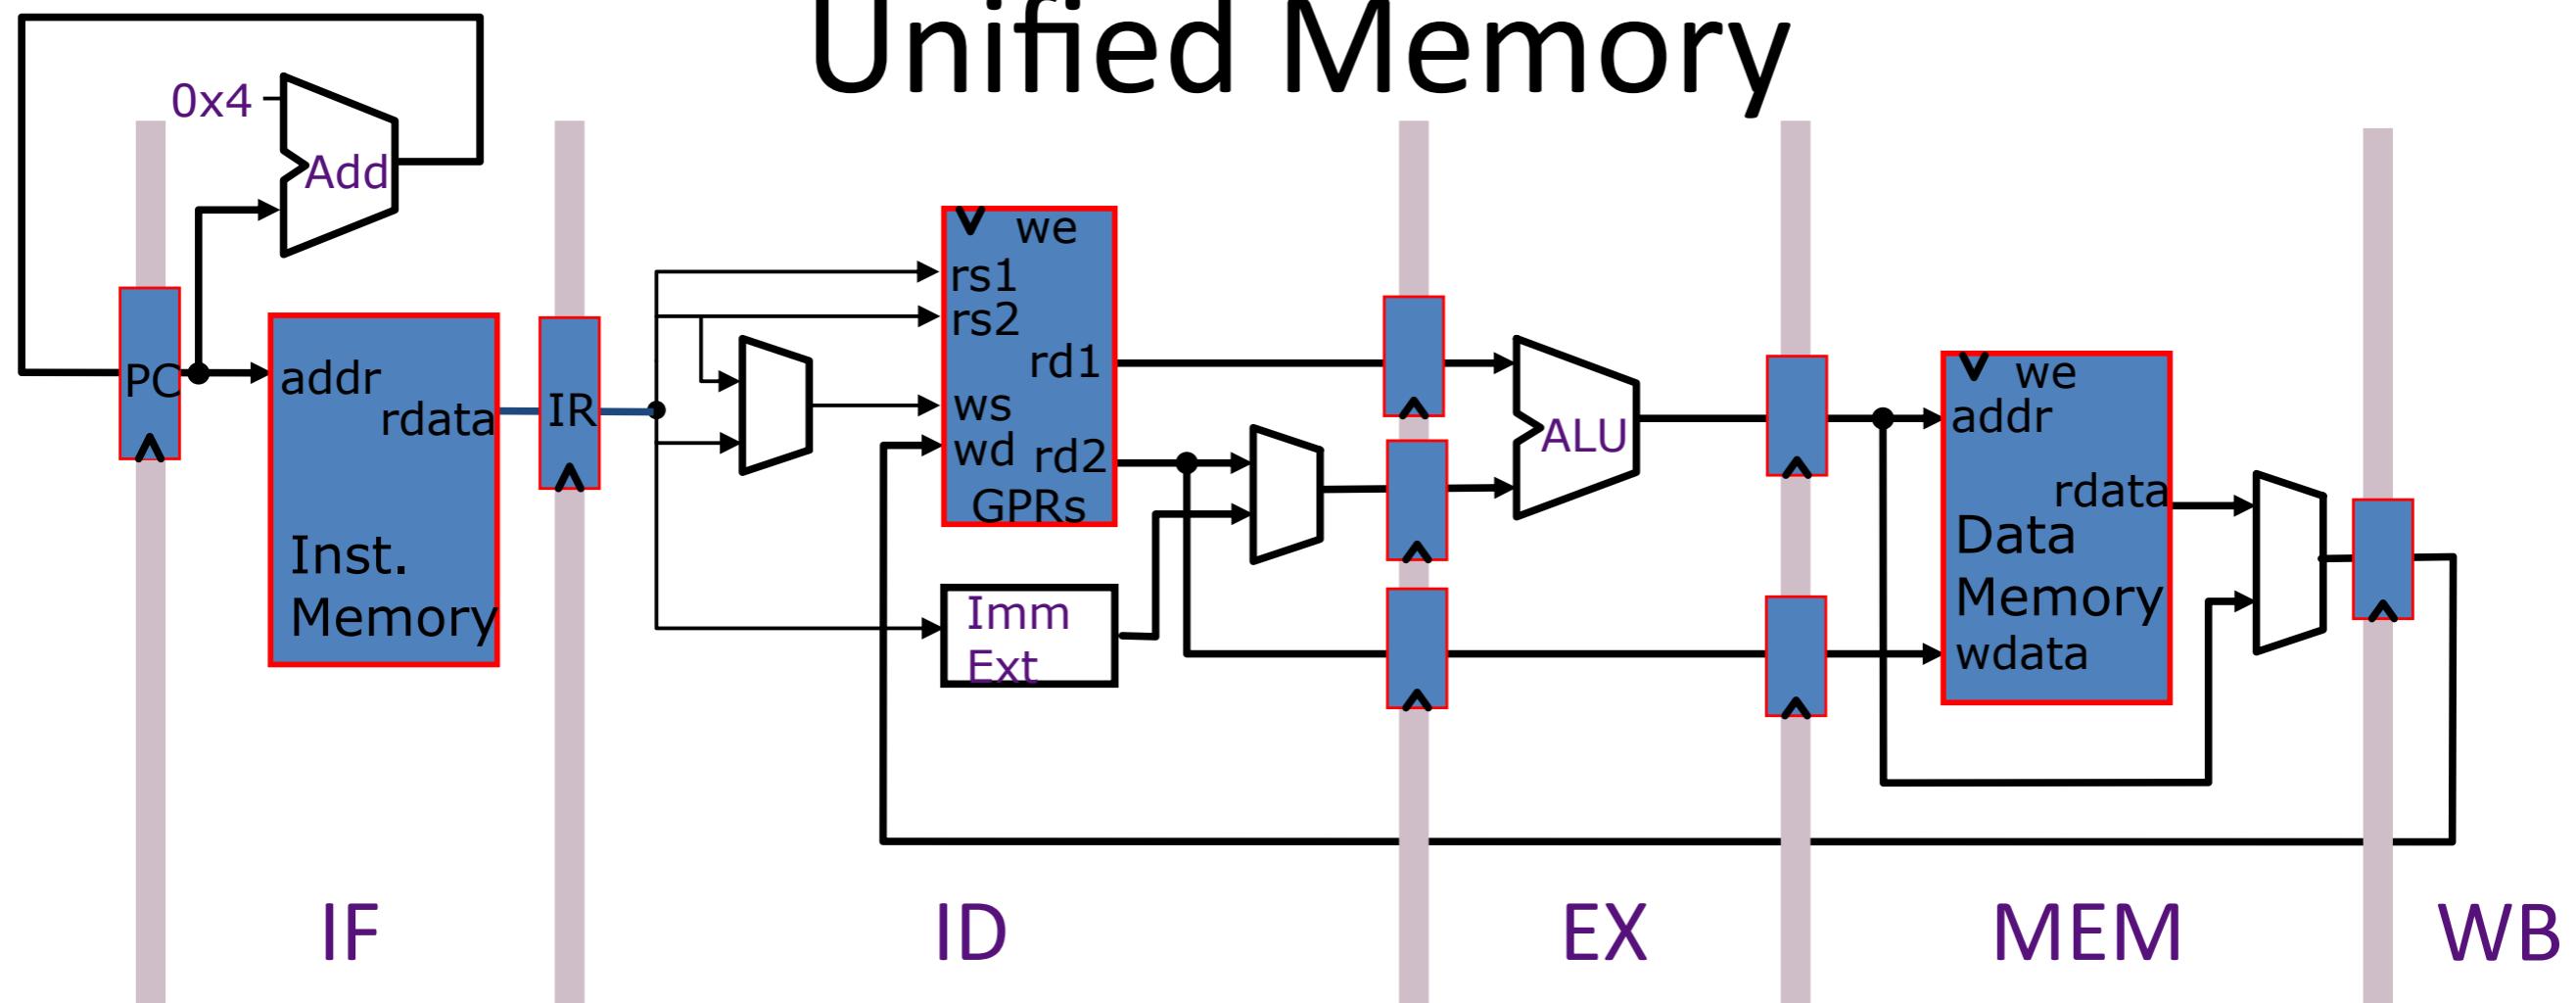
Pipeline Diagrams: Space vs. Time



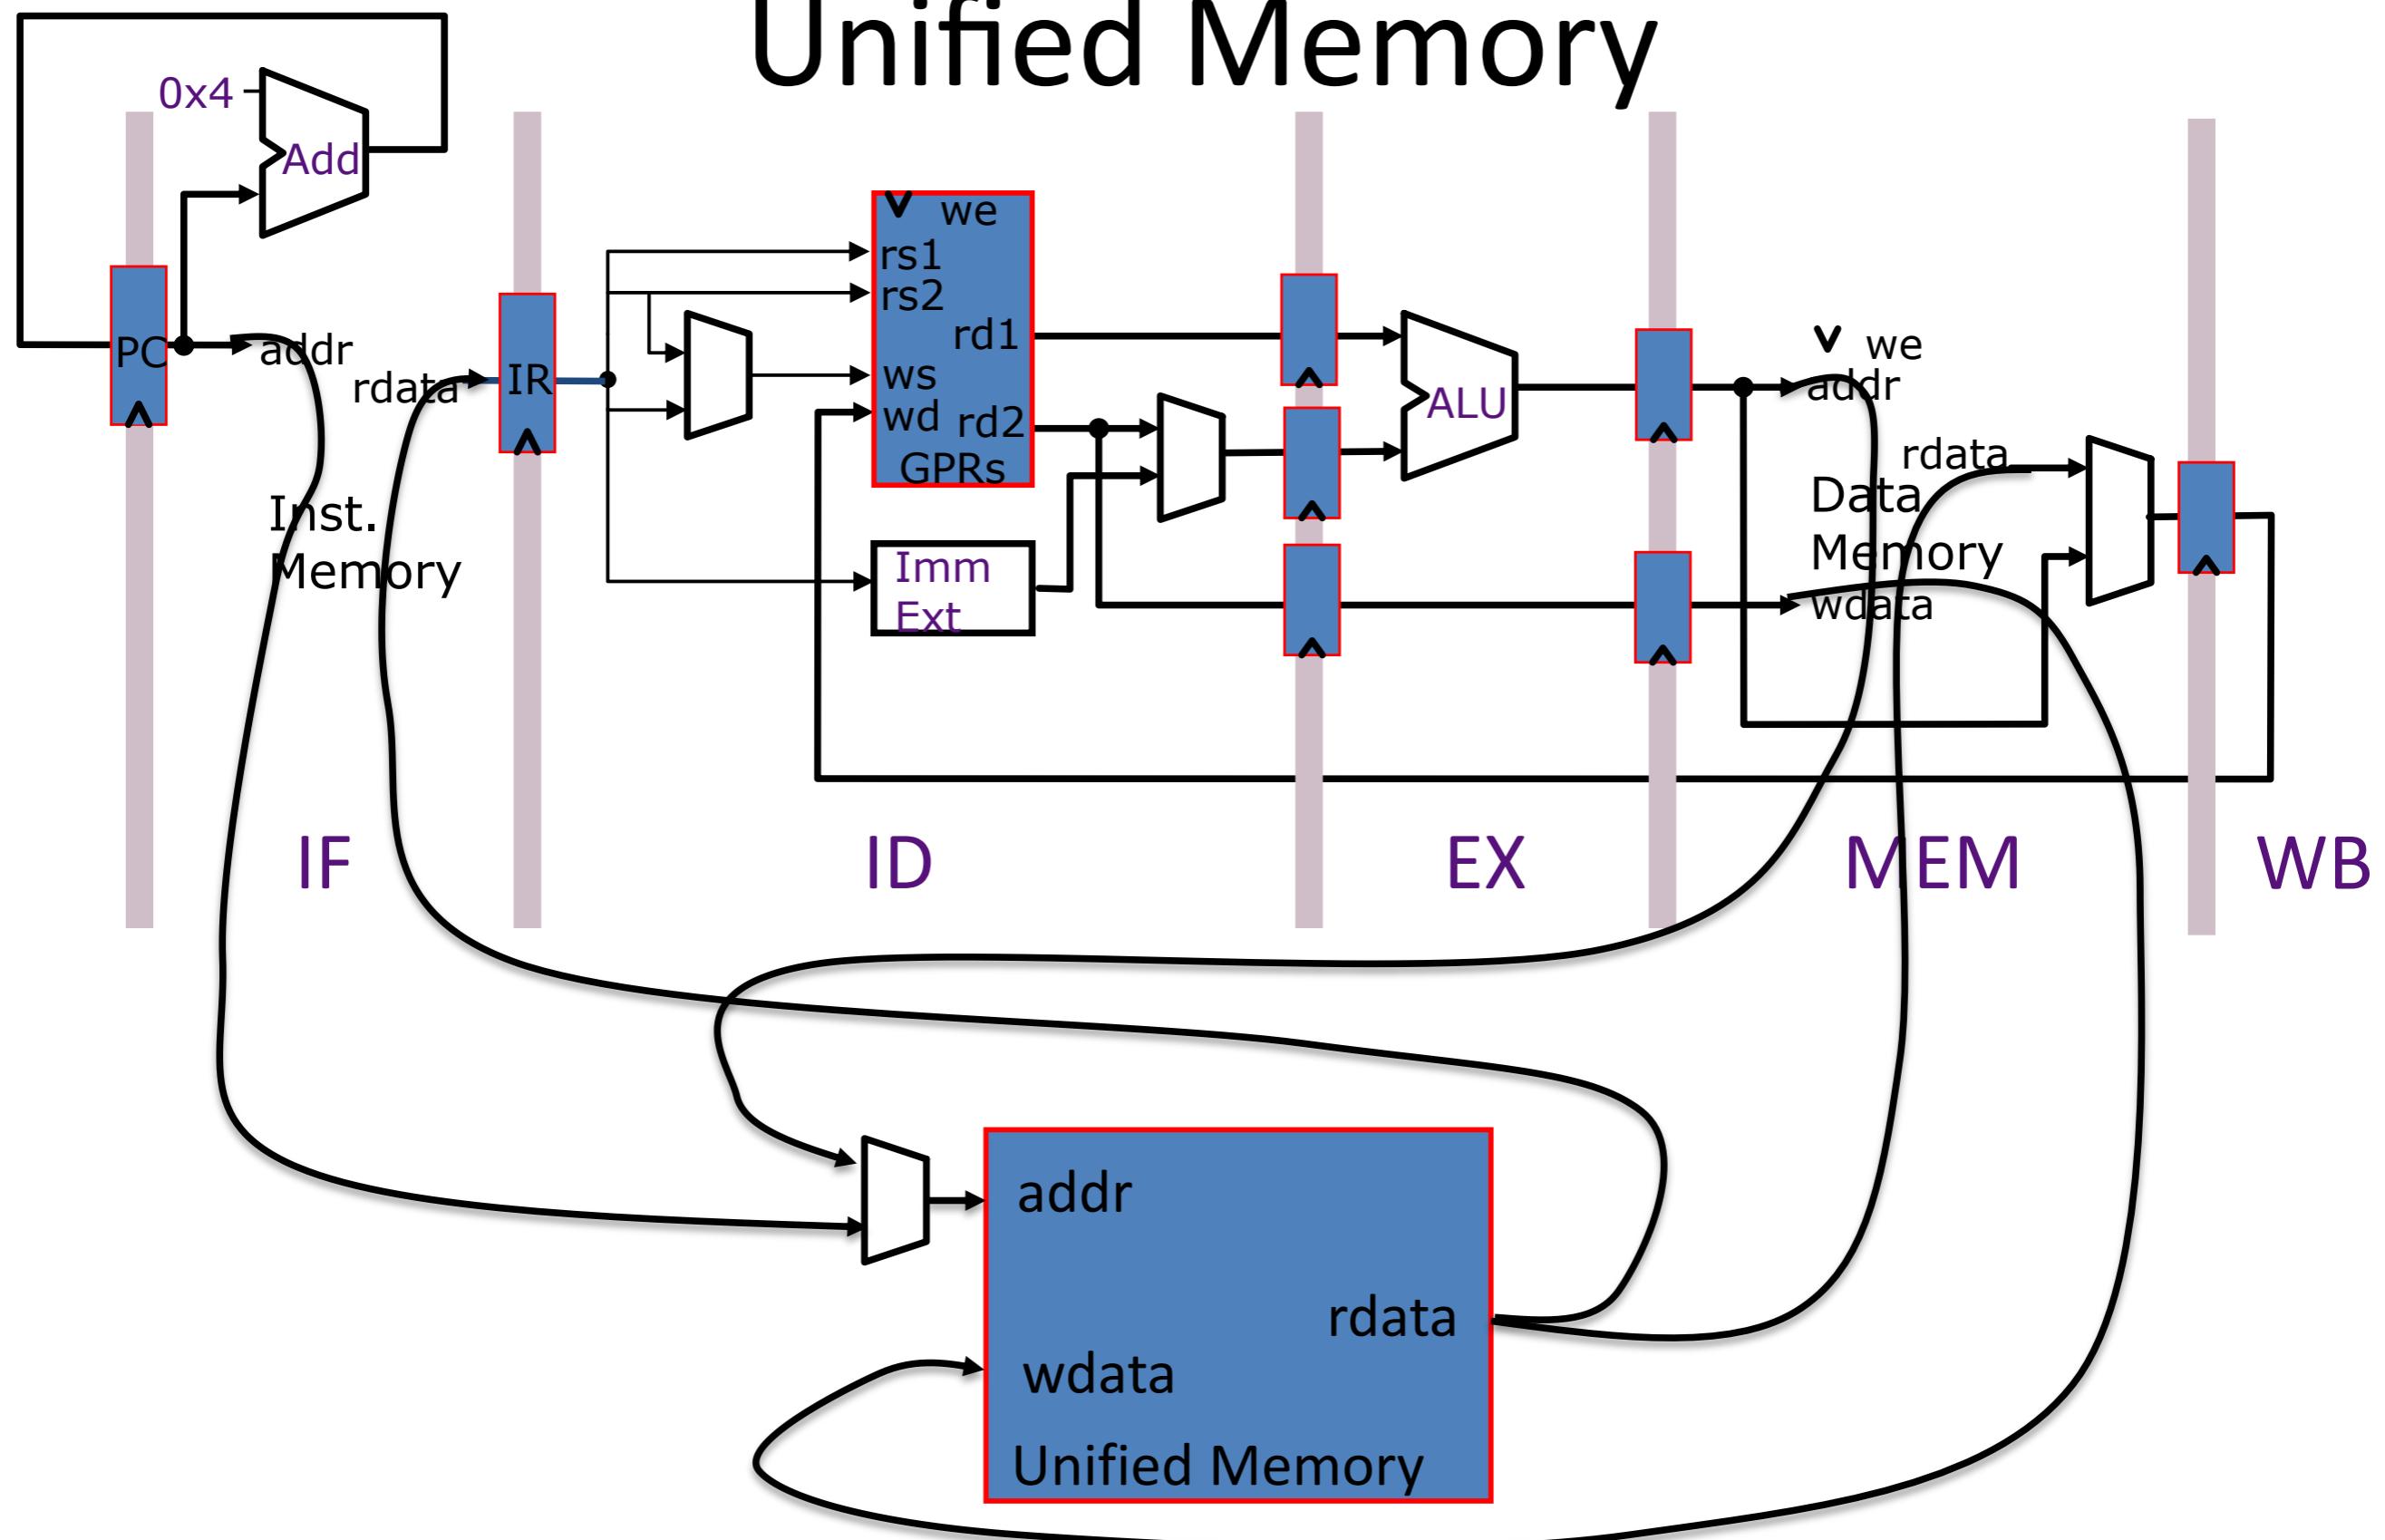
Instructions Interact With Each Other in Pipeline

- **Structural Hazard:** An instruction in the pipeline needs a resource being used by another instruction in the pipeline
- **Data Hazard:** An instruction depends on a data value produced by an earlier instruction
- **Control Hazard:** Whether or not an instruction should be executed depends on a control decision made by an earlier instruction

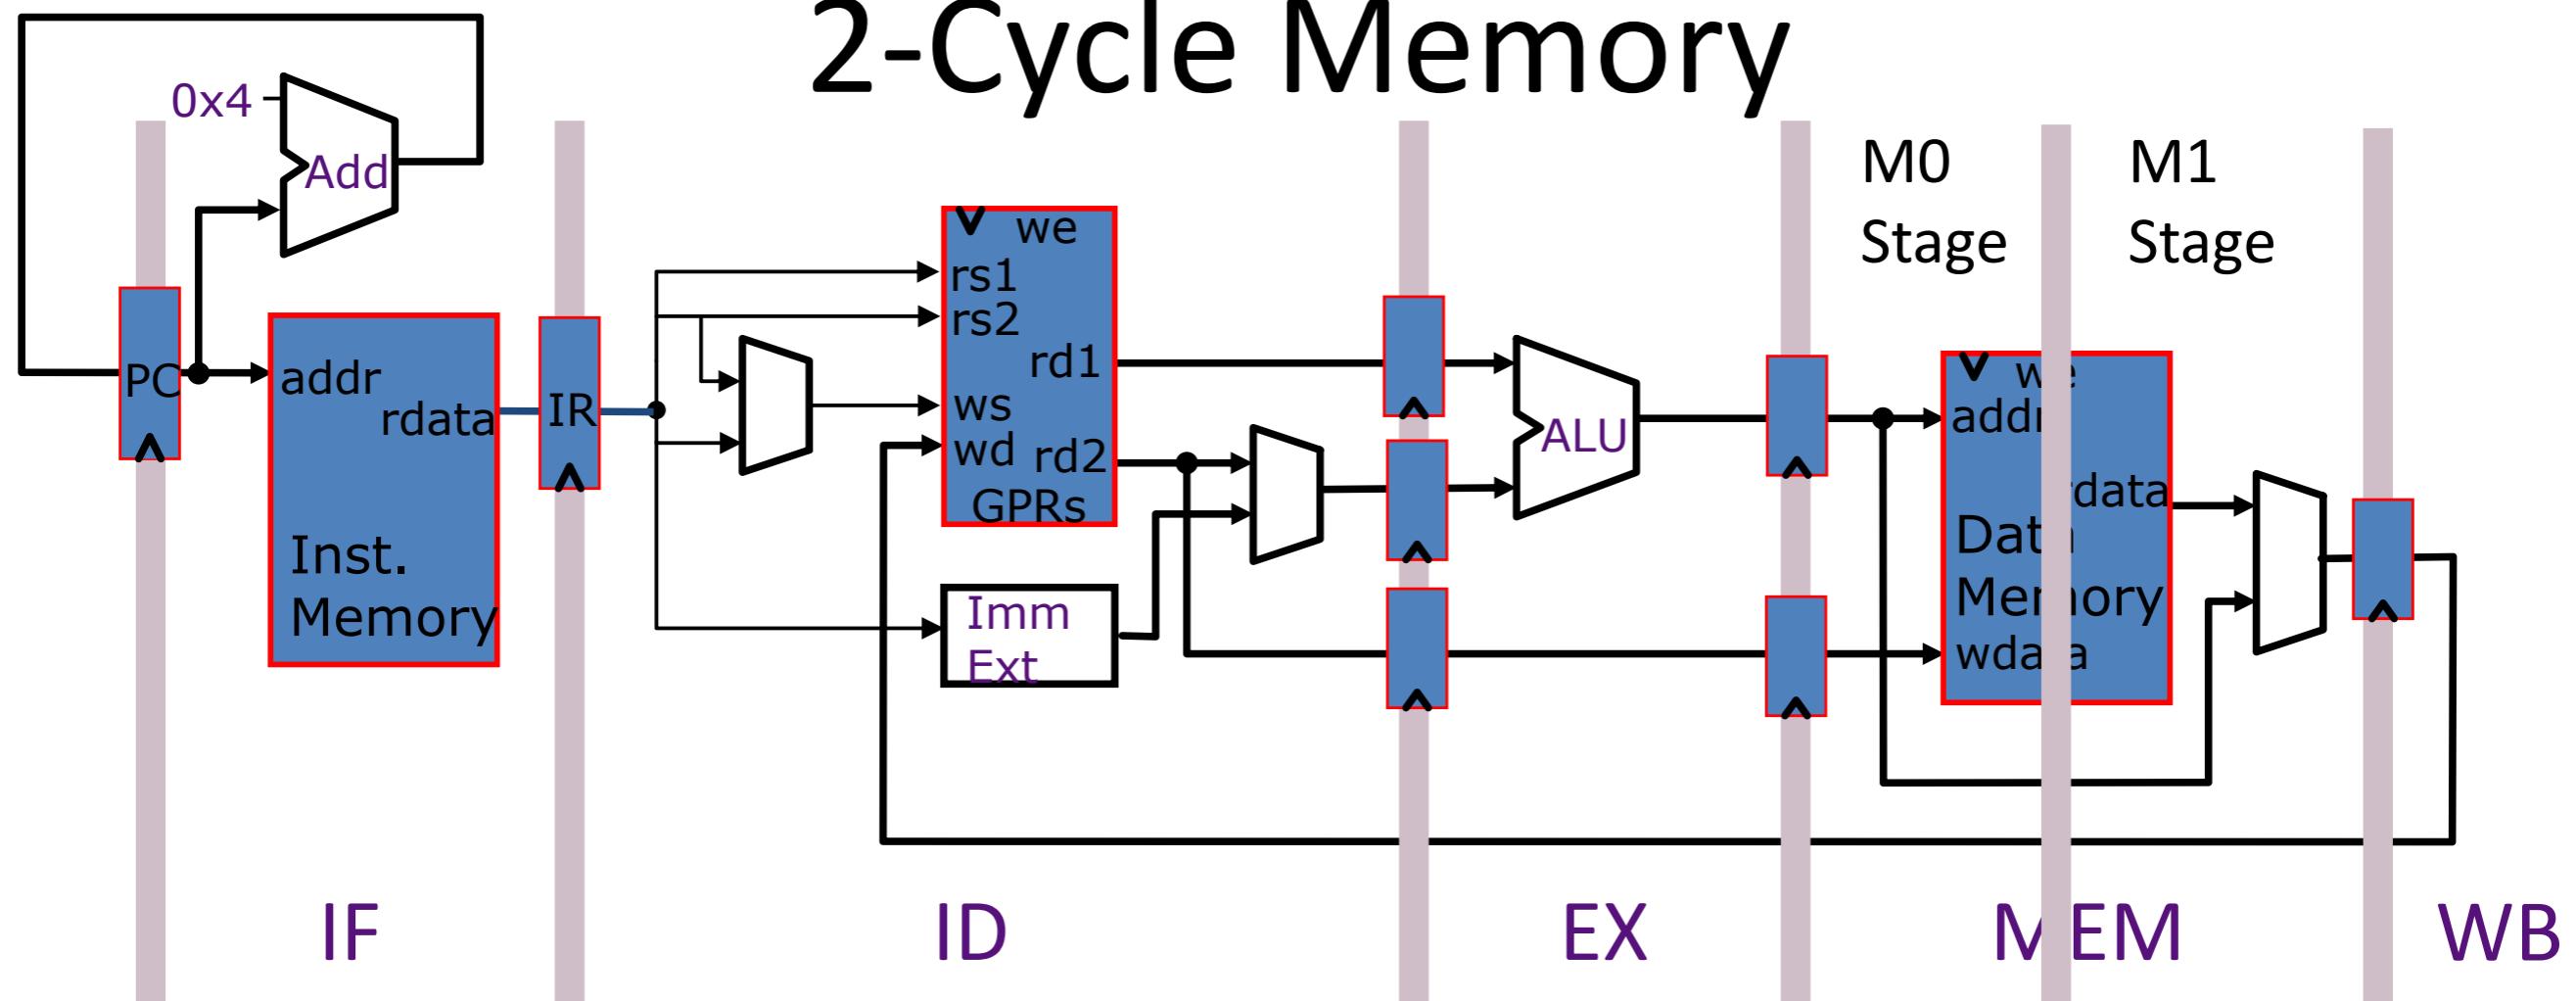
Example Structural Hazard: Unified Memory



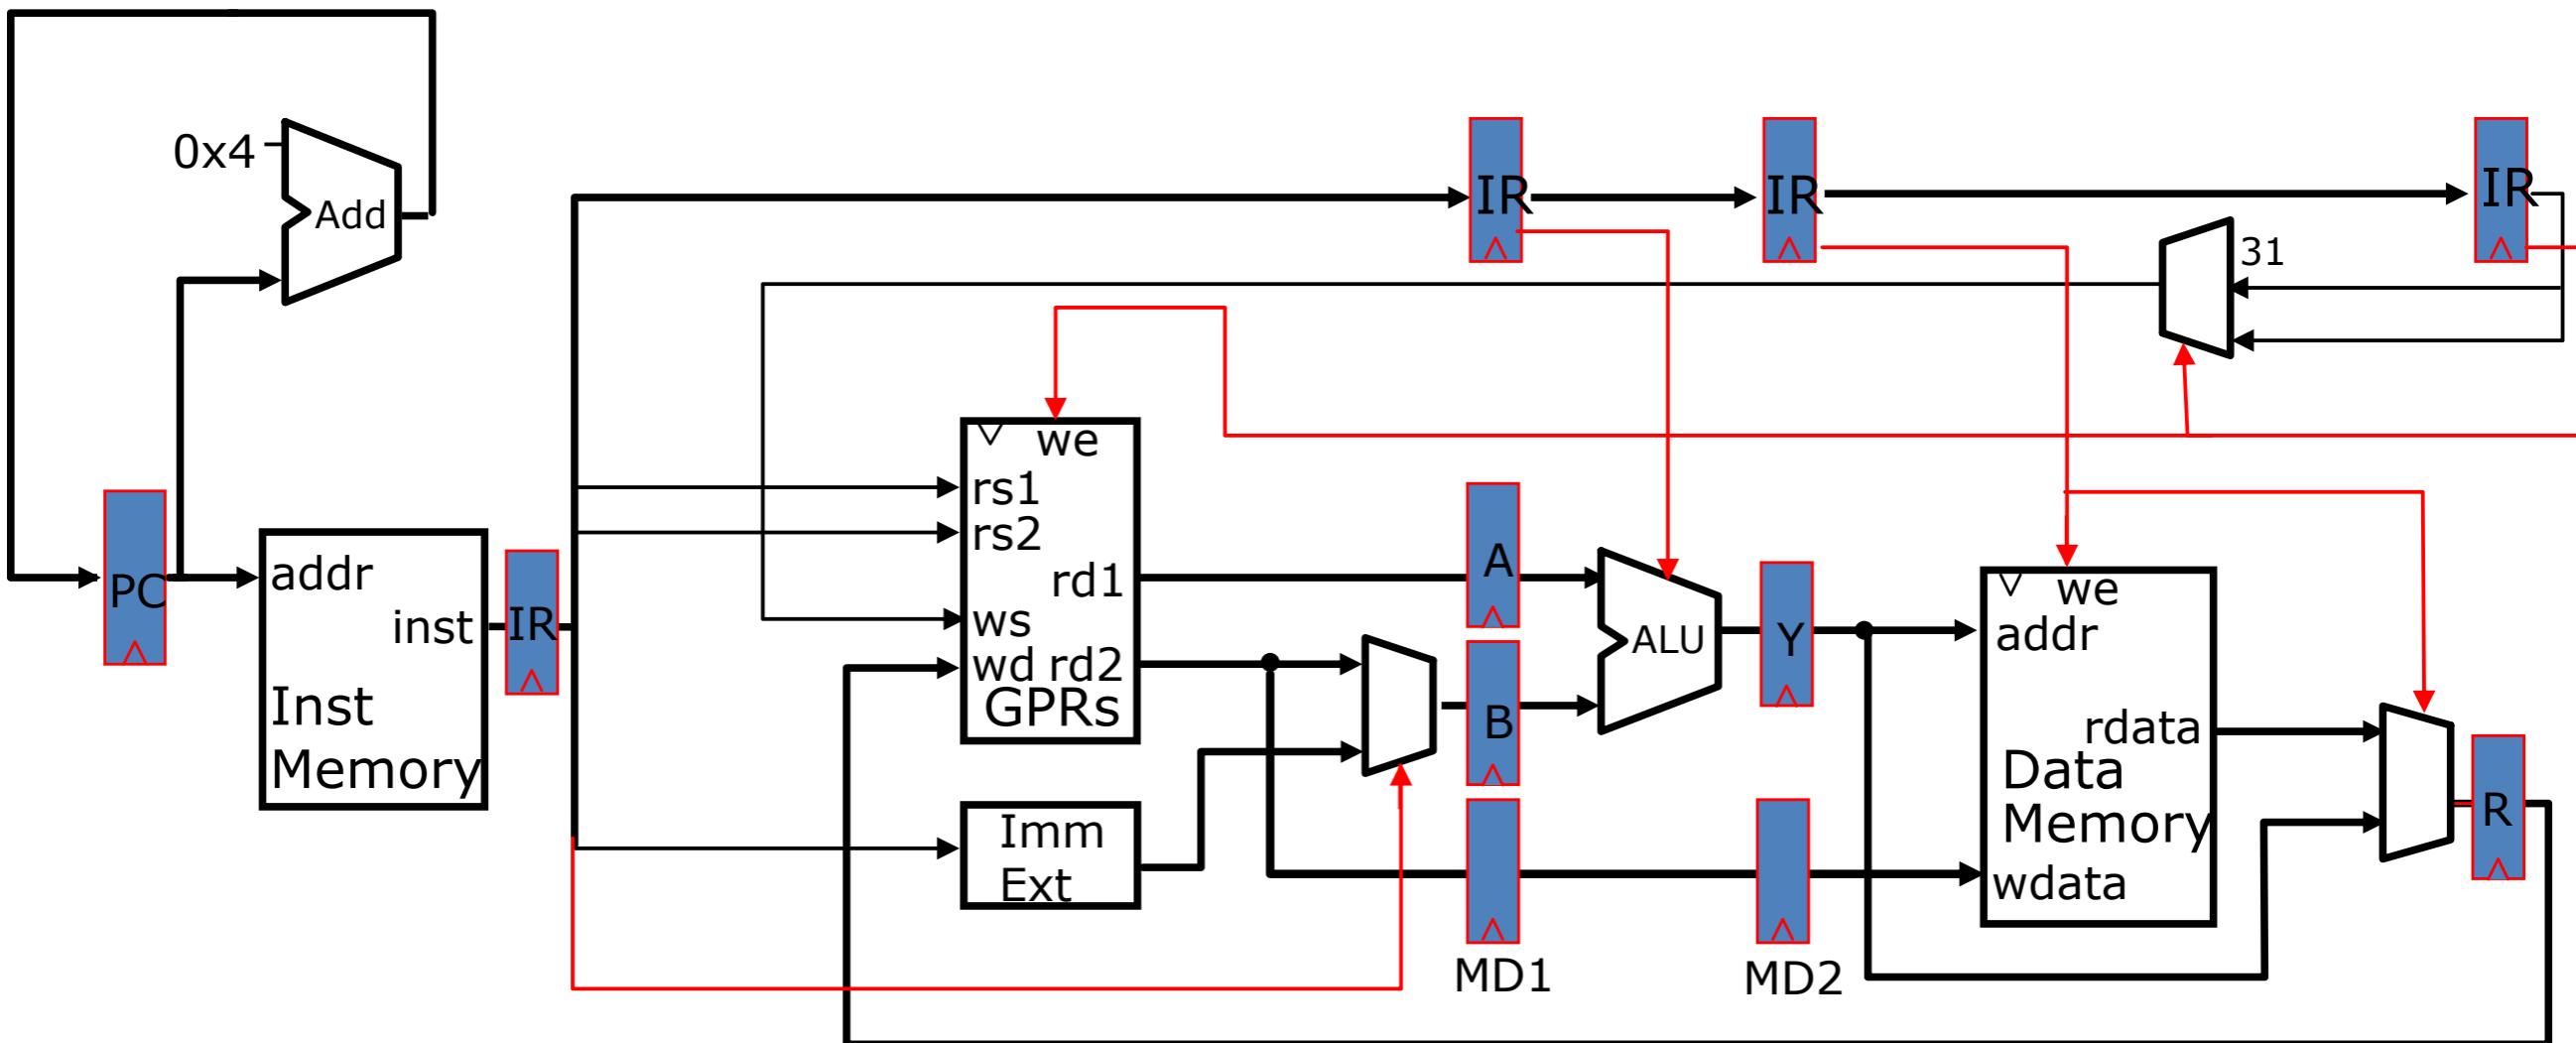
Example Structural Hazard: Unified Memory



Example Structural Hazard: 2-Cycle Memory



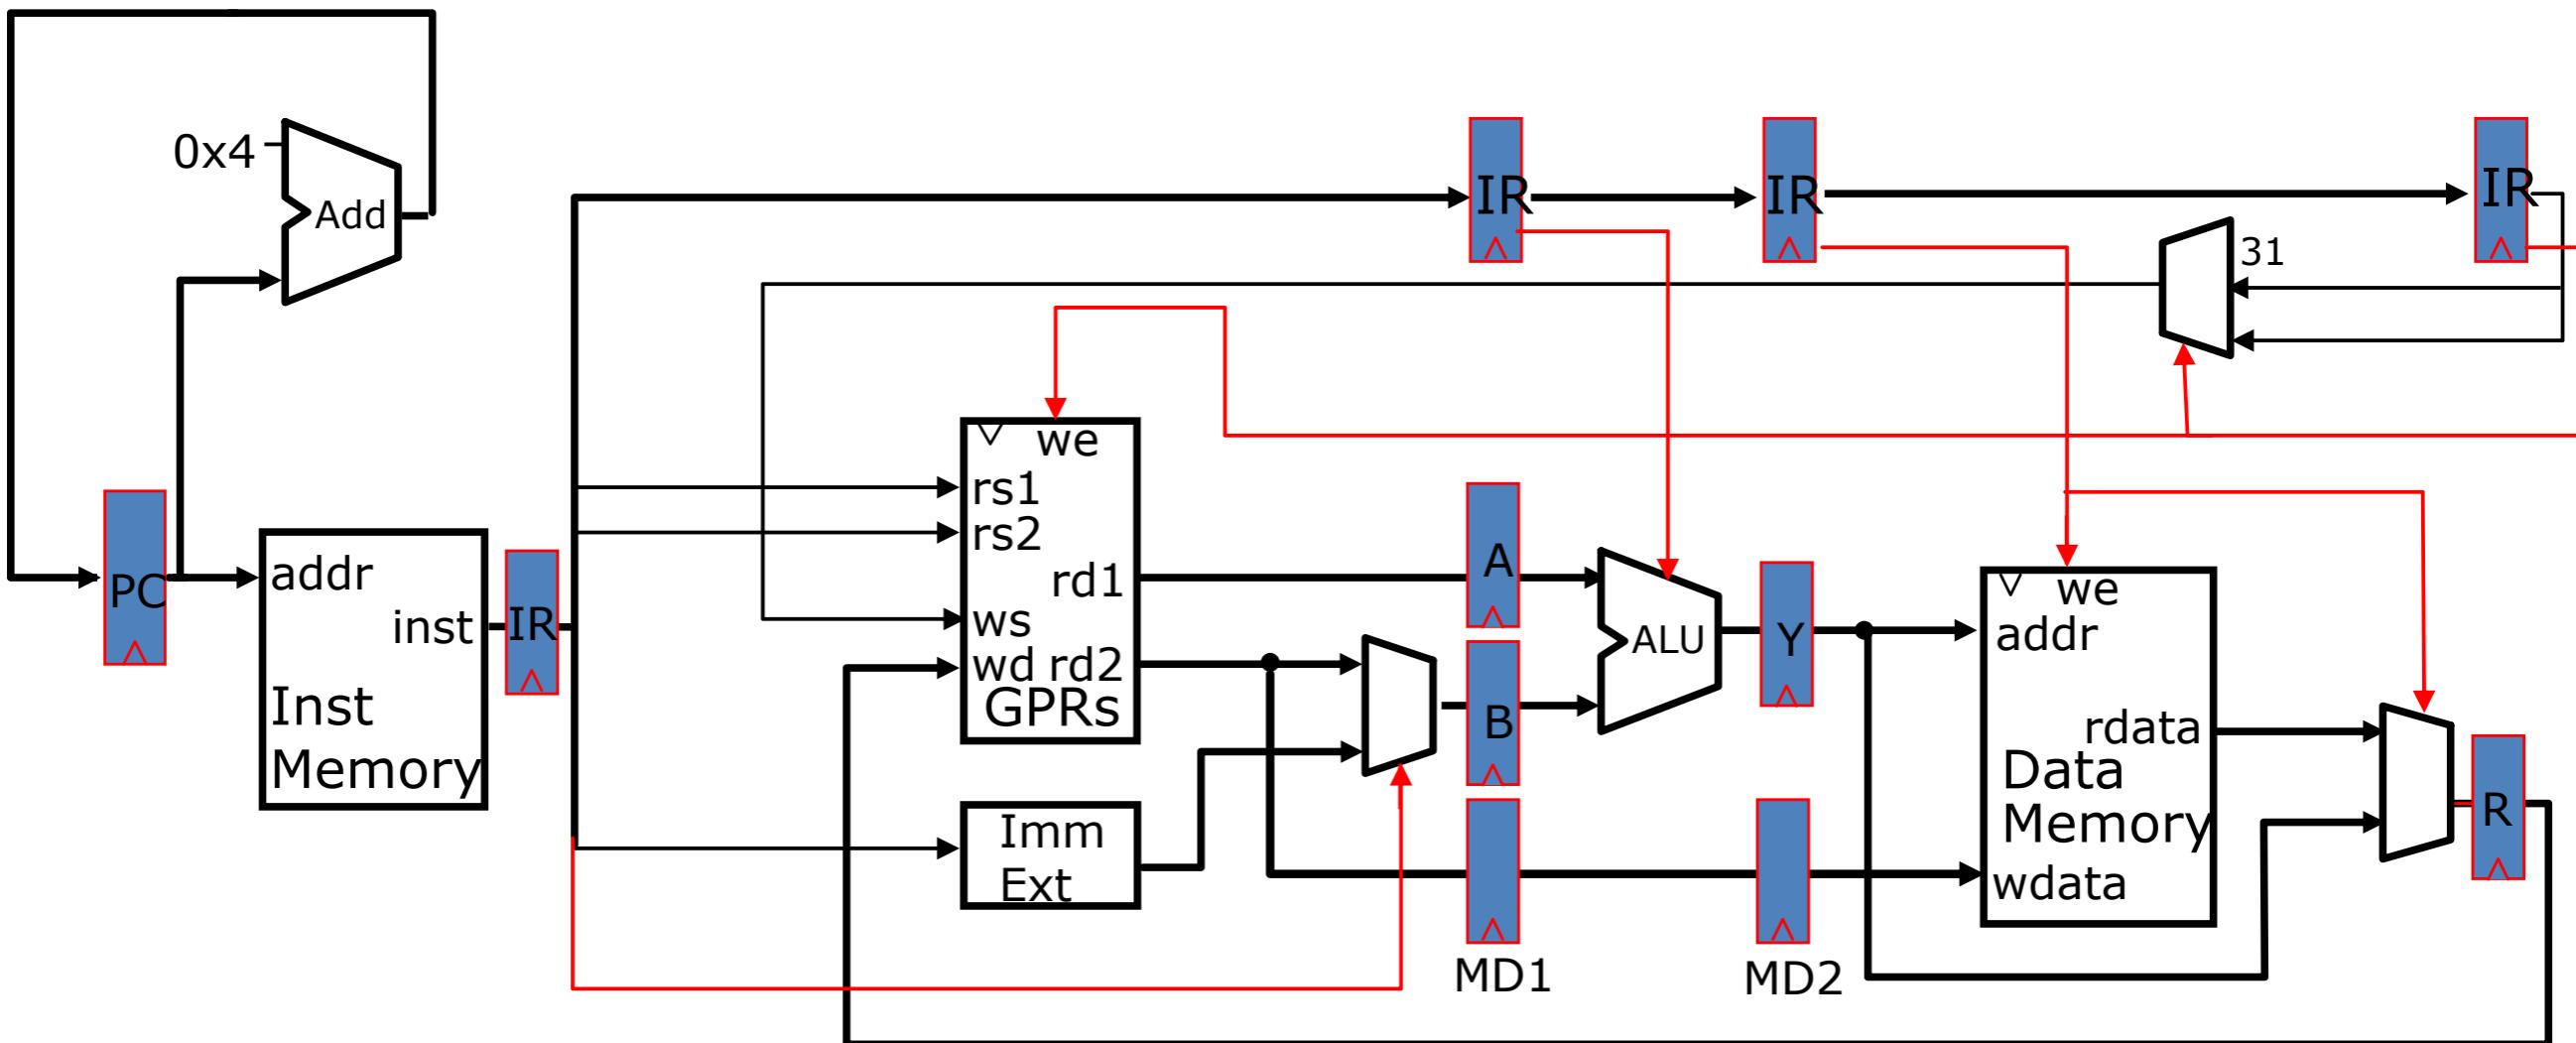
Example Data Hazard



...
r1 \leftarrow r0 + 10 (ADDI R1, R0, #10)
r4 \leftarrow r1 + 17 (ADDI R4, R1, #17)

...

Example Data Hazard



...
 $r_1 \leftarrow r_0 + 10$ (ADDI R1, R0, #10)
 $r_4 \leftarrow r_1 + 17$ (ADDI R4, R1, #17)

r1 is stale. Oops!

Data Dependence

Flow dependence

$$\begin{array}{rcl} r_3 & \leftarrow & r_1 \text{ op } r_2 \\ r_5 & \leftarrow & r_3 \text{ op } r_4 \end{array}$$

Read-after-Write
(RAW)

Anti dependence

$$\begin{array}{rcl} r_3 & \leftarrow & r_1 \text{ op } r_2 \\ r_1 & \leftarrow & r_4 \text{ op } r_5 \end{array}$$

Write-after-Read
(WAR)

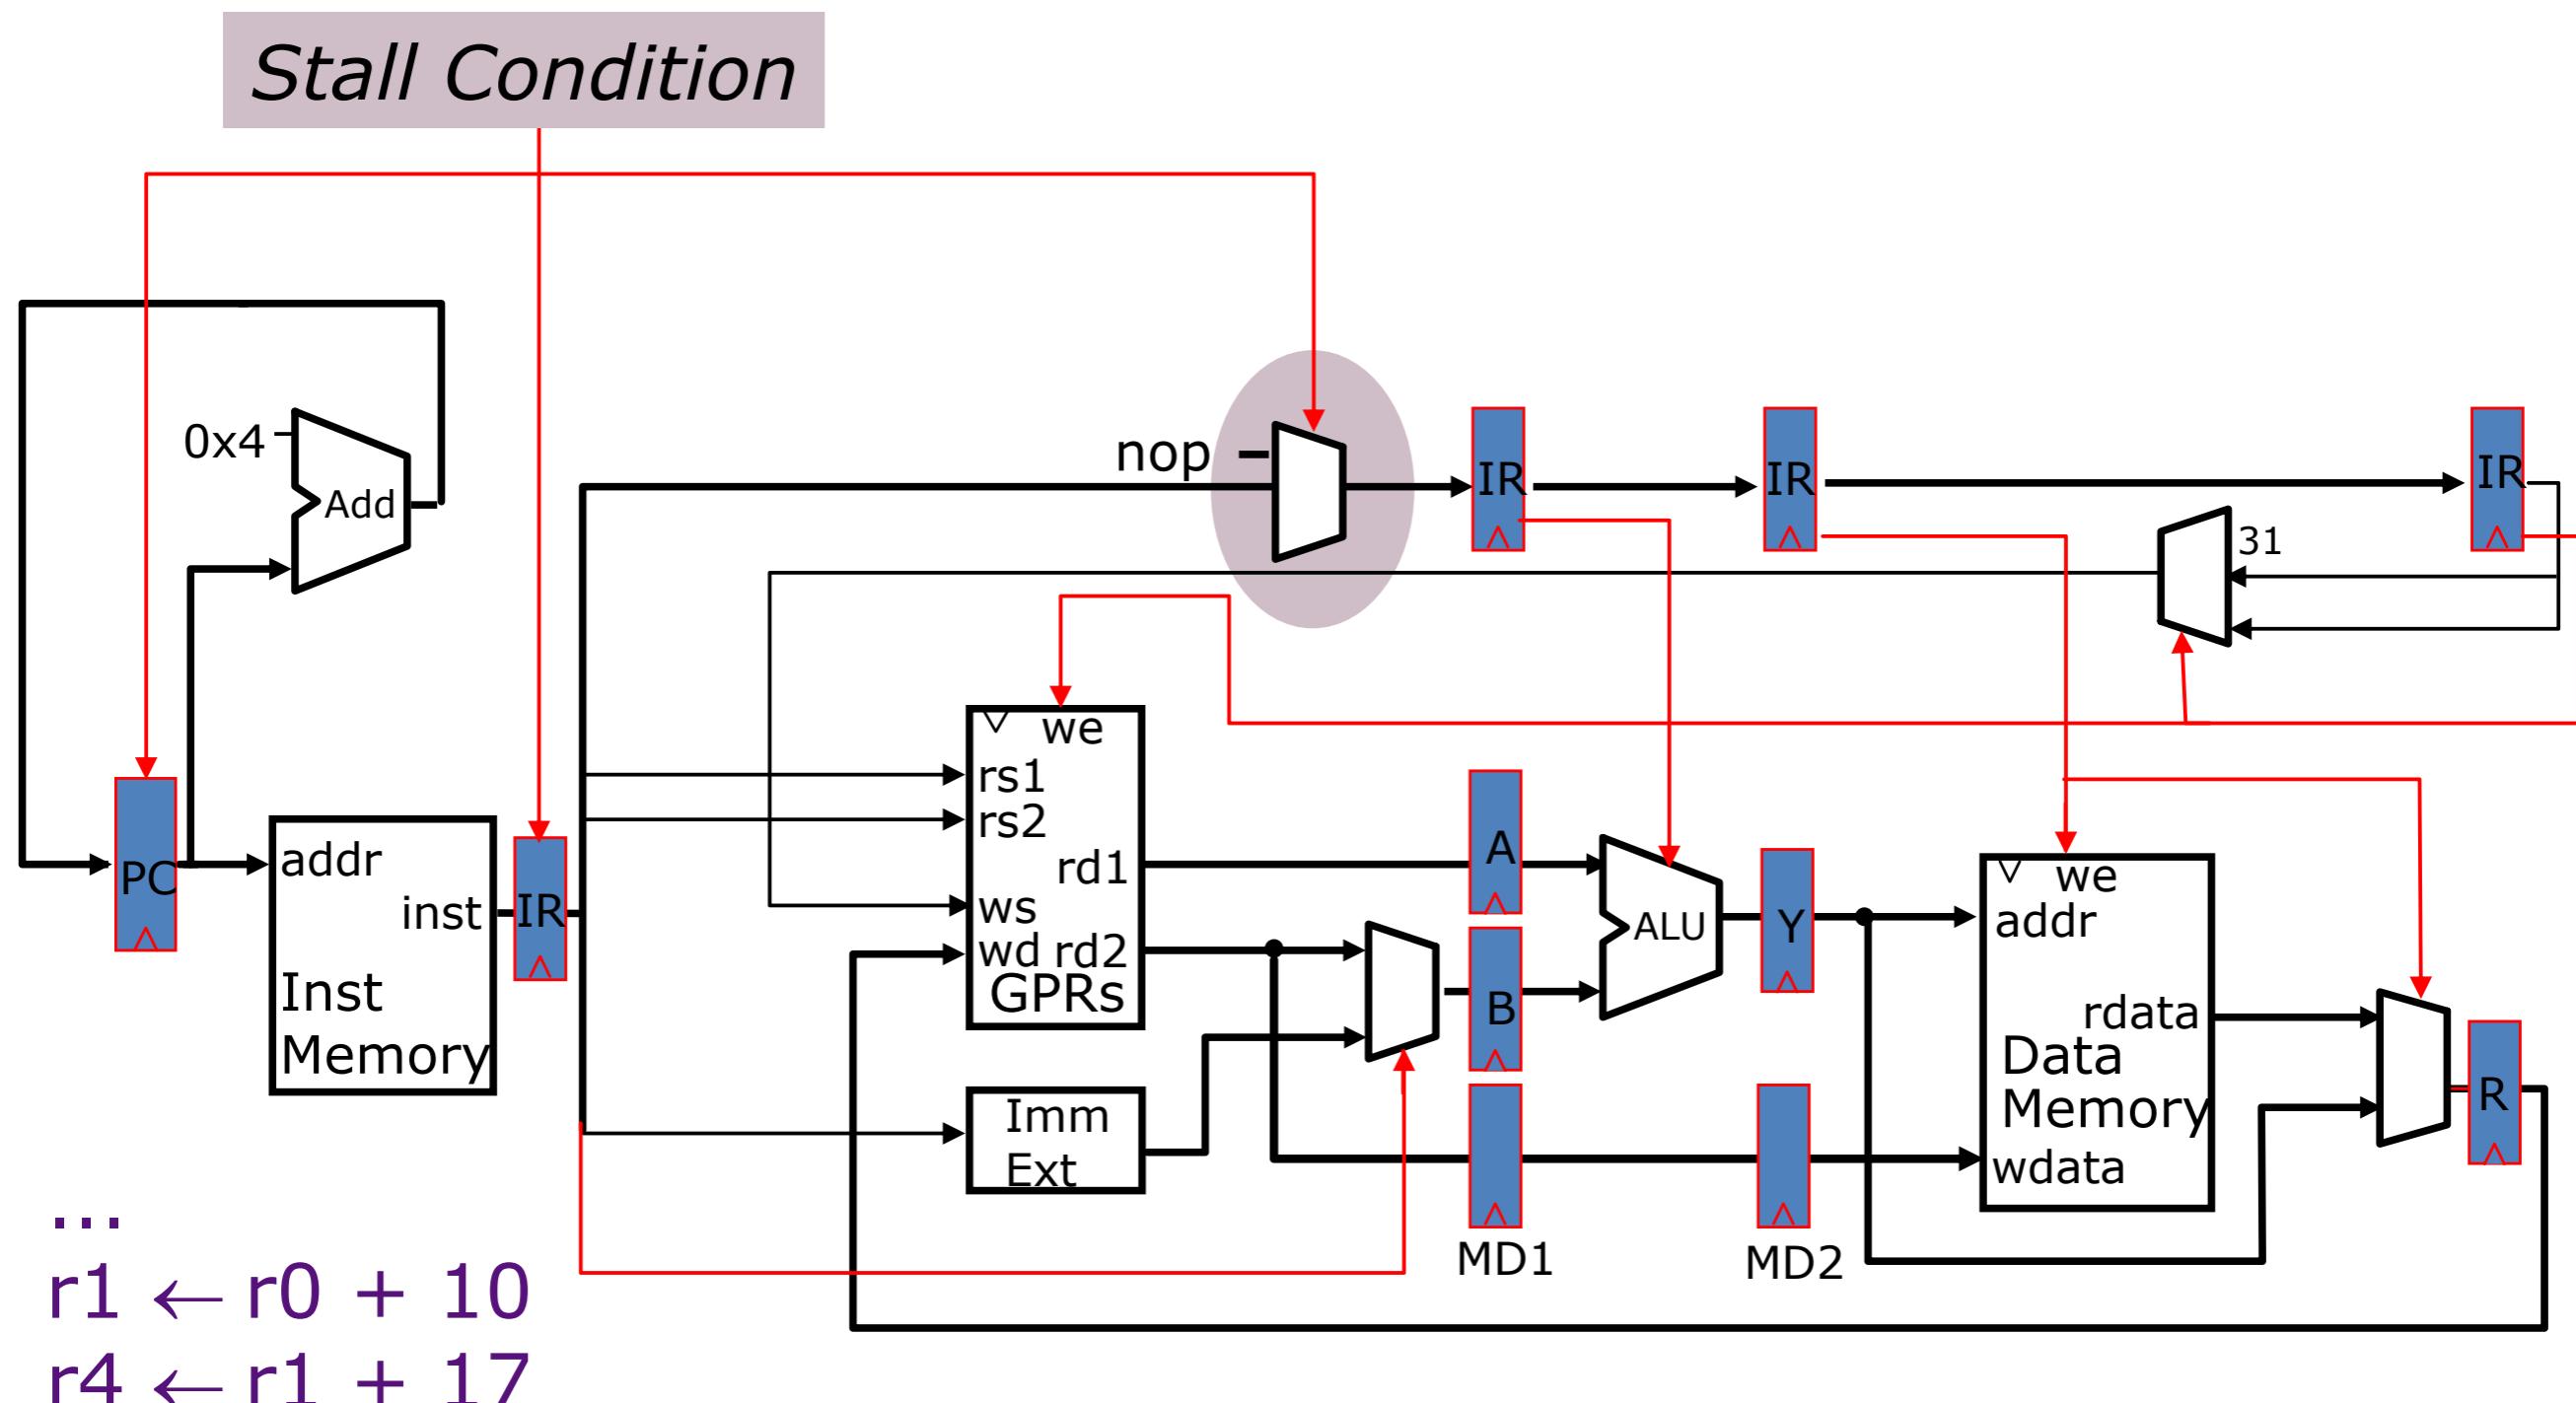
Output-dependence

$$\begin{array}{rcl} r_3 & \leftarrow & r_1 \text{ op } r_2 \\ r_5 & \leftarrow & r_3 \text{ op } r_4 \\ r_3 & \leftarrow & r_6 \text{ op } r_7 \end{array}$$



Write-after-Write
(WAW)

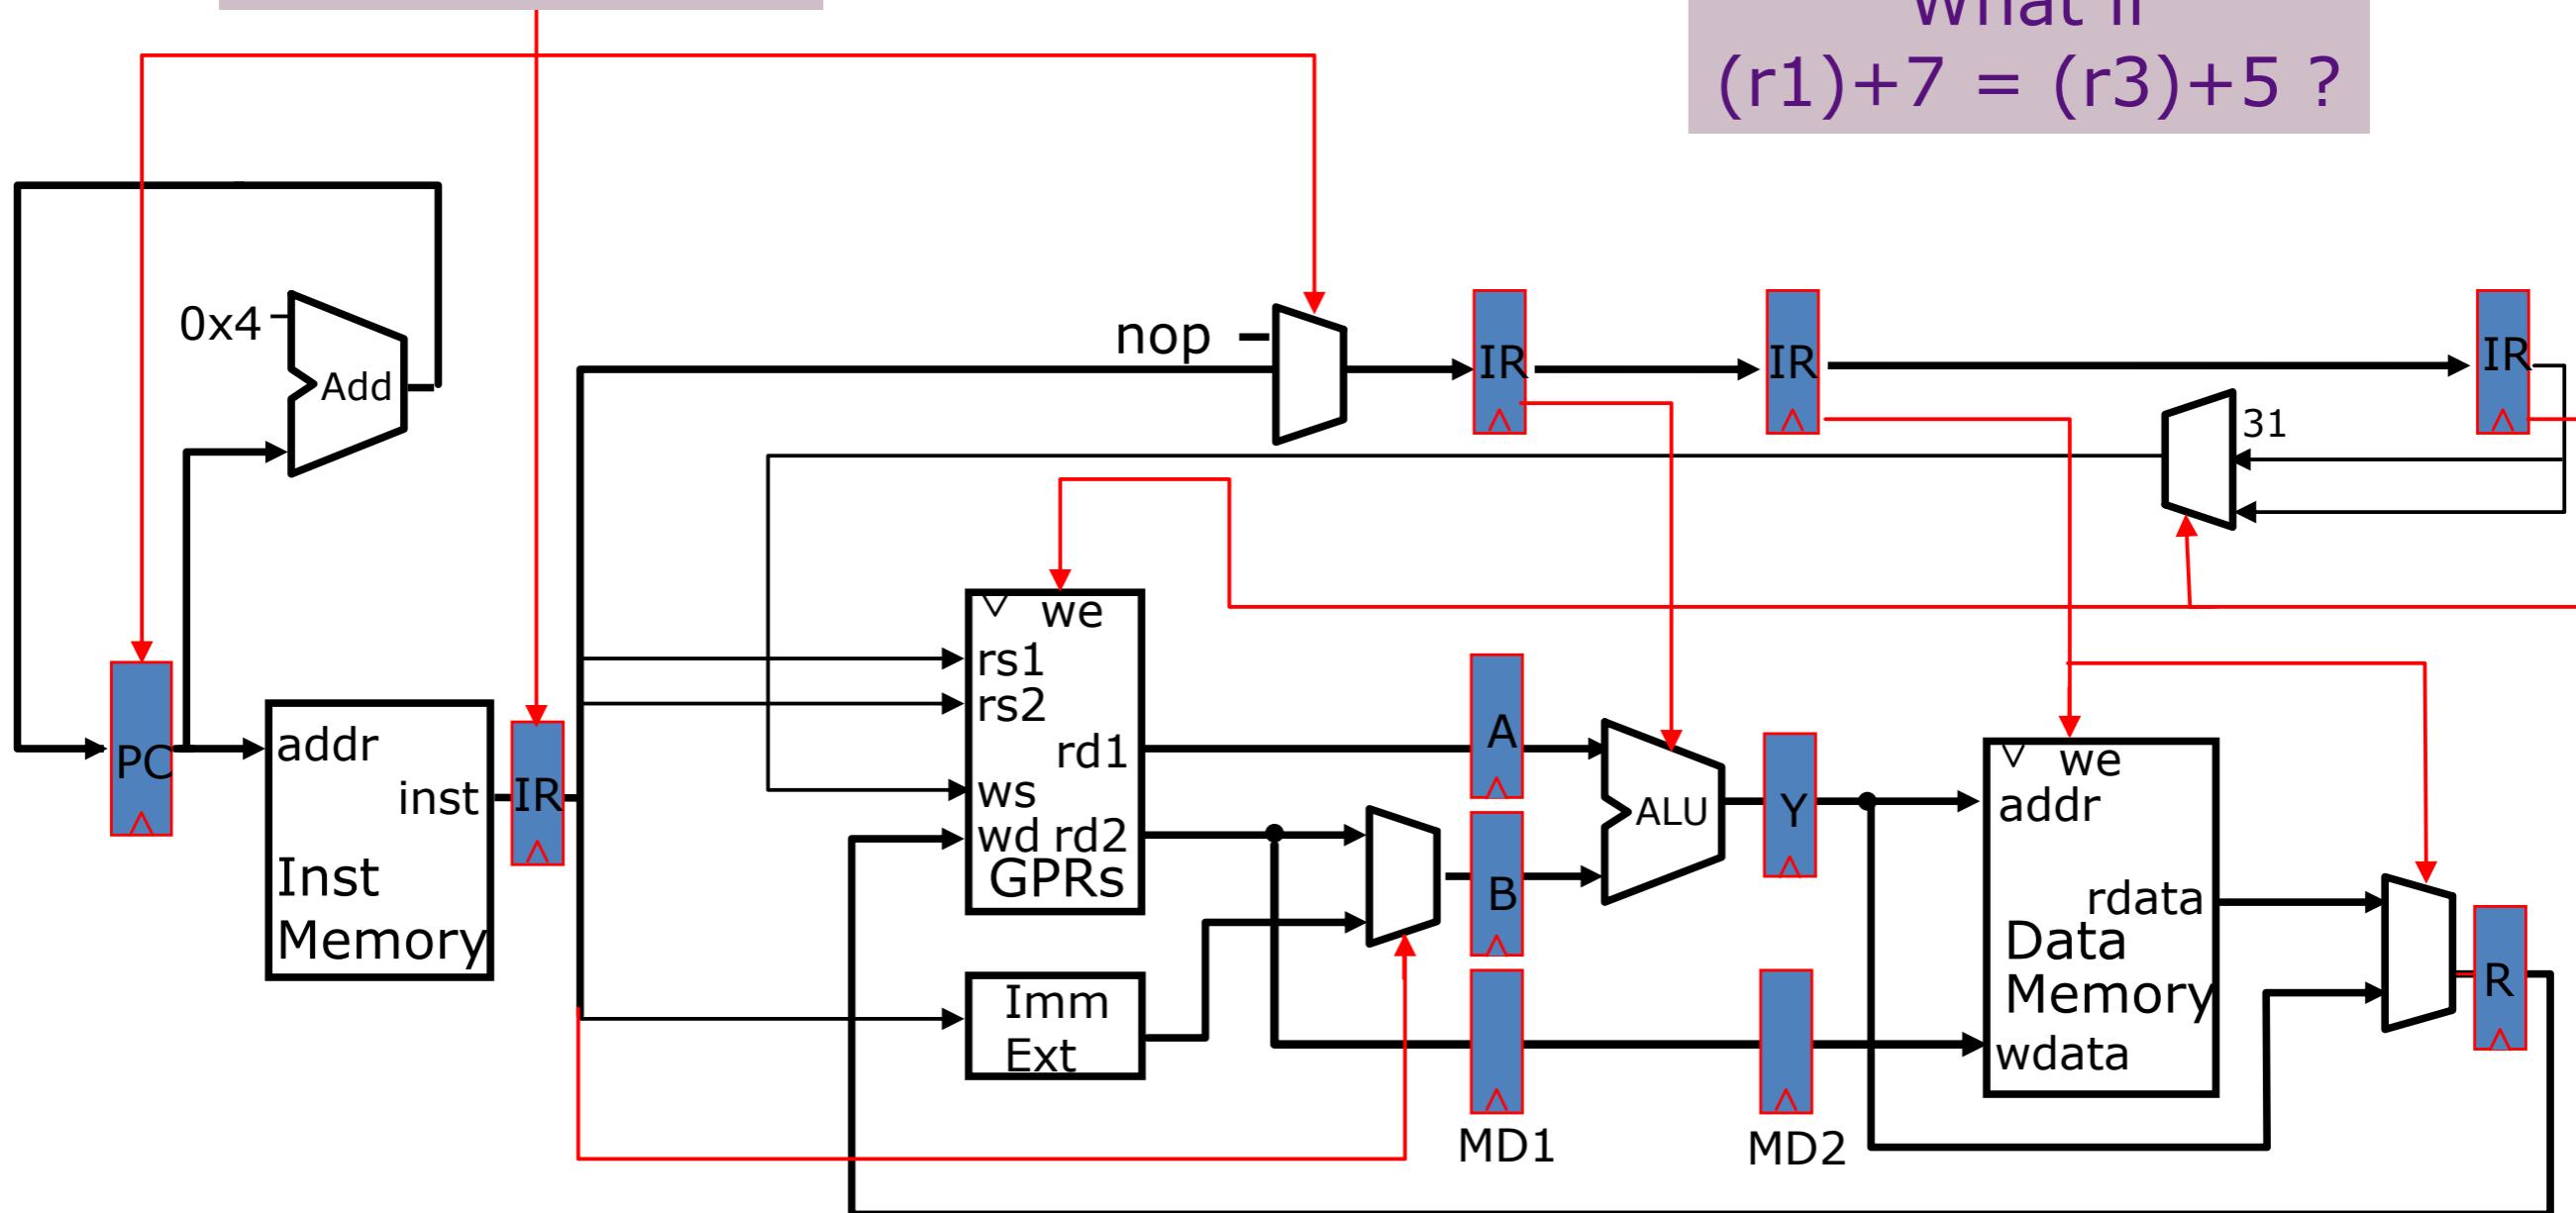
Resolving Data Hazards with Stalls



Hazards due to Loads & Stores

Stall Condition

What if
 $(r1)+7 = (r3)+5 ?$



...
 $M[(r1)+7] \leftarrow (r2)$
 $r4 \leftarrow M[(r3)+5]$
...

*Is there any possible data hazard
in this instruction sequence?*

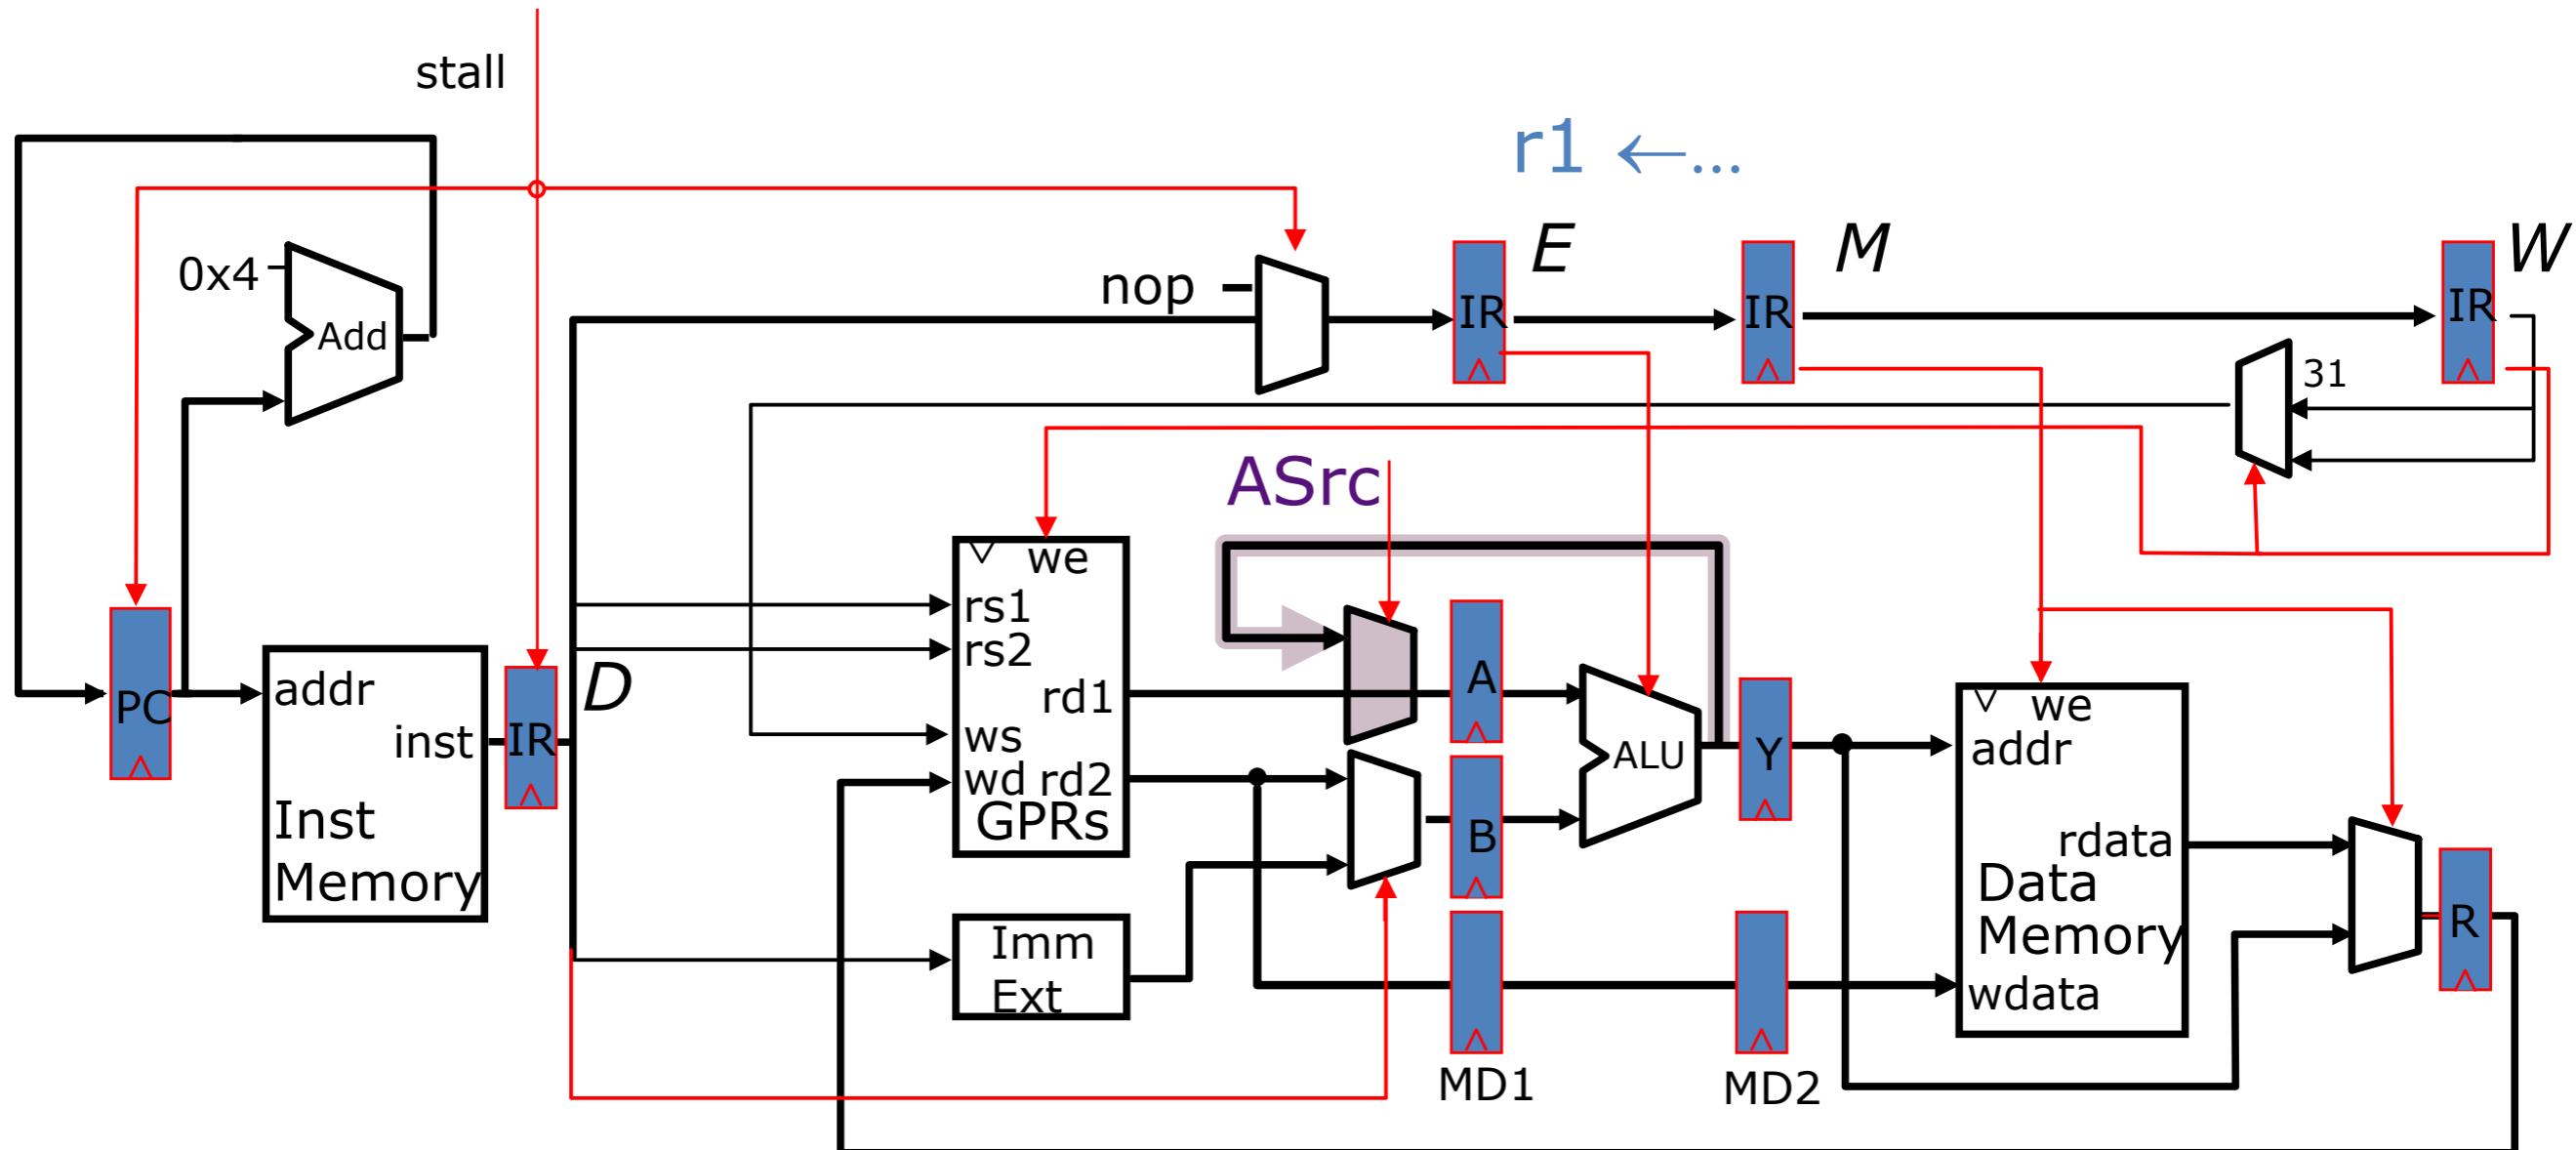
Data Hazards Due to Loads and Store

- Example instruction sequence
 - $\text{Mem}[\text{Regs}[r1] + 7] \leftarrow \text{Regs}[r2]$
 - $\text{Regs}[r4] \leftarrow \text{Mem}[\text{Regs}[r3] + 5]$

Data Hazards Due to Loads and Store

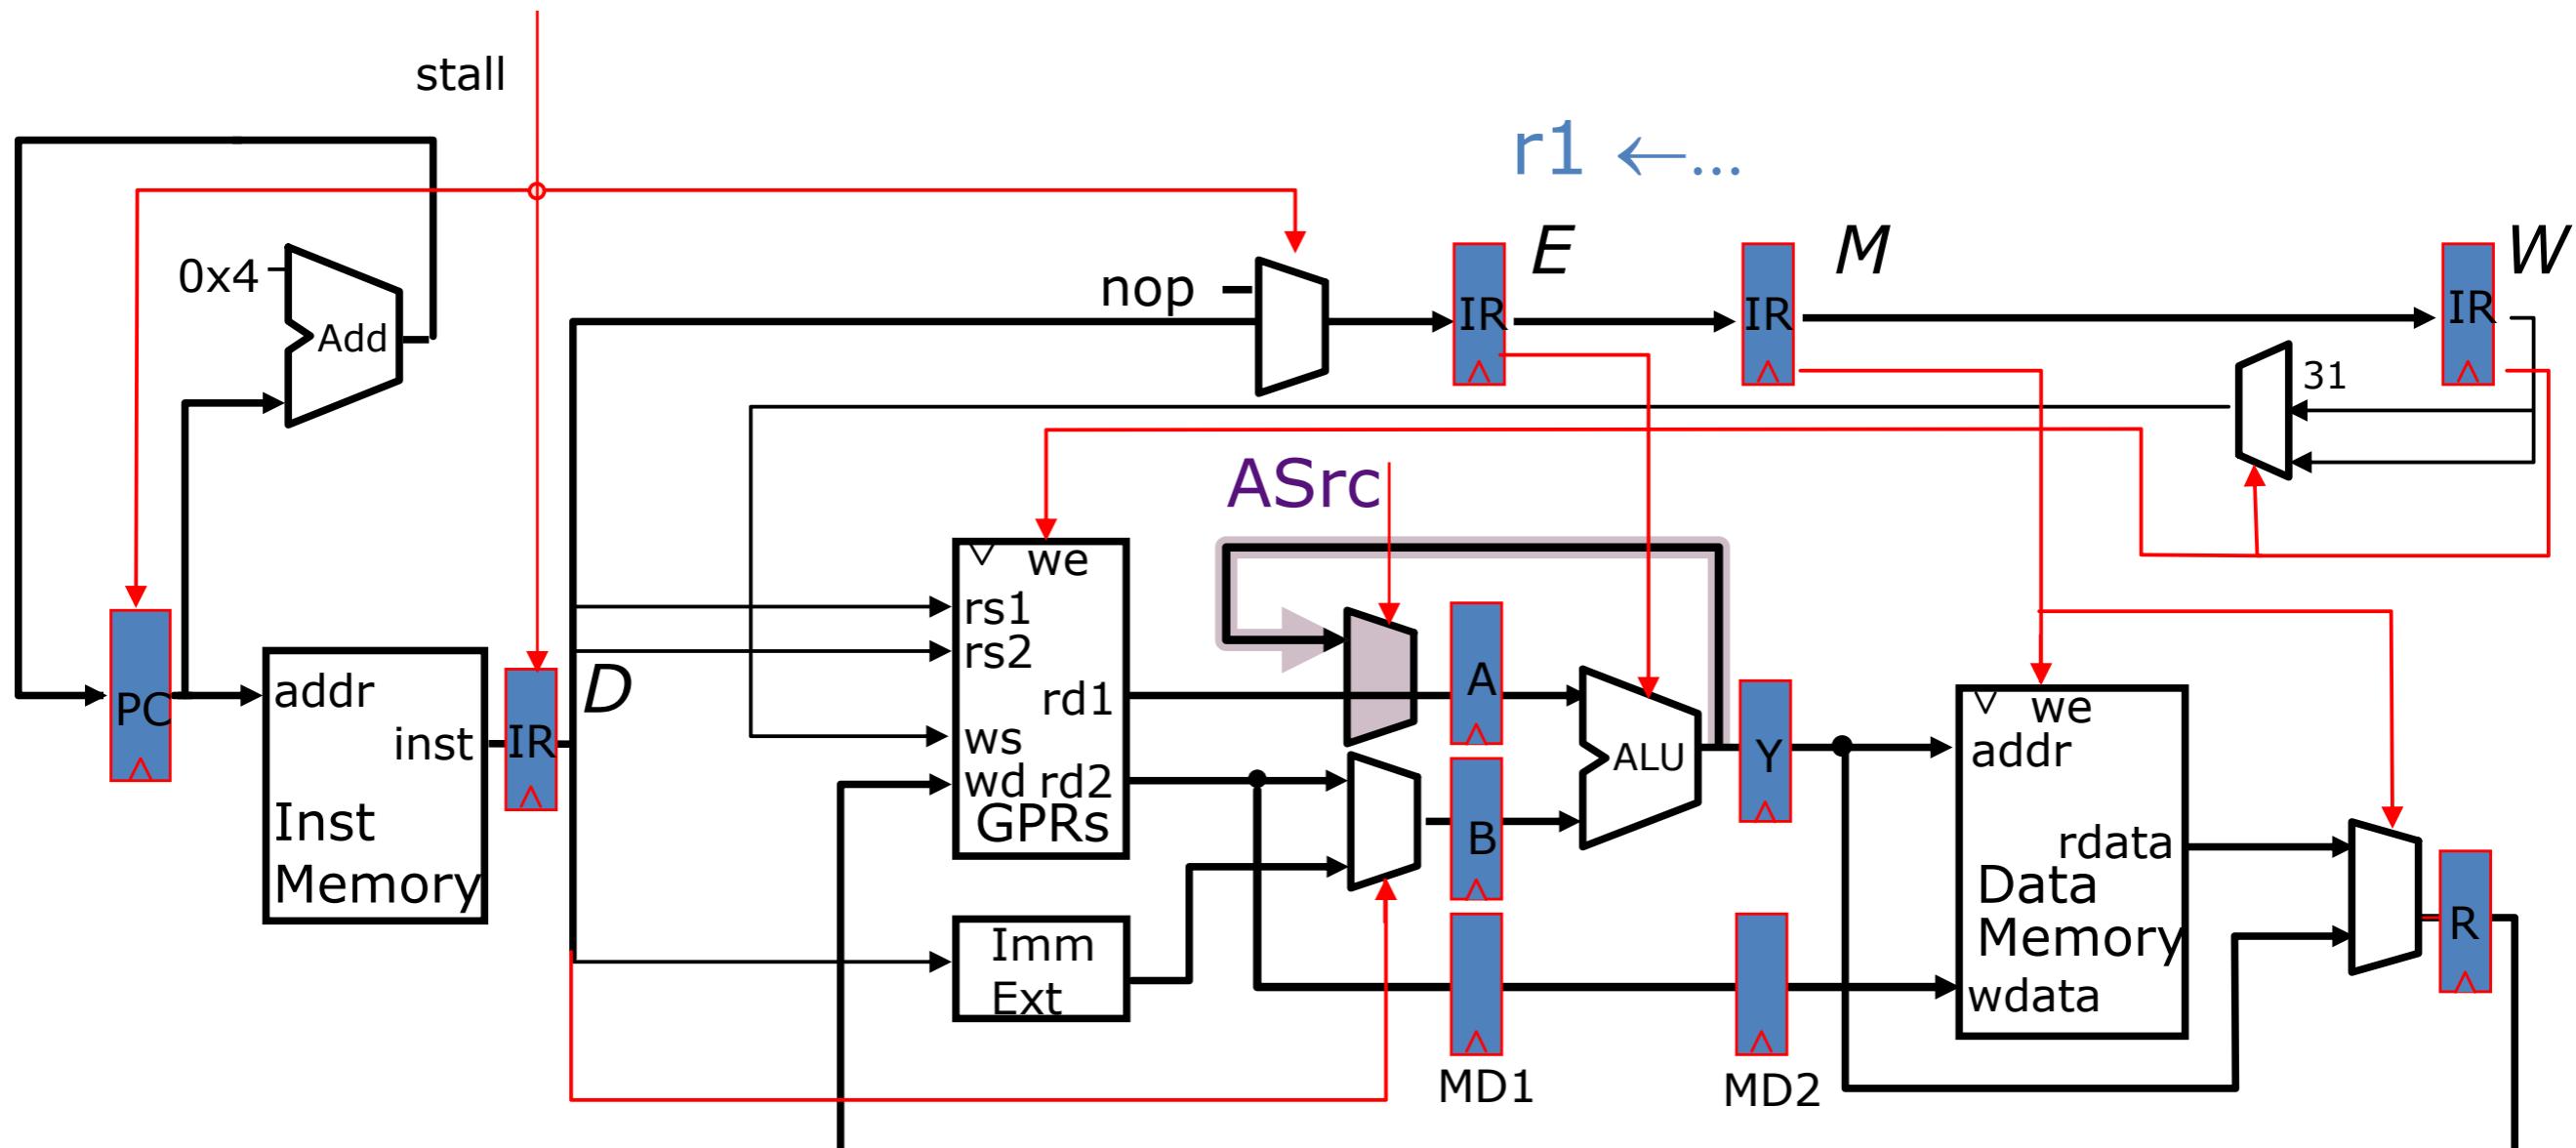
- Example instruction sequence
 - $\text{Mem}[\text{Regs}[r1] + 7] \leftarrow \text{Regs}[r2]$
 - $\text{Regs}[r4] \leftarrow \text{Mem}[\text{Regs}[r3] + 5]$
- What if $\text{Regs}[r1]+7 == \text{Regs}[r3]+5$?
 - Writing and reading to/from the same address
 - Hazard is avoided because our memory system completes writes in a single cycle
 - More realistic memory system will require more careful handling of data hazards due to loads and stores

Adding Bypassing to the Datapath



(I_1) $r1 \leftarrow r0 + 10$
 (I_2) $r4 \leftarrow r1 + 17$

Adding Bypassing to the Datapath



When does this bypass help?

(I_1)	$r1 \leftarrow r0 + 10$	$r1 \leftarrow \text{Mem}[r0 + 10]$	$\text{JAL } 500$
(I_2)	$r4 \leftarrow r1 + 17$	$r4 \leftarrow r1 + 17$	$r4 \leftarrow r31 + 1$

Bypassing in Action

- Consider this 8-stage pipeline



- For the following pairs of instructions, how many stalls will the second instruction experience (with and without bypassing)?

ADD R1+R2→R3

ADD R3+R4→R5

LD [R1]→R2

ADD R2+R3→R4

LD [R1]→R2

SD [R2]←R3

LD [R1]→R2

SD [R3]←R2

Bypassing in Action

- Consider this 8-stage pipeline



- For the following pairs of instructions, how many stalls will the second instruction experience (with and without bypassing)?

ADD R1+R2→R3

without: 5 with: 1

LD [R1]→R2

without: 5 with: 3

LD [R1]→R2

without: 5 with: 3

SD [R2]←R3

without: 5 with: 1

LD [R1]→R2

without: 5 with: 1

SD [R3]←R2

Control Hazards

- What do we need to calculate next PC?

Control Hazards

- What do we need to calculate next PC?
 - For Jumps
 - Opcode, offset and PC
 - For Jump Register
 - Opcode and Register value
 - For Conditional Branches
 - Opcode, PC, Register (for condition), and offset
 - For all other instructions
 - Opcode and PC

Opcode Decoding Bubble

(assuming no branch delay slots for now)

	time									
	t0	t1	t2	t3	t4	t5	t6	t7	...	
(I ₁) r1 ← (r0) + 10	IF ₁	ID ₁	EX ₁	MA ₁	WB ₁					
(I ₂) r3 ← (r2) + 17		IF ₂	IF ₂	ID ₂	EX ₂	MA ₂	WB ₂			
(I ₃)			IF ₃	IF ₃	ID ₃	EX ₃	MA ₃	WB ₃		
(I ₄)				IF ₄	IF ₄	ID ₄	EX ₄	MA ₄	WB ₄	

Opcode Decoding Bubble

(assuming no branch delay slots for now)

	time									
	t0	t1	t2	t3	t4	t5	t6	t7	...	
(I ₁) r1 ← (r0) + 10	IF ₁	ID ₁	EX ₁	MA ₁	WB ₁					
(I ₂) r3 ← (r2) + 17		IF ₂	IF ₂	ID ₂	EX ₂	MA ₂	WB ₂			
(I ₃)			IF ₃	IF ₃	ID ₃	EX ₃	MA ₃	WB ₃		
(I ₄)				IF ₄	IF ₄	ID ₄	EX ₄	MA ₄	WB ₄	

Resource Usage	time										
	t0	t1	t2	t3	t4	t5	t6	t7	...		
	IF	I ₁	nop	I ₂	nop	I ₃	nop	I ₄			
	ID		I ₁	nop	I ₂	nop	I ₃	nop	I ₄		
	EX			I ₁	nop	I ₂	nop	I ₃	nop	I ₄	
	MA				I ₁	nop	I ₂	nop	I ₃	nop	I ₄
	WB					I ₁	nop	I ₂	nop	I ₃	nop

nop ⇒ *pipeline bubble*

Opcode Decoding Bubble

(assuming no branch delay slots for now)

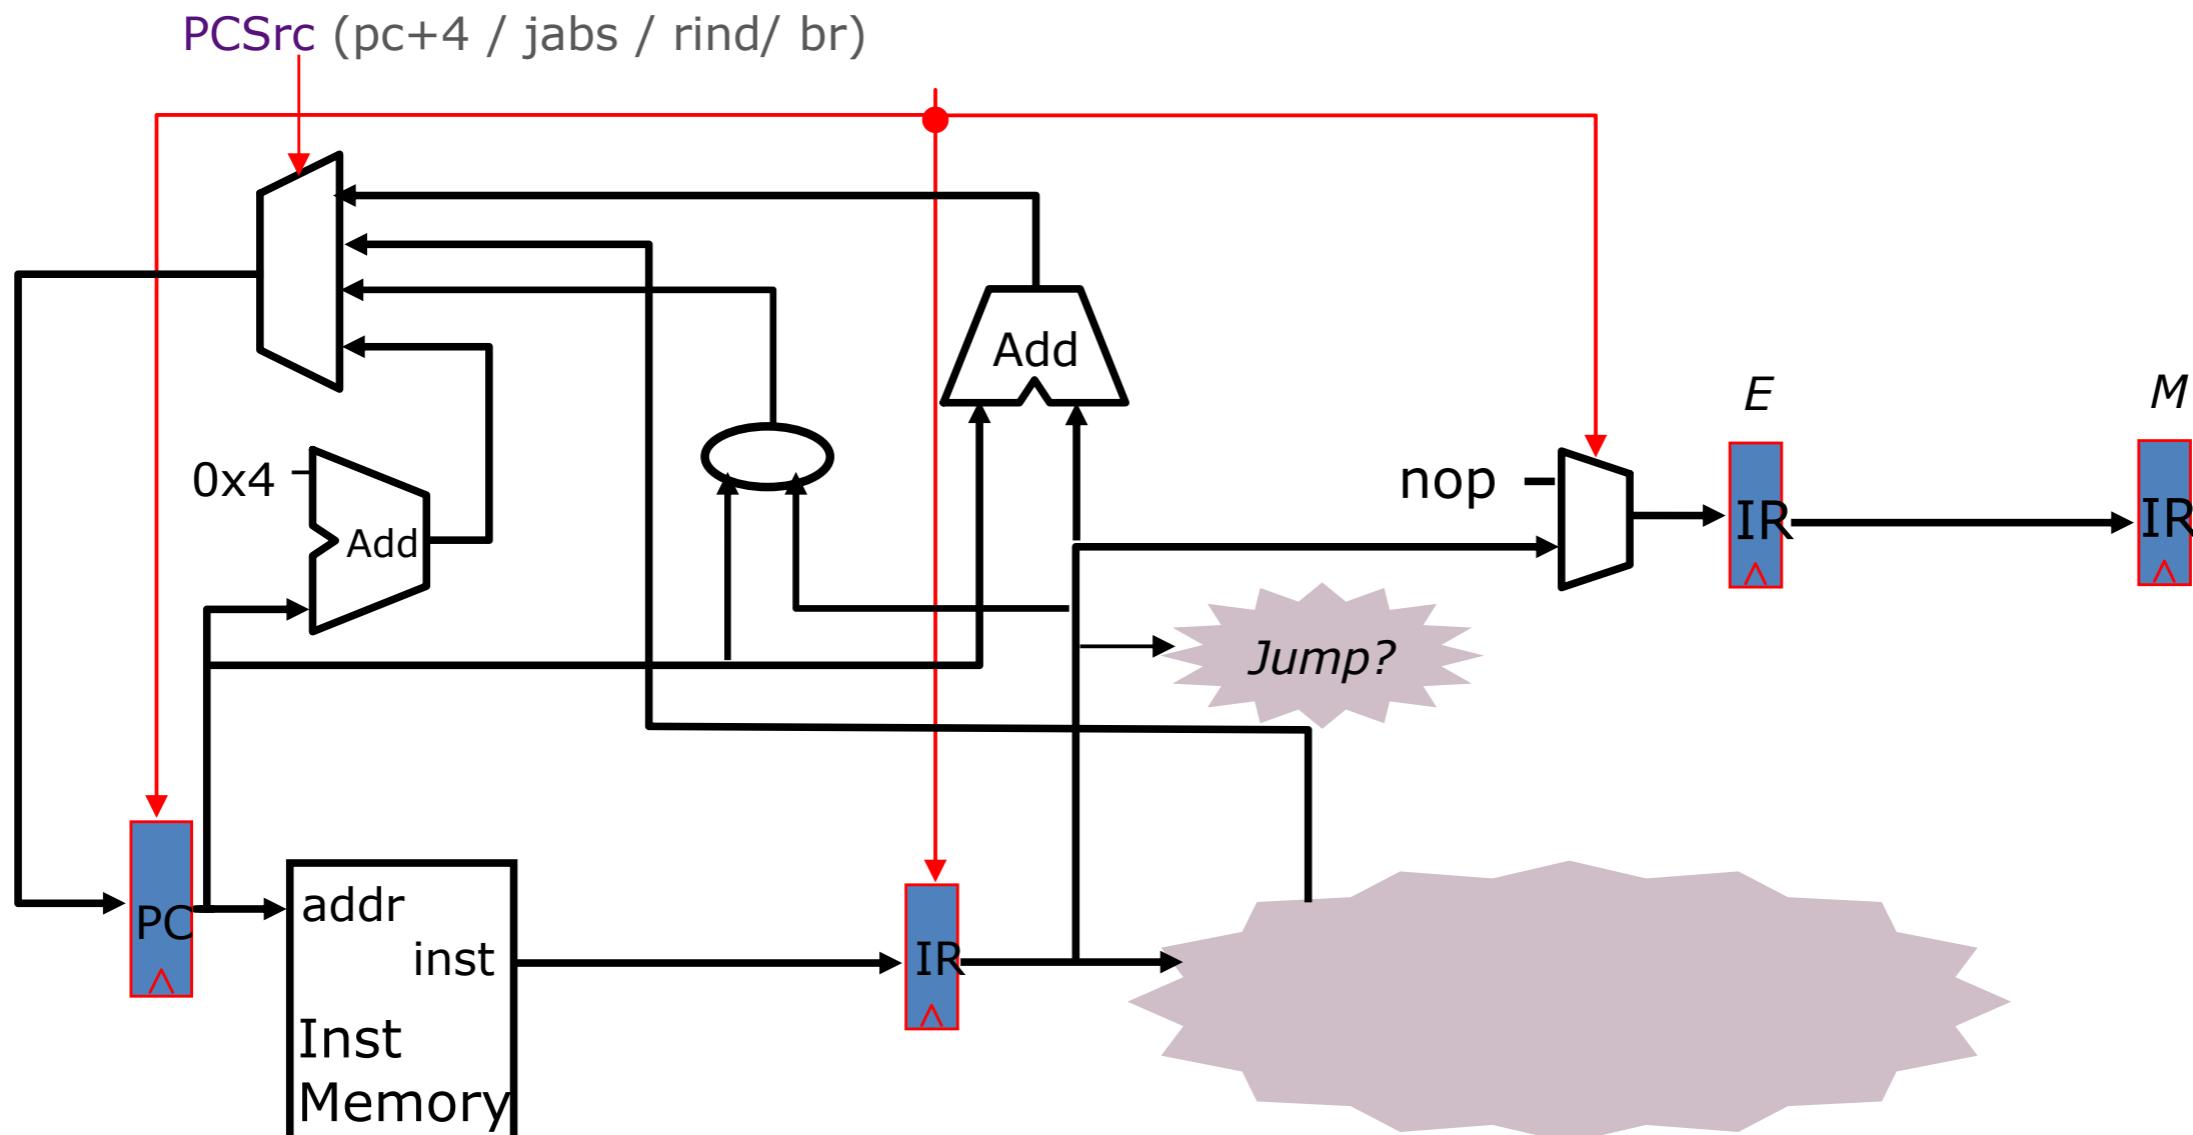
	time									
	t0	t1	t2	t3	t4	t5	t6	t7	...	
(I ₁) r1 ← (r0) + 10	IF ₁	ID ₁	EX ₁	MA ₁	WB ₁					
(I ₂) r3 ← (r2) + 17		IF ₂	IF ₂	ID ₂	EX ₂	MA ₂	WB ₂			
(I ₃)			IF ₃	IF ₃	ID ₃	EX ₃	MA ₃	WB ₃		
(I ₄)				IF ₄	IF ₄	ID ₄	EX ₄	MA ₄	WB ₄	

	time									
	t0	t1	t2	t3	t4	t5	t6	t7	...	
Resource Usage	IF	I ₁	nop	I ₂	nop	I ₃	nop	I ₄		
	ID		I ₁	nop	I ₂	nop	I ₃	nop	I ₄	
	EX			I ₁	nop	I ₂	nop	I ₃	nop	I ₄
	MA				I ₁	nop	I ₂	nop	I ₃	nop
	WB					I ₁	nop	I ₂	nop	I ₃

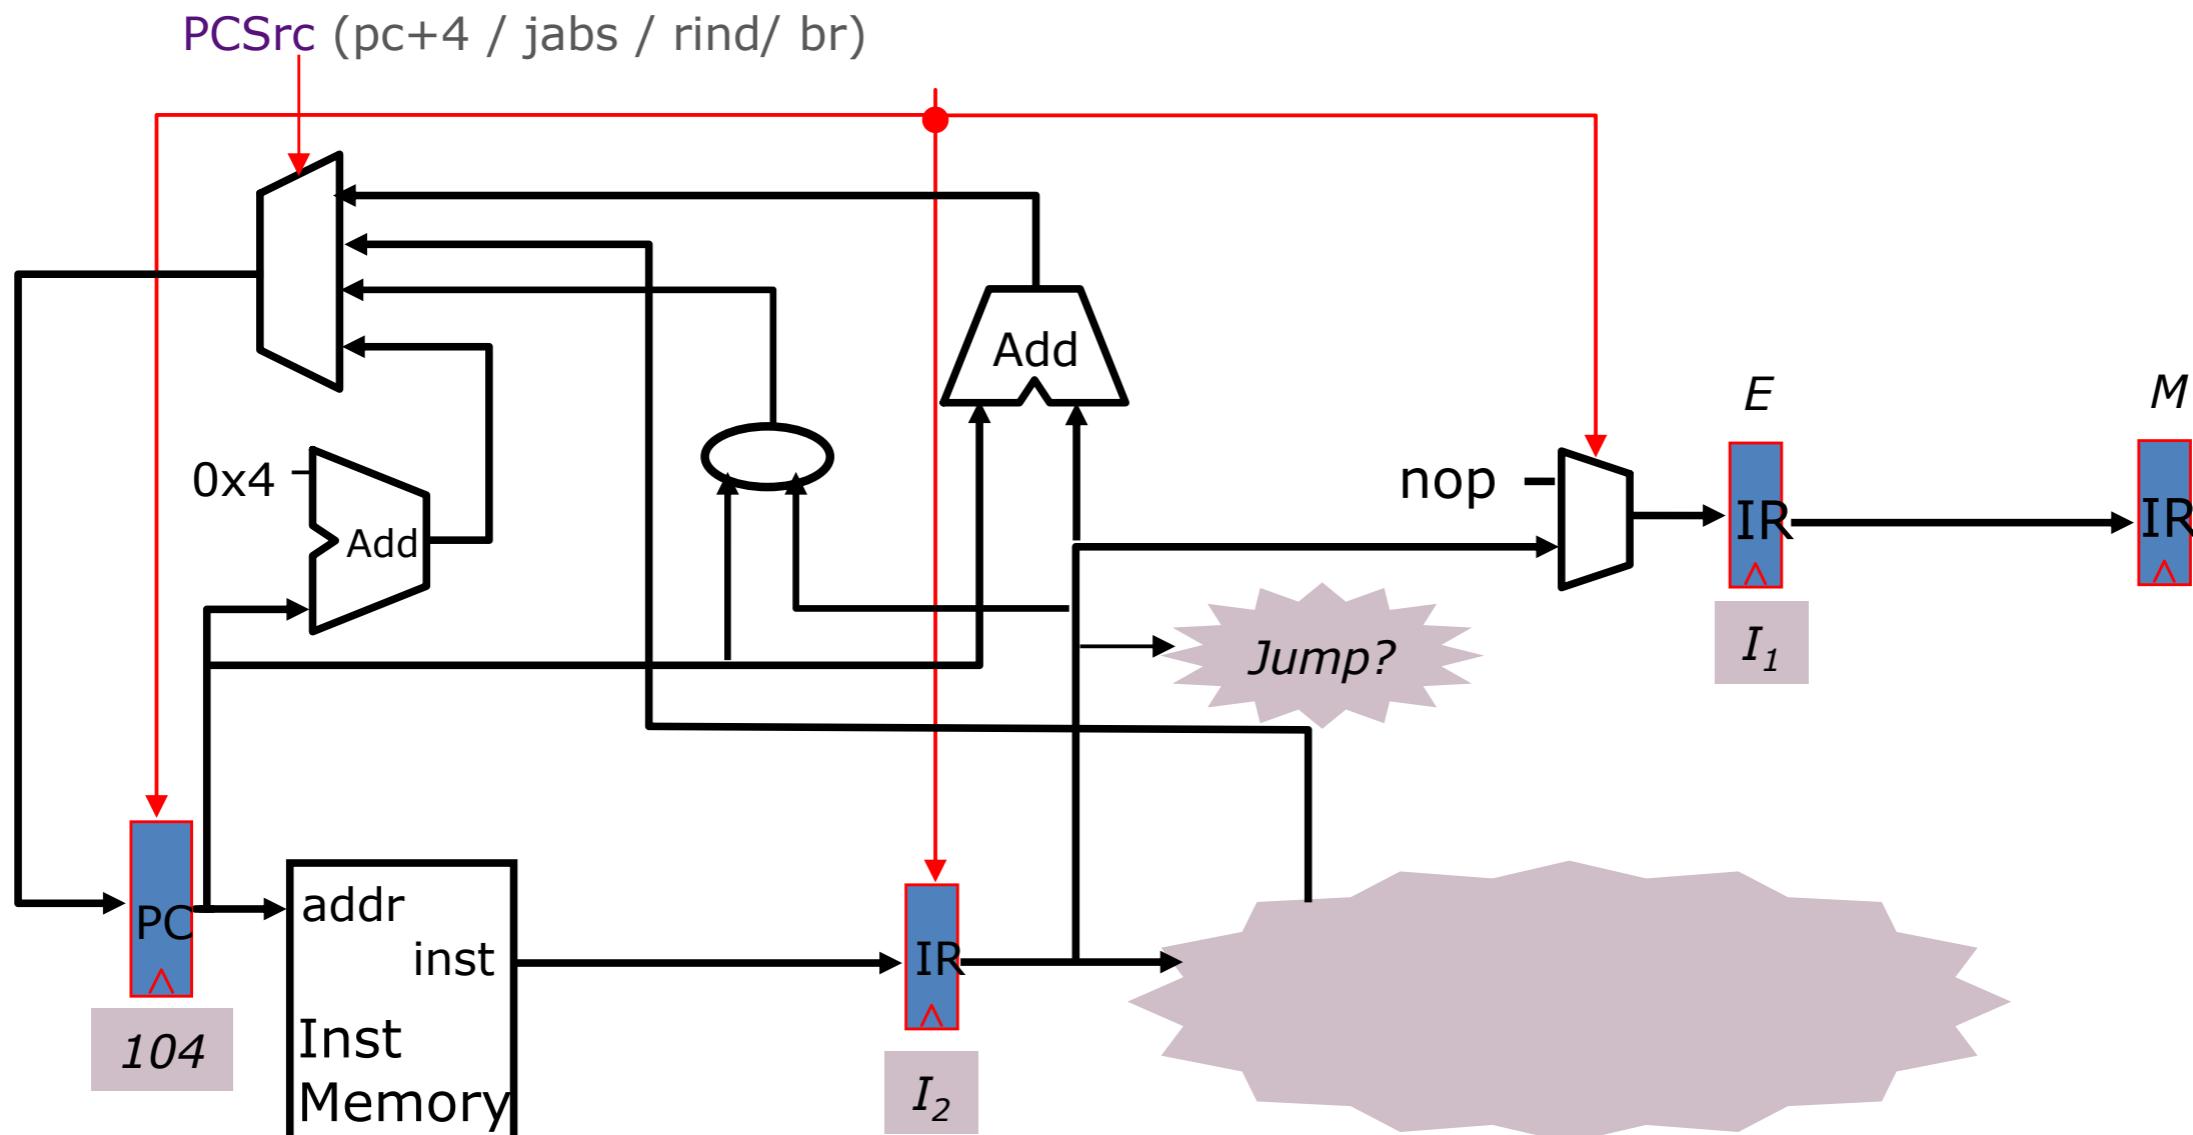
CPI = 2!

nop ⇒ *pipeline bubble*

Speculate next address is PC+4



Speculate next address is PC+4



I₁	096	ADD
I₂	100	J 304
I₃	104	ADD
I₄	304	ADD

kill

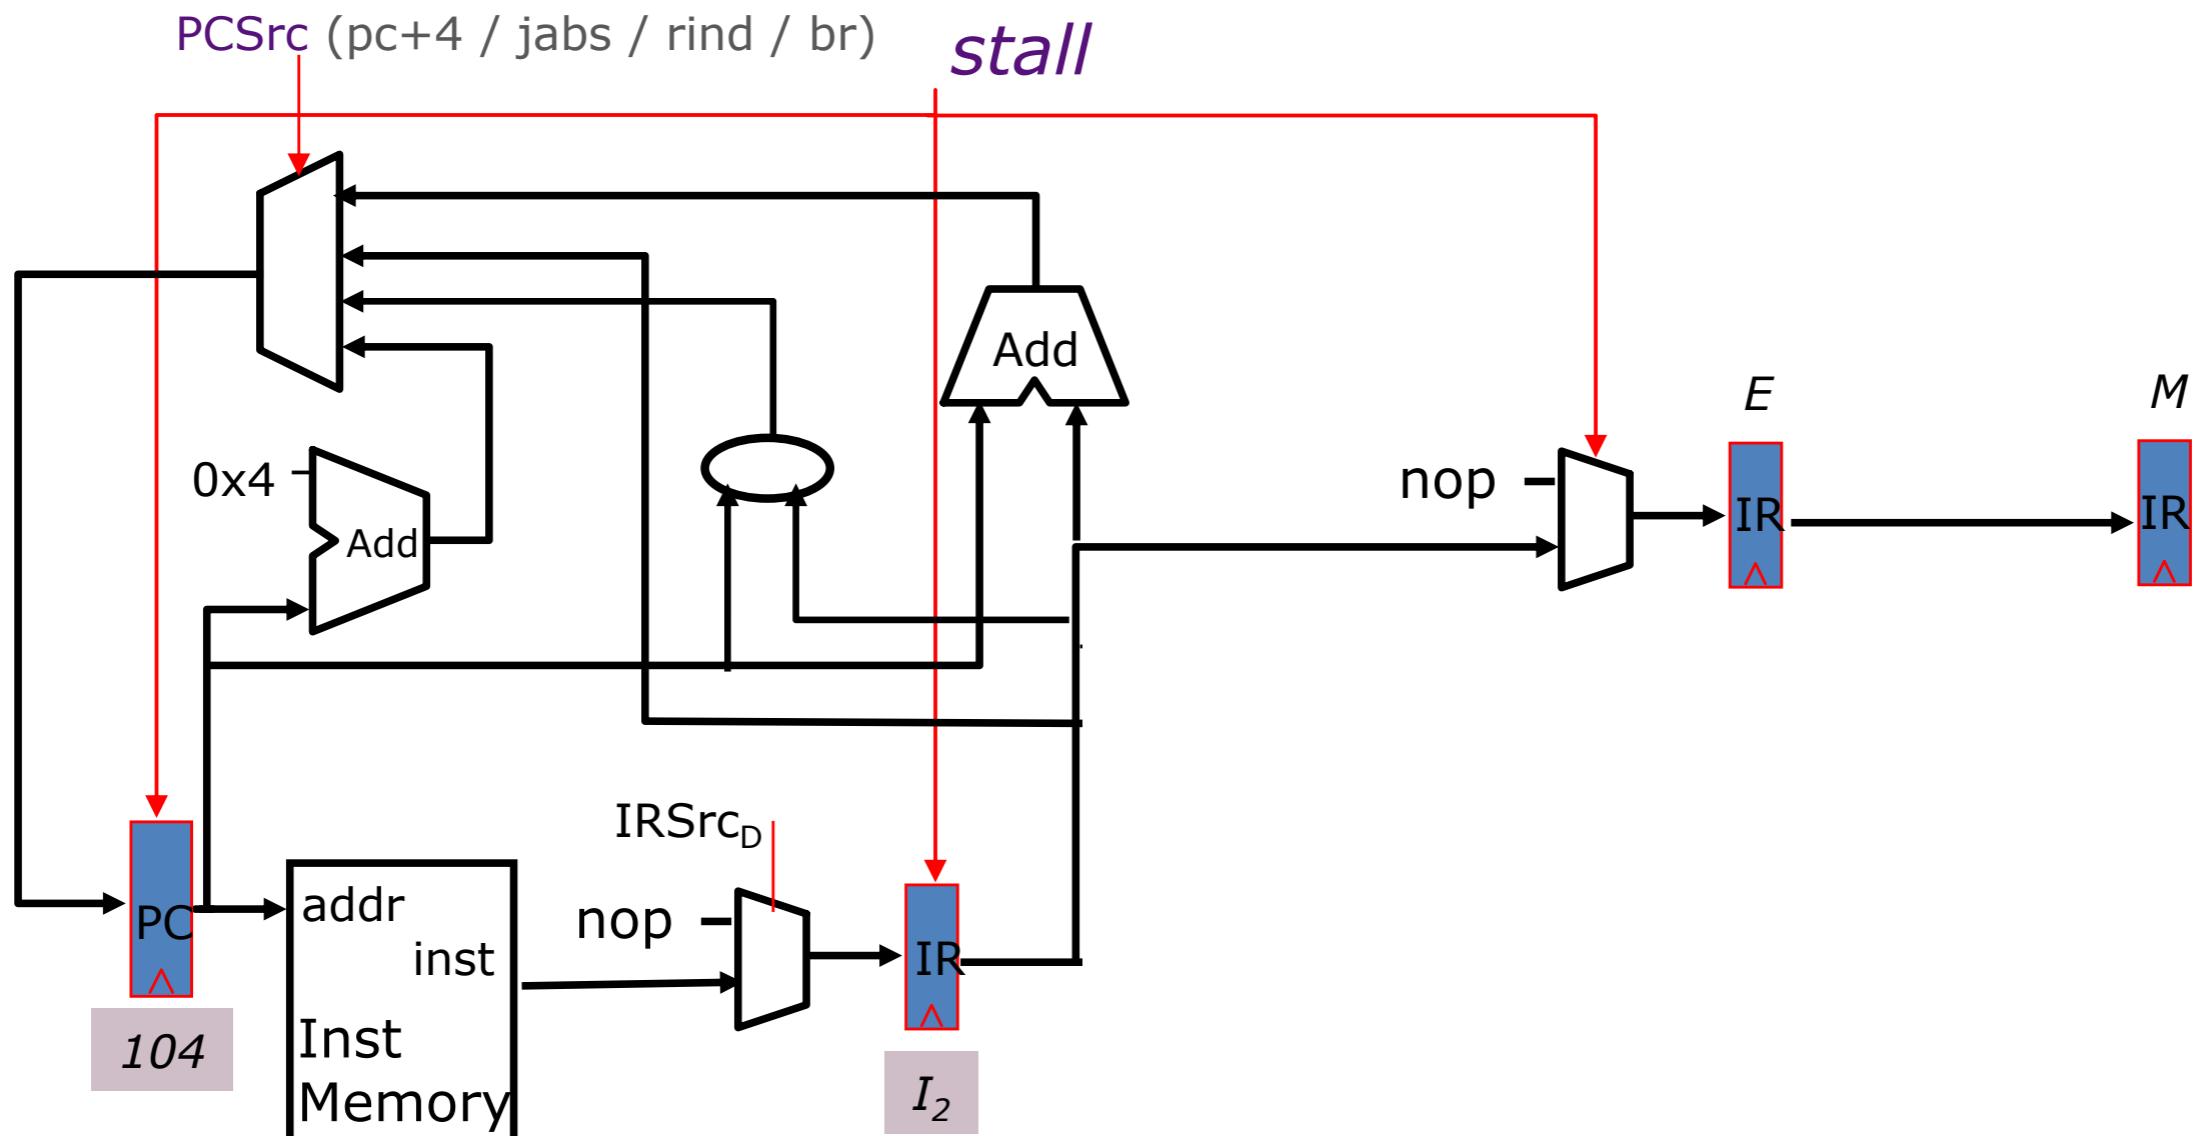
Jump Pipeline Diagrams

	<i>time</i>								
	t0	t1	t2	t3	t4	t5	t6	t7	...
(I ₁) 096: ADD	IF ₁	ID ₁	EX ₁	MA ₁	WB ₁				
(I ₂) 100: J 304		IF ₂	ID ₂	EX ₂	MA ₂	WB ₂			
(I ₃) 104: ADD		IF ₃	nop	nop	nop	nop			
(I ₄) 304: ADD		IF ₄	ID ₄	EX ₄	MA ₄	WB ₄			

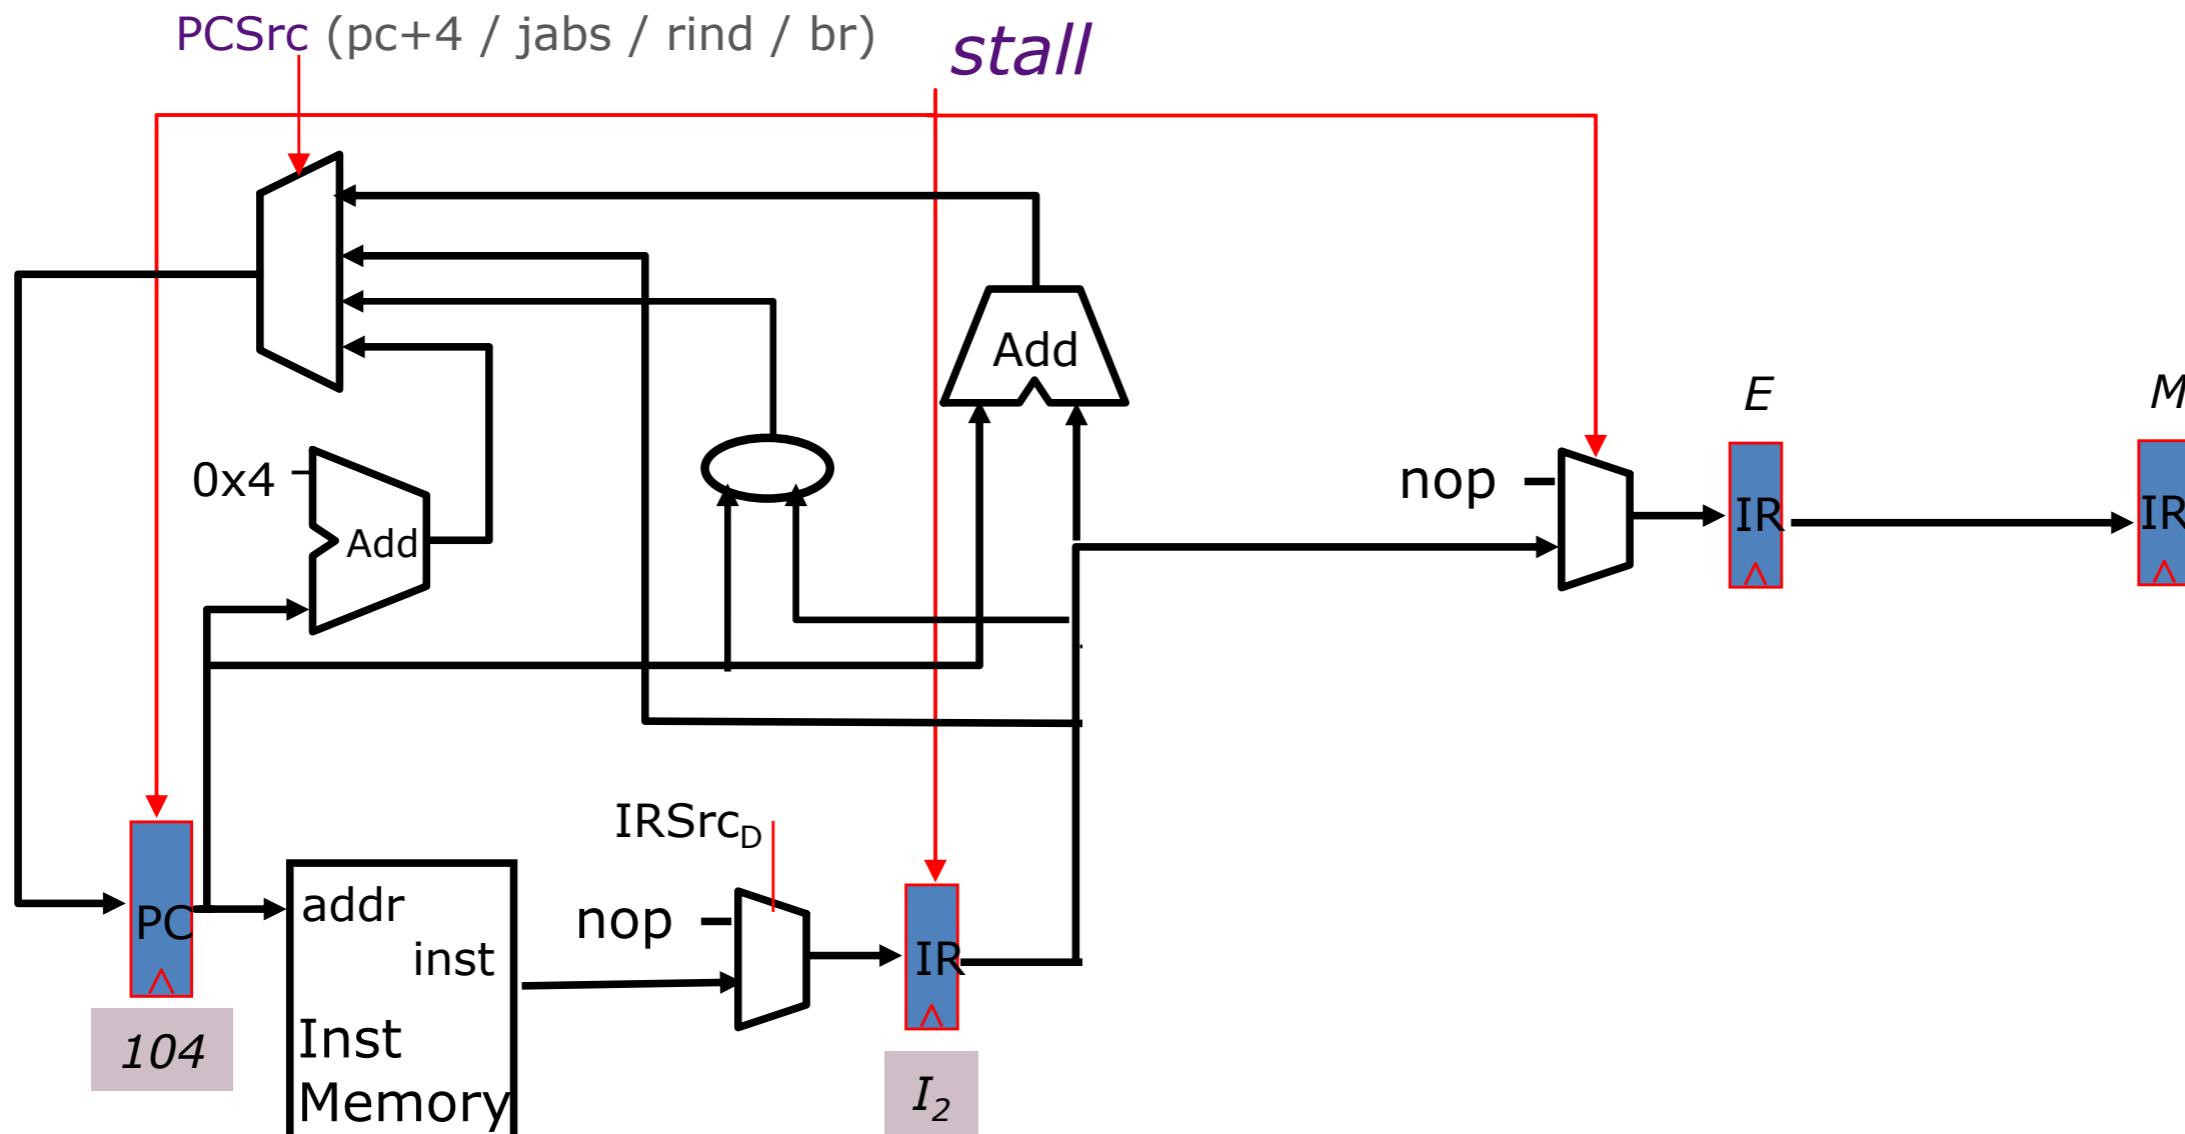
	<i>time</i>								
	t0	t1	t2	t3	t4	t5	t6	t7	...
Resource Usage	IF	I ₁	I ₂	I ₃	I ₄	I ₅			
	ID		I ₁	I ₂	nop	I ₄	I ₅		
	EX			I ₁	I ₂	nop	I ₄	I ₅	
	MA				I ₁	I ₂	nop	I ₄	I ₅
	WB					I ₁	I ₂	nop	I ₄ I ₅

nop \Rightarrow *pipeline bubble*

Pipelining Conditional Branches



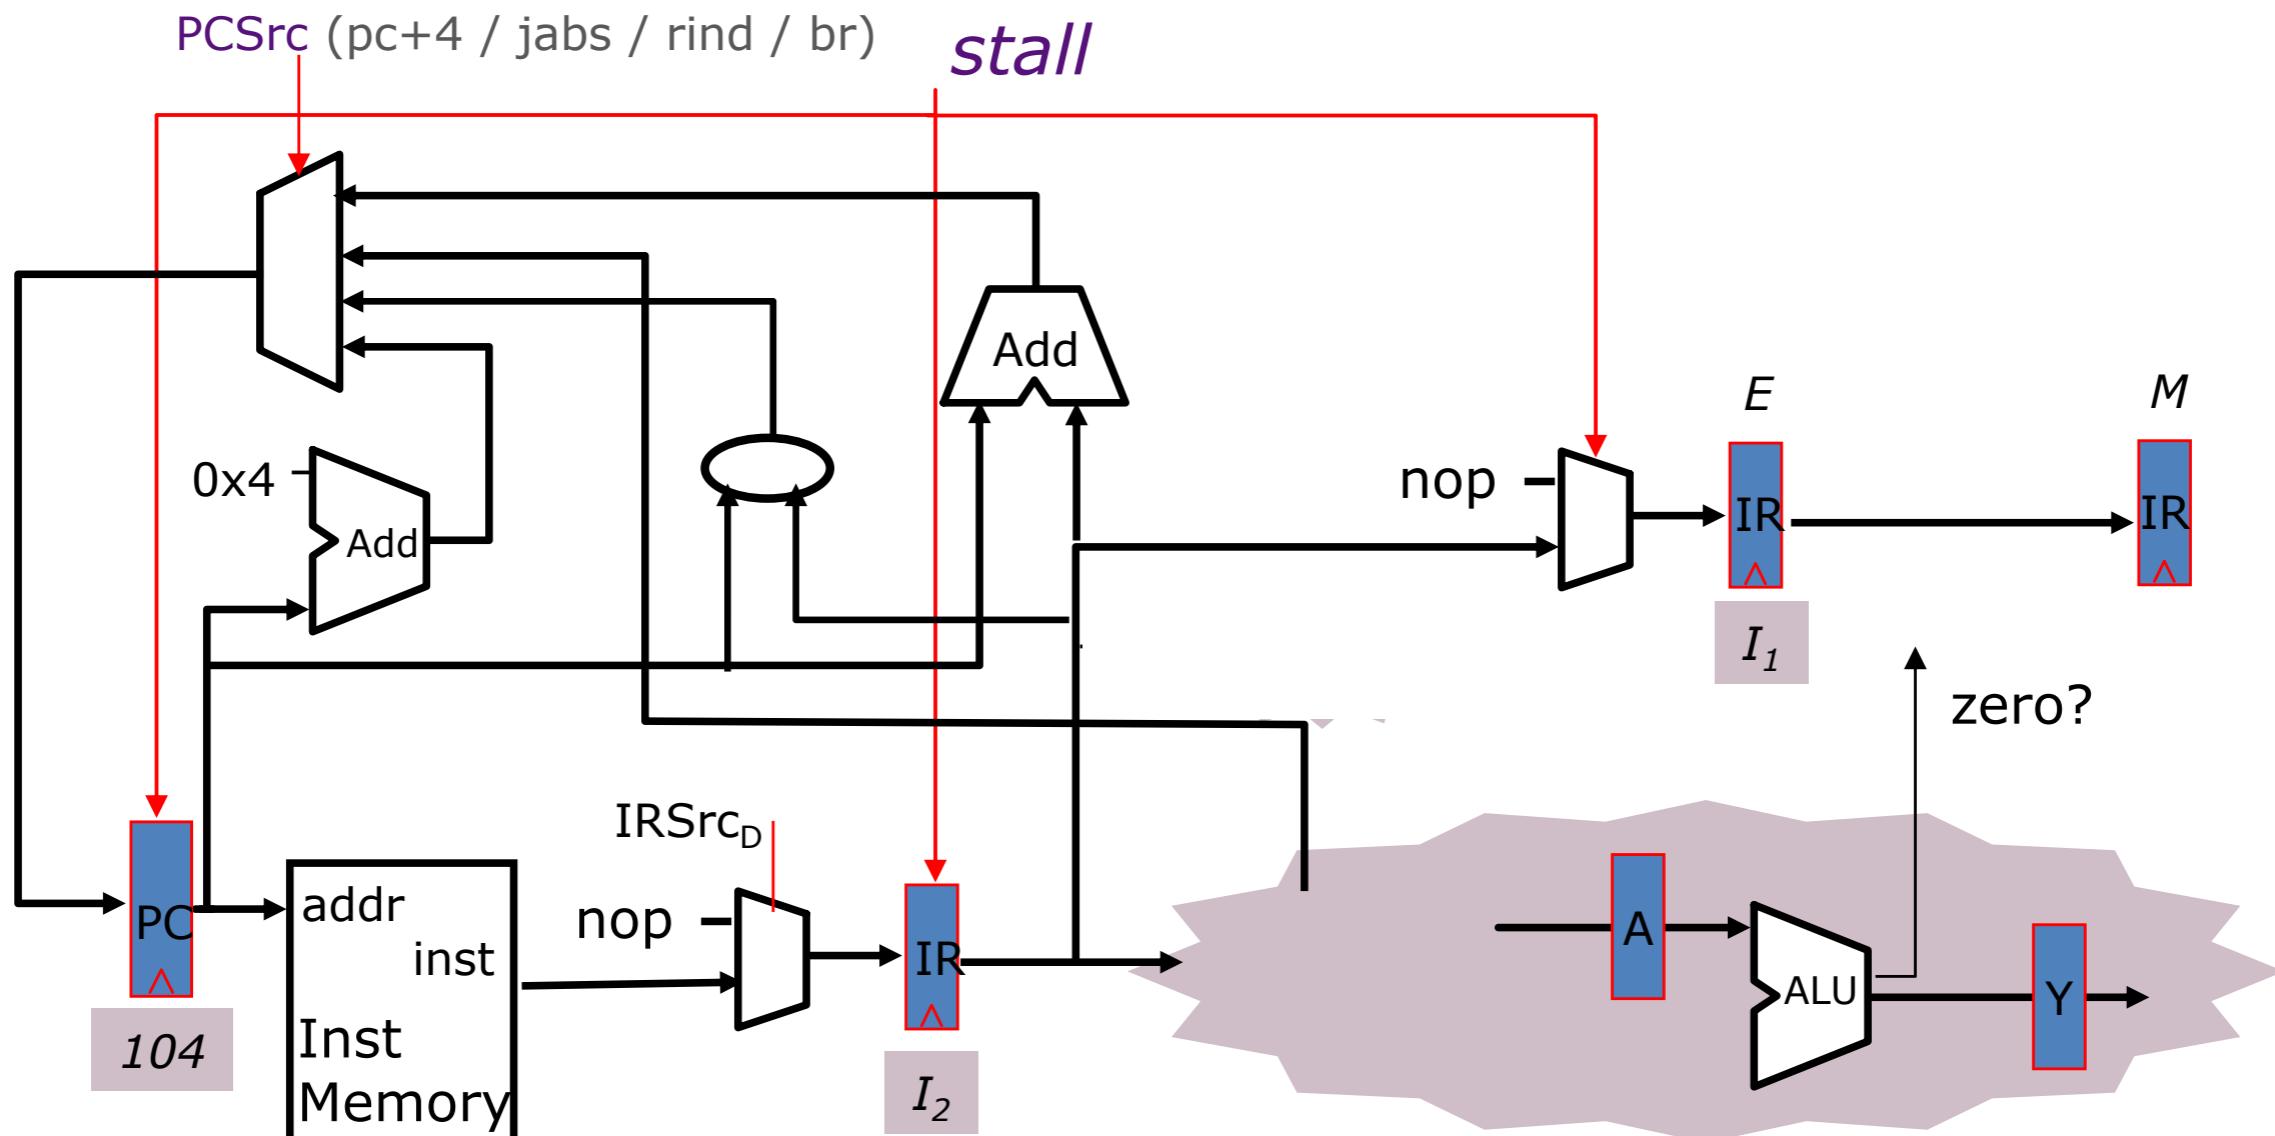
Pipelining Conditional Branches



I ₁	096	ADD
I ₂	100	BEQZ r1 +200
I ₃	104	ADD
I ₄	108	...
	304	ADD

Branch condition is not known until the execute stage
what action should be taken in the decode stage ?

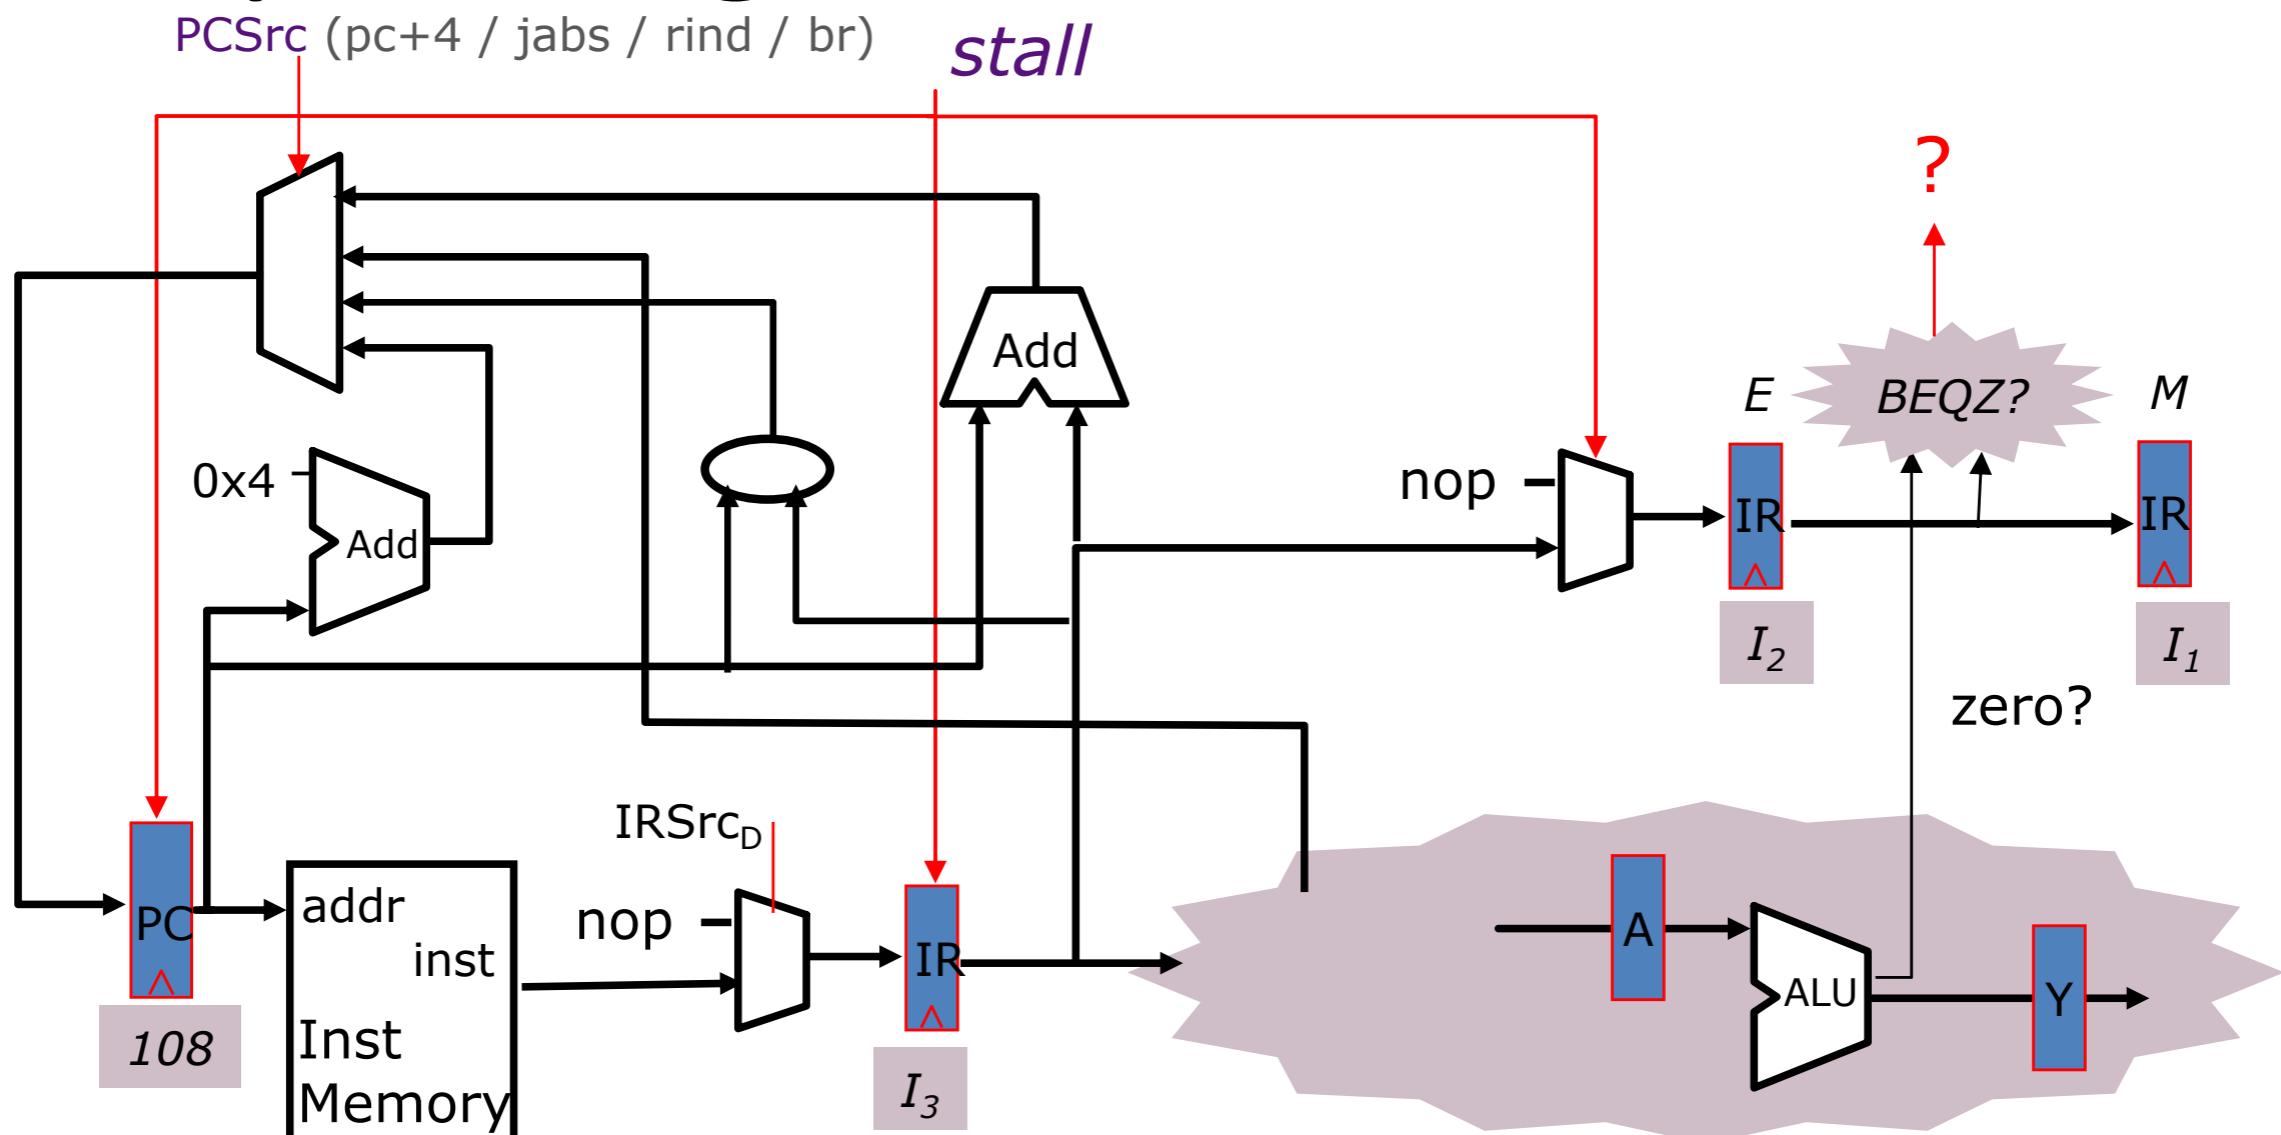
Pipelining Conditional Branches



I_1	096	ADD
I_2	100	BEQZ r1 +200
I_3	104	ADD
	108	...
I_4	304	ADD

Branch condition is not known until the execute stage
what action should be taken in the decode stage ?

Pipelining Conditional Branches



If the branch is taken

- kill the two following instructions
 - the instruction at the decode stage is not valid

```
096      ADD
100      BEQZ r1 +200
104      ADD
108      ...
304      ADD
```

Branch Pipeline Diagrams

(resolved in execute stage)

	<i>time</i>									
	t0	t1	t2	t3	t4	t5	t6	t7	...	
(I ₁) 096: ADD		IF ₁	ID ₁	EX ₁	MA ₁	WB ₁				
(I ₂) 100: BEQZ +200			IF ₂	ID ₂	EX ₂	MA ₂	WB ₂			
(I ₃) 104: ADD				IF ₃	ID ₃	nop	nop	nop		
(I ₄) 108:					IF ₄	nop	nop	nop	nop	
(I ₅) 304: ADD						IF ₅	ID ₅	EX ₅	MA ₅	WB ₅

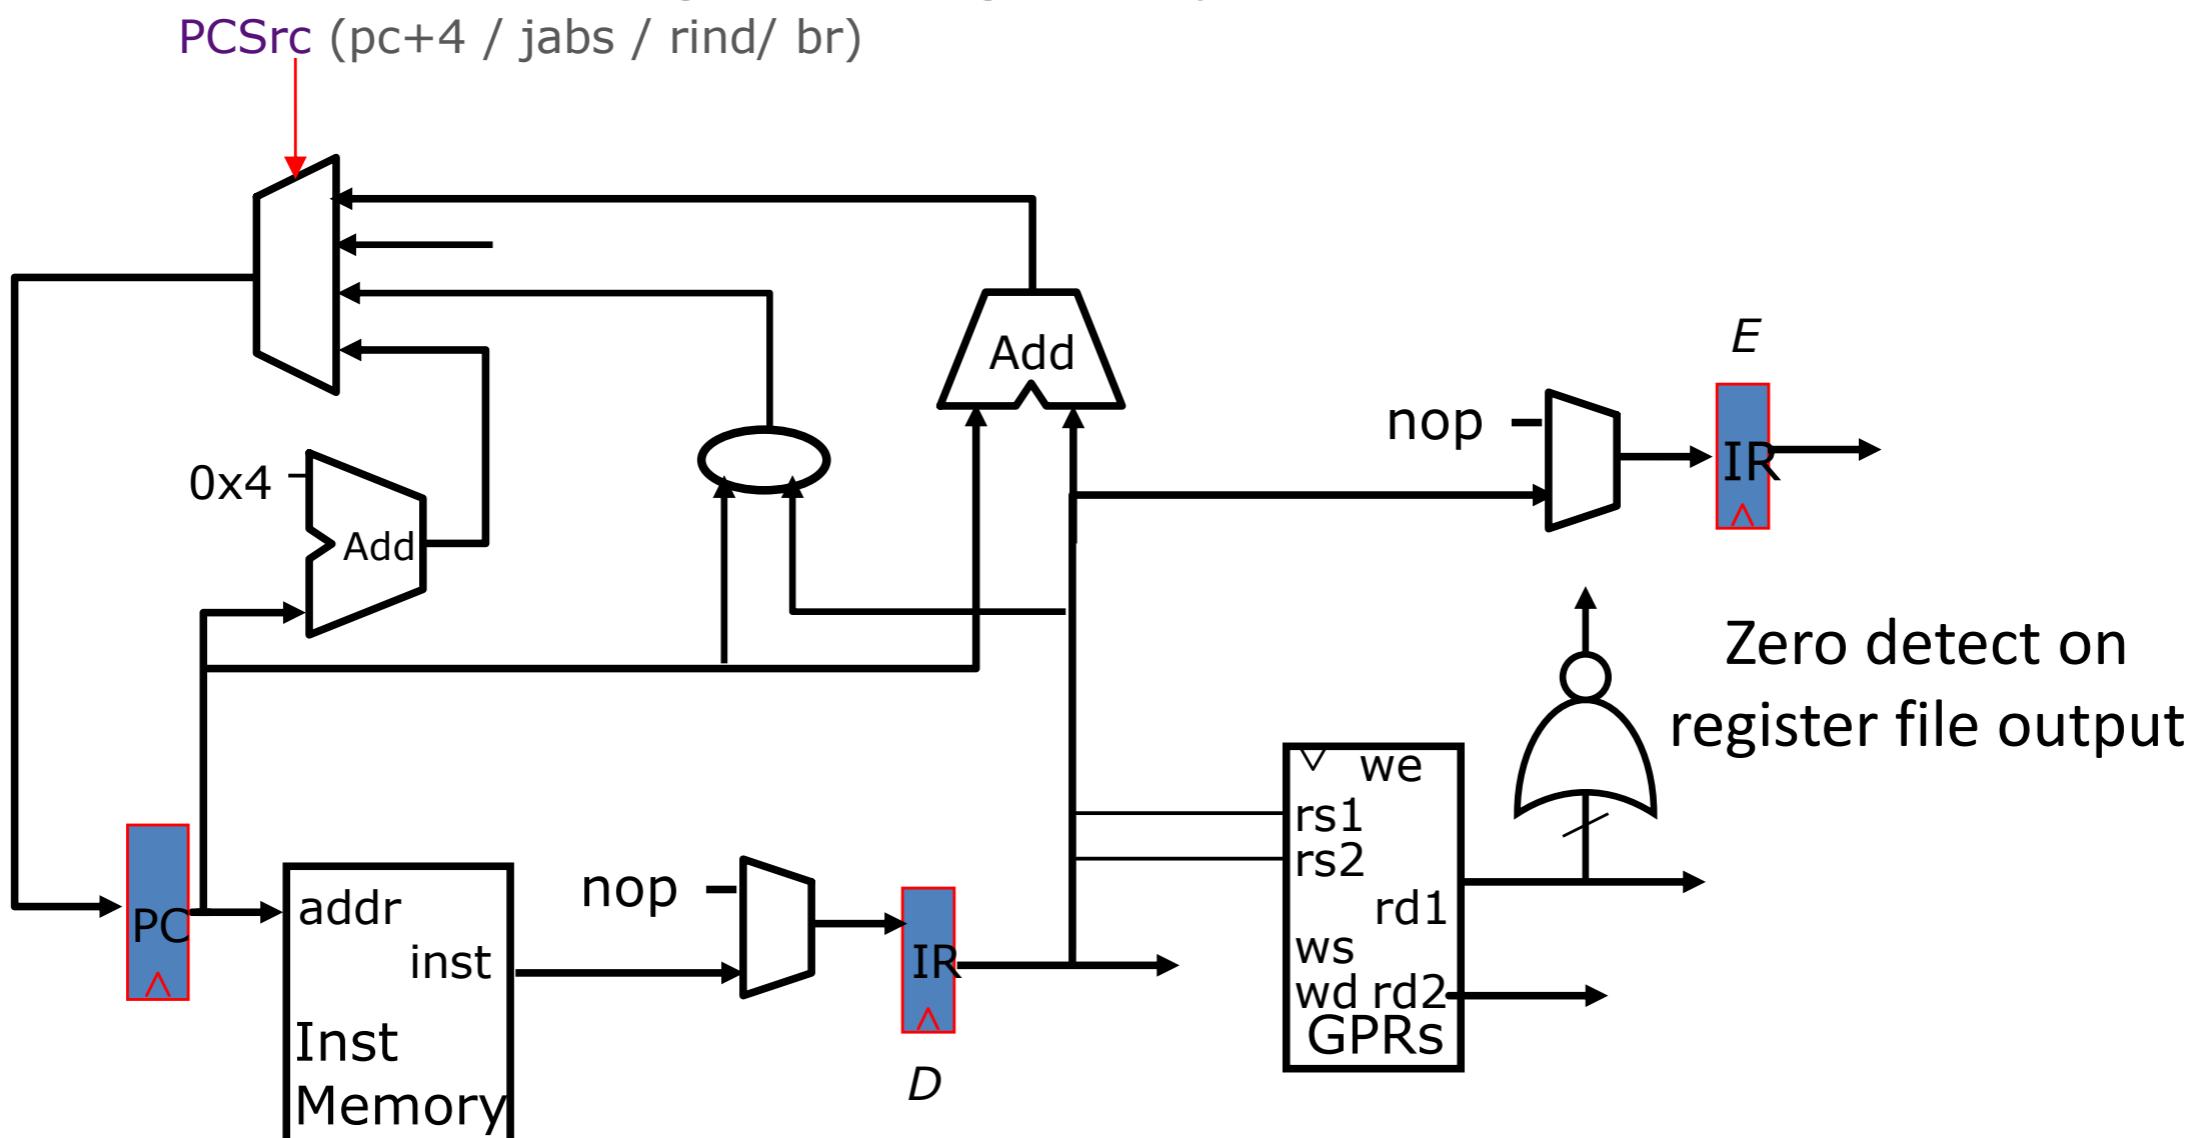
	<i>time</i>									
	t0	t1	t2	t3	t4	t5	t6	t7	...	
Resource Usage	IF	I ₁	I ₂	I ₃	I ₄	I ₅				
	ID		I ₁	I ₂	I ₃	nop	I ₅			
	EX			I ₁	I ₂	nop	nop	I ₅		
	MA				I ₁	I ₂	nop	nop	I ₅	
	WB					I ₁	I ₂	nop	nop	I ₅

nop \Rightarrow *pipeline bubble*

Reducing Branch Penalty

(resolve in decode stage)

- One pipeline bubble can be removed if an extra comparator is used in the Decode stage
 - But might elongate cycle time



Pipeline diagram now same as for jumps

Branch Delay Slots

(expose control hazard to software)

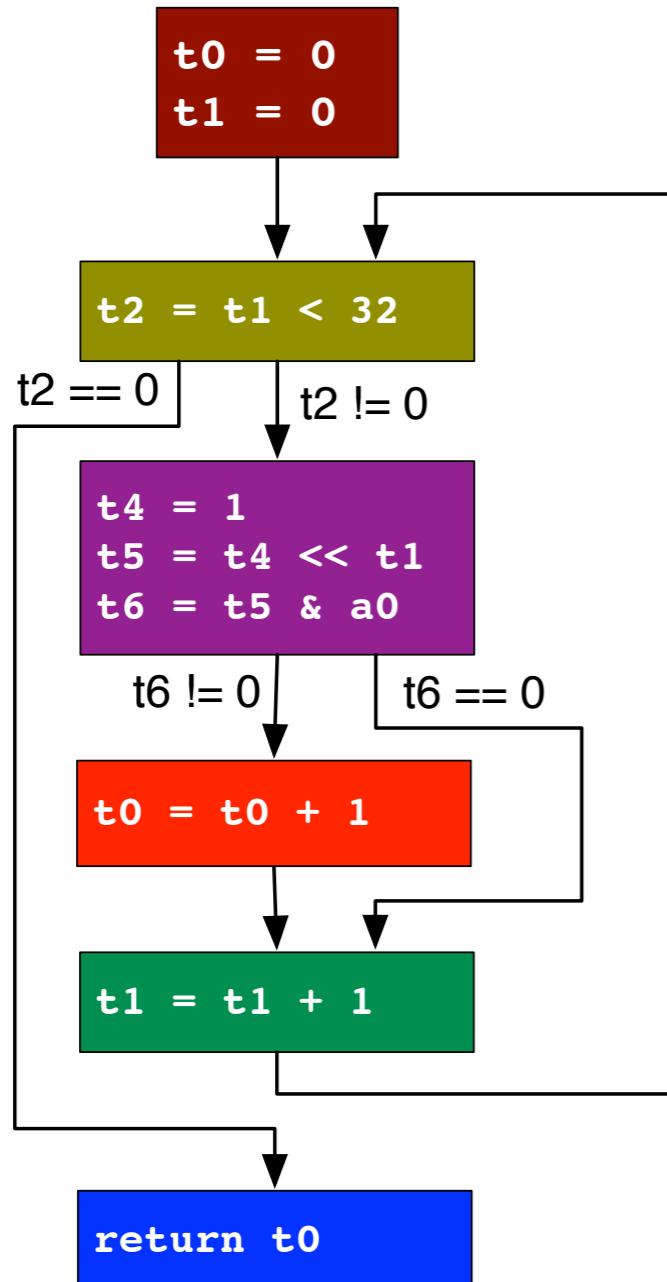
- Change the ISA semantics so that the instruction that follows a jump or branch is always executed
 - gives compiler the flexibility to put in a useful instruction where normally a pipeline bubble would have resulted.

I ₁	096	ADD
I ₂	100	BEQZ r1 +200
I ₃	104	ADD
I ₄	304	ADD

*Delay slot instruction executed
regardless of branch outcome*

- Other techniques include more advanced branch prediction, which can dramatically reduce the branch penalty... *to come later*

In the Compiler



```
popcount:  
    ori $v0, $zero, 0  
    ori $t1, $zero, 0  
  
top:  
    slti $t2, $t1, 32  
    beq $t2, $zero, end  
    nop  
    addi $t3, $zero, 1  
    sllv $t3, $t3, $t1  
    and $t3, $a0, $t3  
    beq $t3, $zero, notone  
    nop  
    addi $v0, $v0, 1  
  
notone:  
    beq $zero, $zero, top  
    addi $t1, $t1, 1  
  
end:  
    jr $ra  
    nop
```

Control flow graph

Assembly

Branch Pipeline Diagrams

(branch delay slot)

	<i>time</i>								
	t0	t1	t2	t3	t4	t5	t6	t7	...
(I ₁) 096: ADD		IF ₁	ID ₁	EX ₁	MA ₁	WB ₁			
(I ₂) 100: BEQZ +200			IF ₂	ID ₂	EX ₂	MA ₂	WB ₂		
(I ₃) 104: ADD				IF ₃	ID ₃	EX ₃	MA ₃	WB ₃	
(I ₄) 304: ADD					IF ₄	ID ₄	EX ₄	MA ₄	WB ₄

	<i>time</i>								
	t0	t1	t2	t3	t4	t5	t6	t7	...
Resource Usage	IF	I ₁	I ₂	I ₃	I ₄				
	ID		I ₁	I ₂	I ₃	I ₄			
	EX			I ₁	I ₂	I ₃	I ₄		
	MA				I ₁	I ₂	I ₃	I ₄	
	WB					I ₁	I ₂	I ₃	I ₄

Scheduling in Action

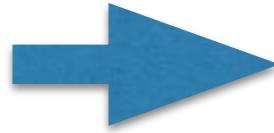
- Suppose every branch brings two stalls, how do you schedule the following instructions?

```
SUB t2, t3, t5  
ADD t4, t3, t5  
bneqz t4, 400  
ADD t1, t3, t8  
ADD t3, t5, t9  
beqz t4, 500
```

Scheduling in Action

- Suppose every branch brings two stalls, how do you schedule the following instructions?

SUB t2, t3, t5
ADD t4, t3, t5
bneqz t4, 400
ADD t1, t3, t8
ADD t3, t5, t9
beqz t4, 500

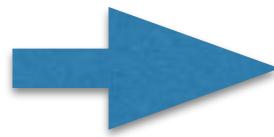


ADD t4, t3, t5
bneqz t4, 400
SUB t2, t3, t5
beqz t4, 500
ADD t1, t3, t8
ADD t3, t5, t9

Scheduling in Action

- Suppose every branch brings two stalls, how do you schedule the following instructions?

SUB t2, t3, t5
ADD t4, t3, t5
bneqz t4, 400
ADD t1, t3, t8
ADD t3, t5, t9
beqz t4, 500



ADD t4, t3, t5
bneqz t4, 400
SUB t2, t3, t5
beqz t4, 500 ← stall
ADD t1, t3, t8
ADD t3, t5, t9

Exceptions

- Causes
 - Arithmetic overflow
 - In an 8-bit machine, what if you do $255+1$?
 - Undefined instruction
 - System call
- When to handle
 - when detected
- Who should handle
 - Process

Interrupts

- Causes
 - external events
 - arrival of network package, hard disk, ...
- When to handle
 - when convenient except for high priority ones
- Who should handle
 - System

Precise exceptions/interrupts

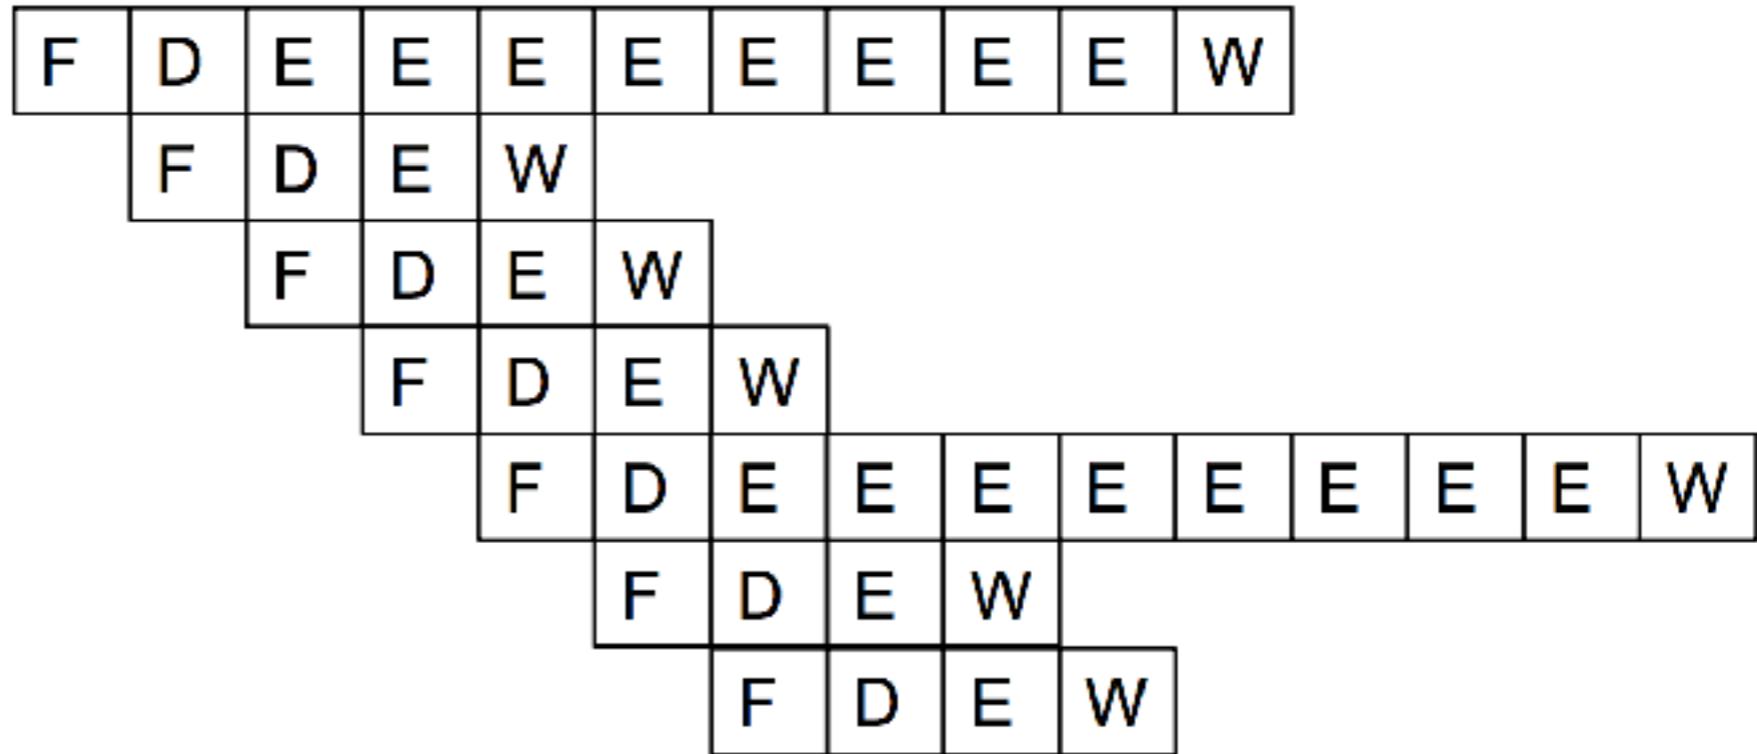
- The architectural state should be consistent when the exception/interrupt is ready to be handled
 - All previous instructions should be completely retired.
 - No later instruction should be retired.

Retire = commit = finish execution and update arch. state

Multi-cycle execute

FMUL R4 \leftarrow R1, R2
ADD R3 \leftarrow R1, R2

FMUL R2 \leftarrow R5, R6
ADD R4 \leftarrow R5, R6

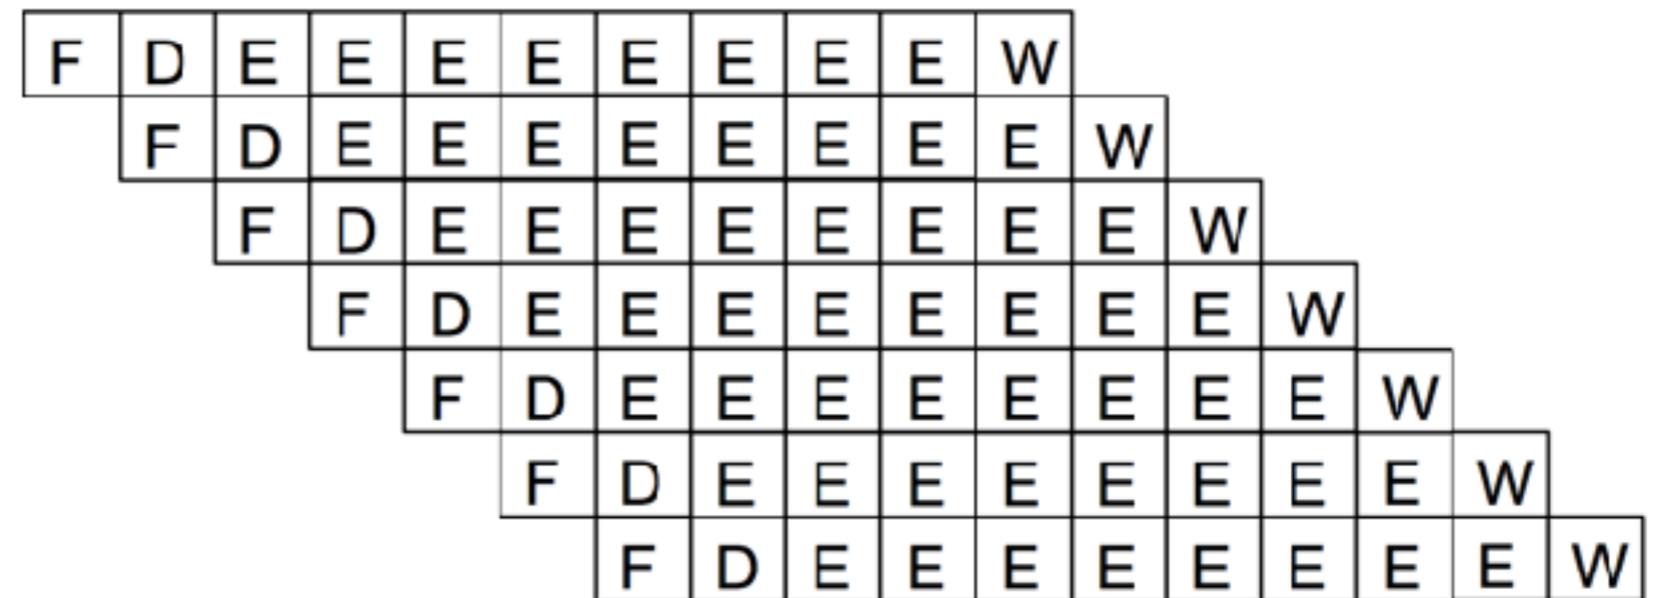


- instructions may take multiple cycles in ALU
- What's wrong here?
 - what if floating point ALU incurs an exception?

Ensuring precise exceptions in pipelining

- Idea 1: make each operation take the same amount of time

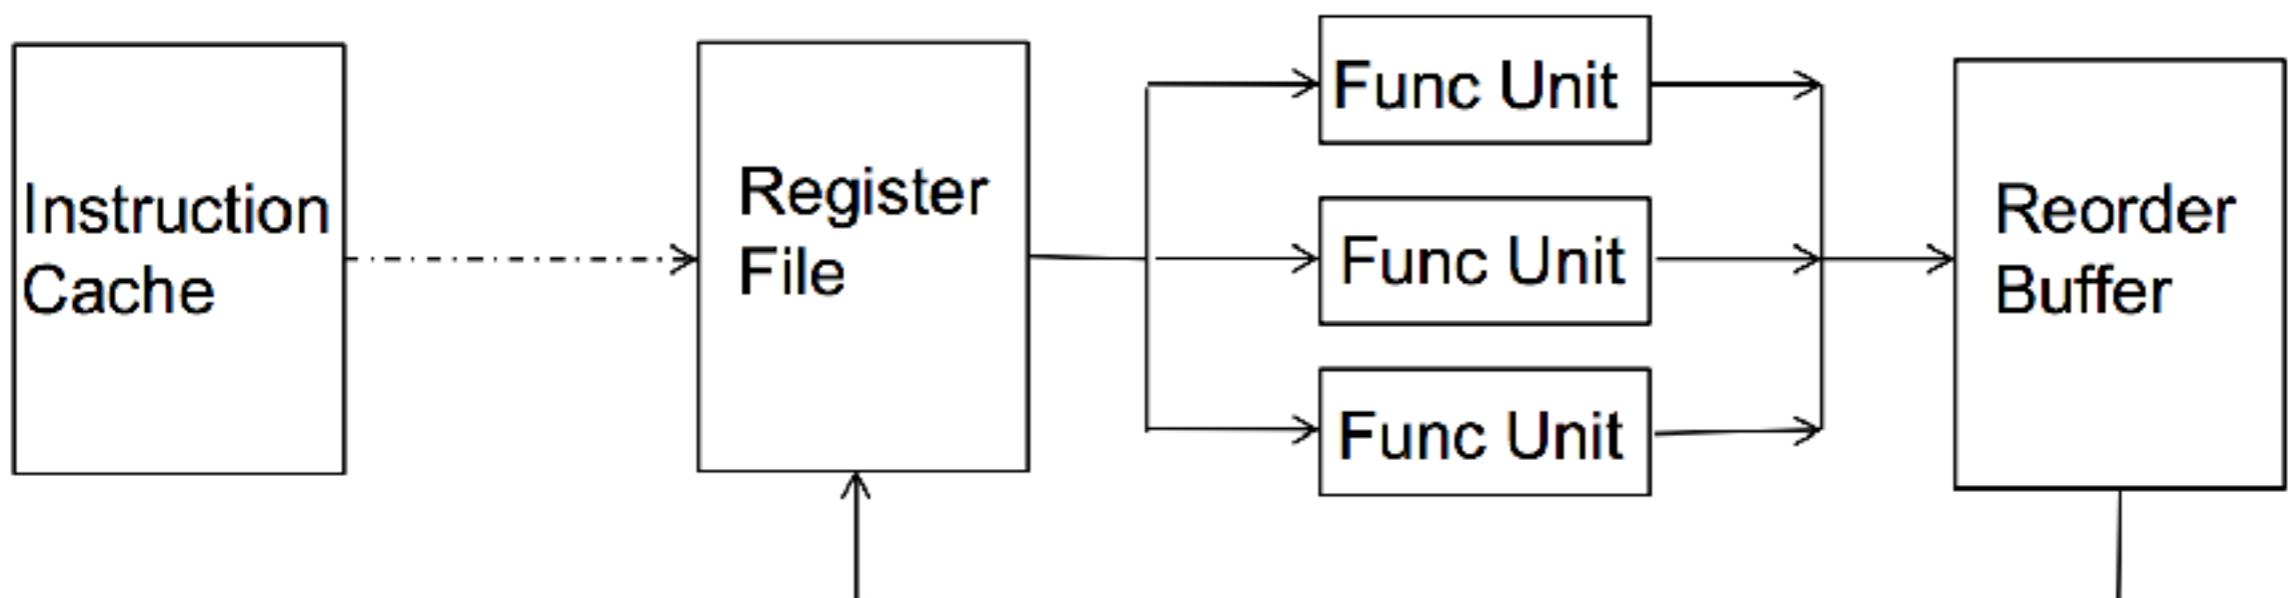
FMUL R3 \leftarrow R1, R2
ADD R4 \leftarrow R1, R2



- Downside
 - worse-case latency determines all instructions' latency
 - chance is high for structural hazards

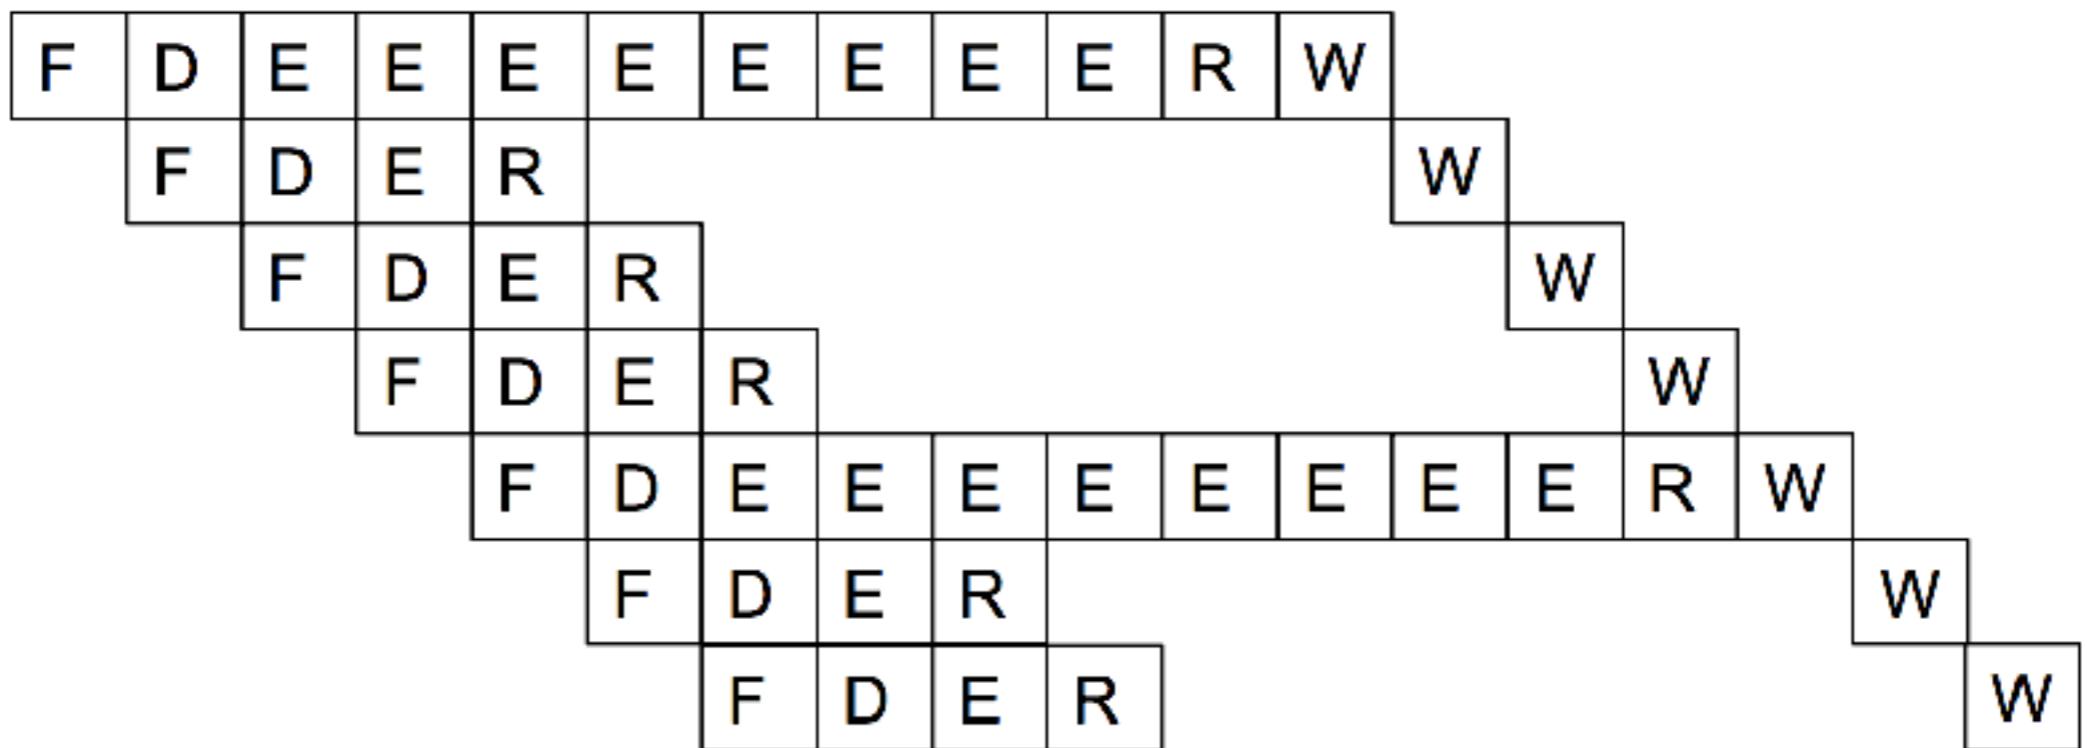
Ensuring precise exceptions in pipelining

- Idea 2: Reorder buffer (ROB)
 - Complete instructions out-of-order, but reorder them before making results visible to architectural state
 - When instruction is decoded it reserves an entry in the ROB
 - When instruction completes, it writes result into ROB entry
 - When instruction oldest in ROB and it has completed without exceptions, its result moved to reg. file



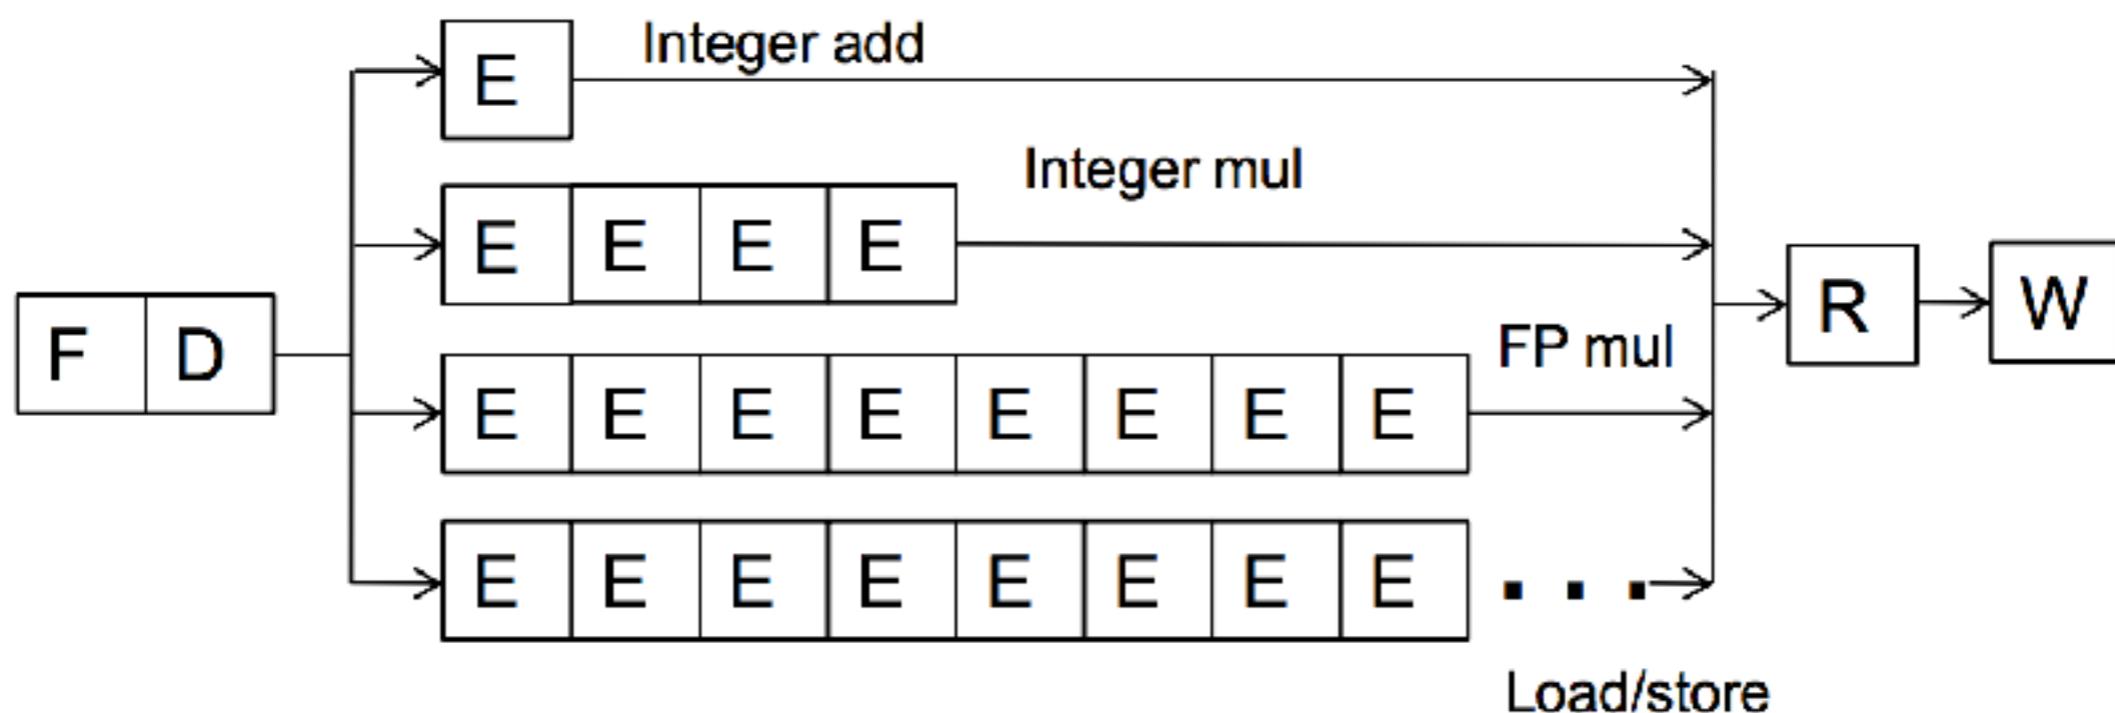
Reorder buffer instruction flow

- Results first written to ROB, then to register file at commit time



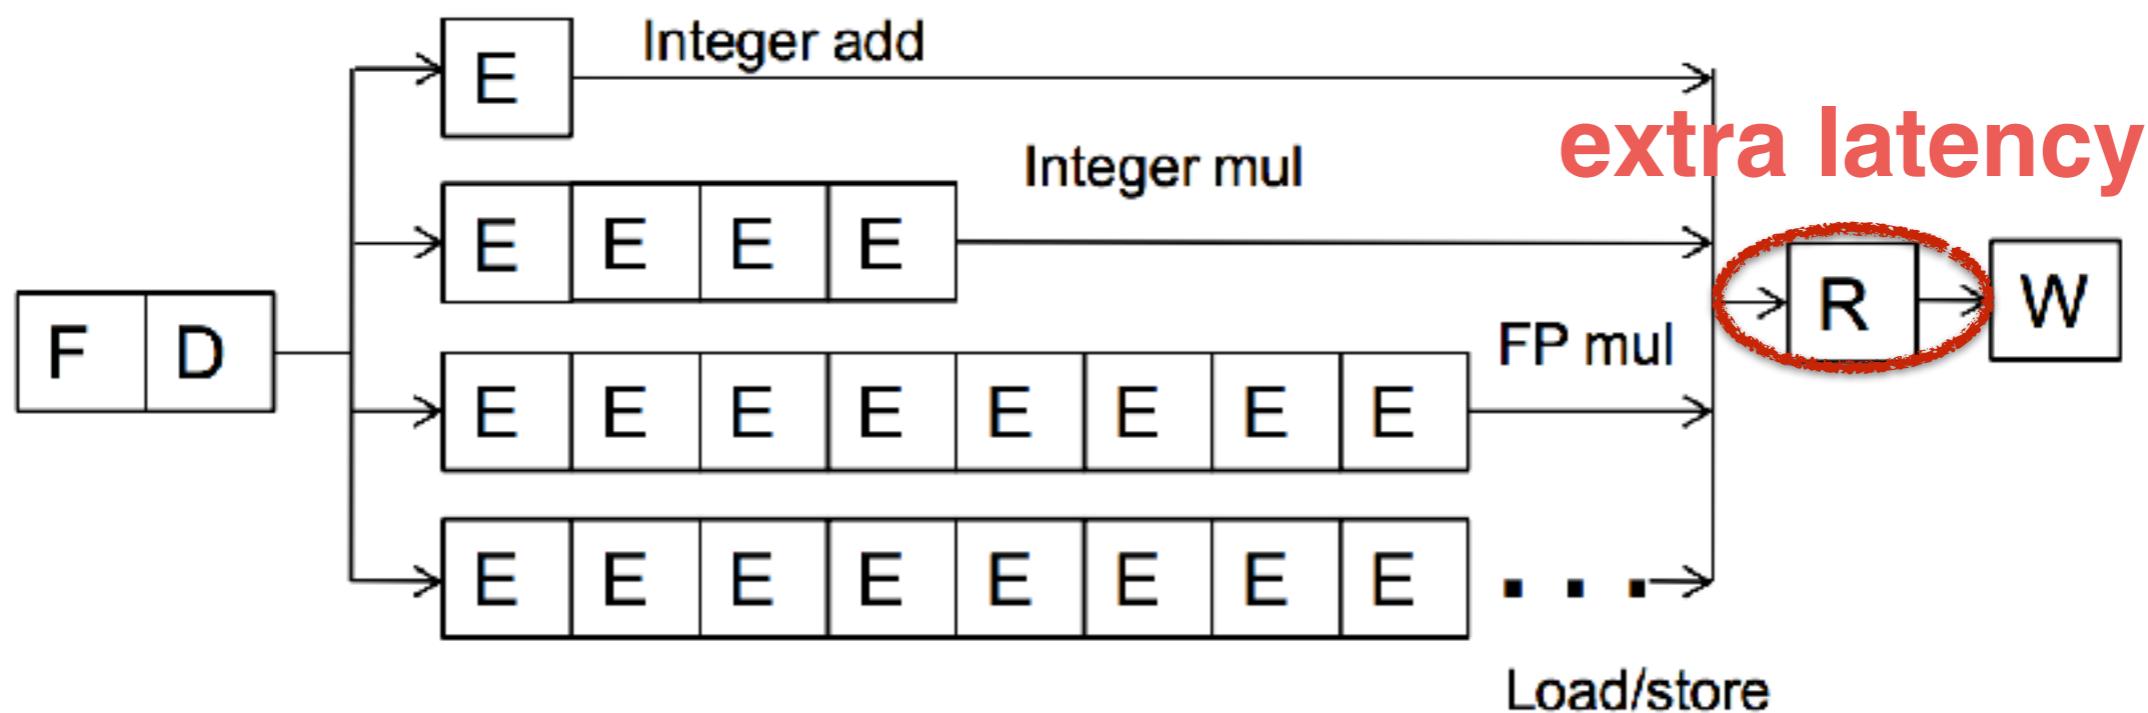
In-order pipeline with reorder buffer

- **Decode (D):** Access regfile/ROB, allocate entry in ROB, and dispatch instruction
- **Execute (E):** Instructions can complete out-of-order
- **Completion (R):** Write result to reorder buffer
- **Retirement/Commit (W):** Check for exceptions; if none, write result to architectural register file or memory; else, flush pipeline and start from exception handler
- In-order dispatch/execution, out-of-order completion, in-order retirement



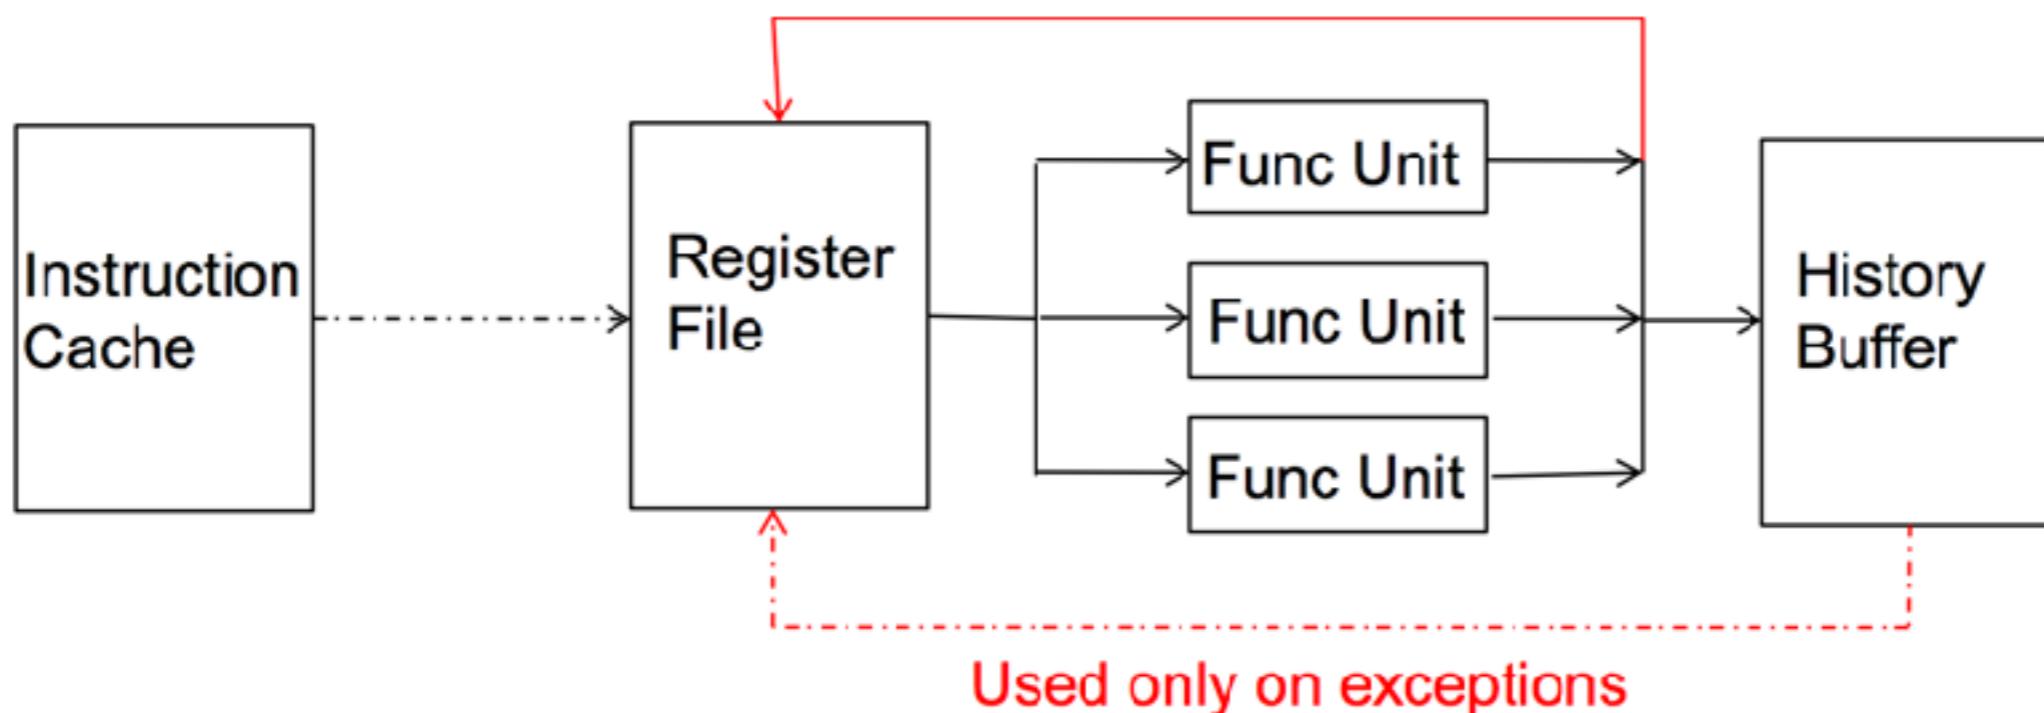
In-order pipeline with reorder buffer

- **Decode (D):** Access regfile/ROB, allocate entry in ROB, and dispatch instruction
- **Execute (E):** Instructions can complete out-of-order
- **Completion (R):** Write result to reorder buffer
- **Retirement/Commit (W):** Check for exceptions; if none, write result to architectural register file or memory; else, flush pipeline and start from exception handler
- In-order dispatch/execution, out-of-order completion, in-order retirement



Ensuring precise exceptions in pipelining

- Idea 3: History buffer (HB)
 - When instruction is decoded, it reserves an HB entry
 - When the instruction completes, it stores the old value of its destination in the HB
 - When instruction is oldest and no exceptions interrupts, the HB entry discarded
 - When instruction is oldest and an exception needs to be handled, old values in the HB are written back into the architectural state from tail to head



**Superscalar: the essential concept that differentiates
undergrad and graduate students**

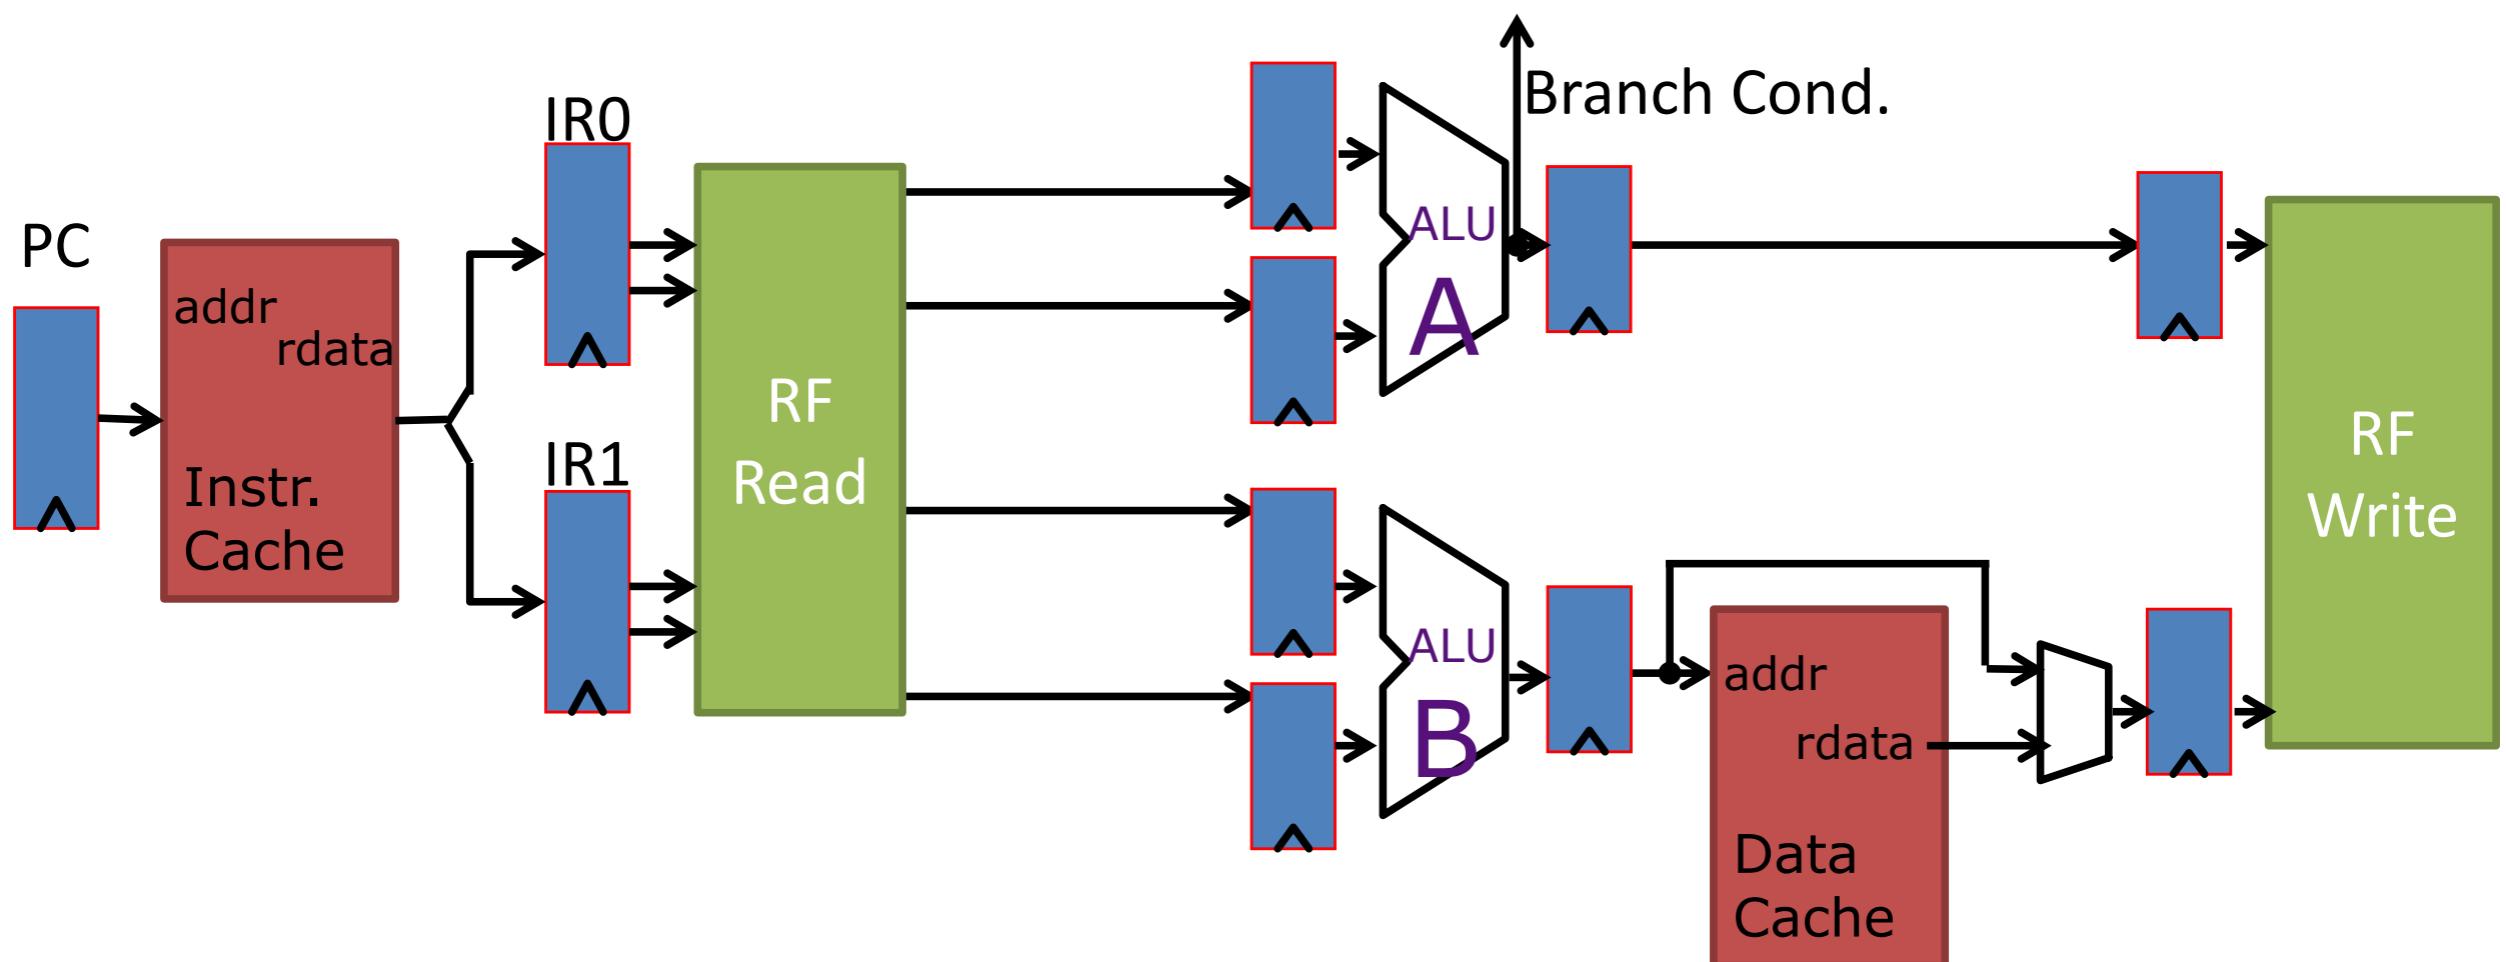
Introduction to Superscalar Processor

- Processors studied so far are fundamentally limited to $CPI \geq 1$

Introduction to Superscalar Processor

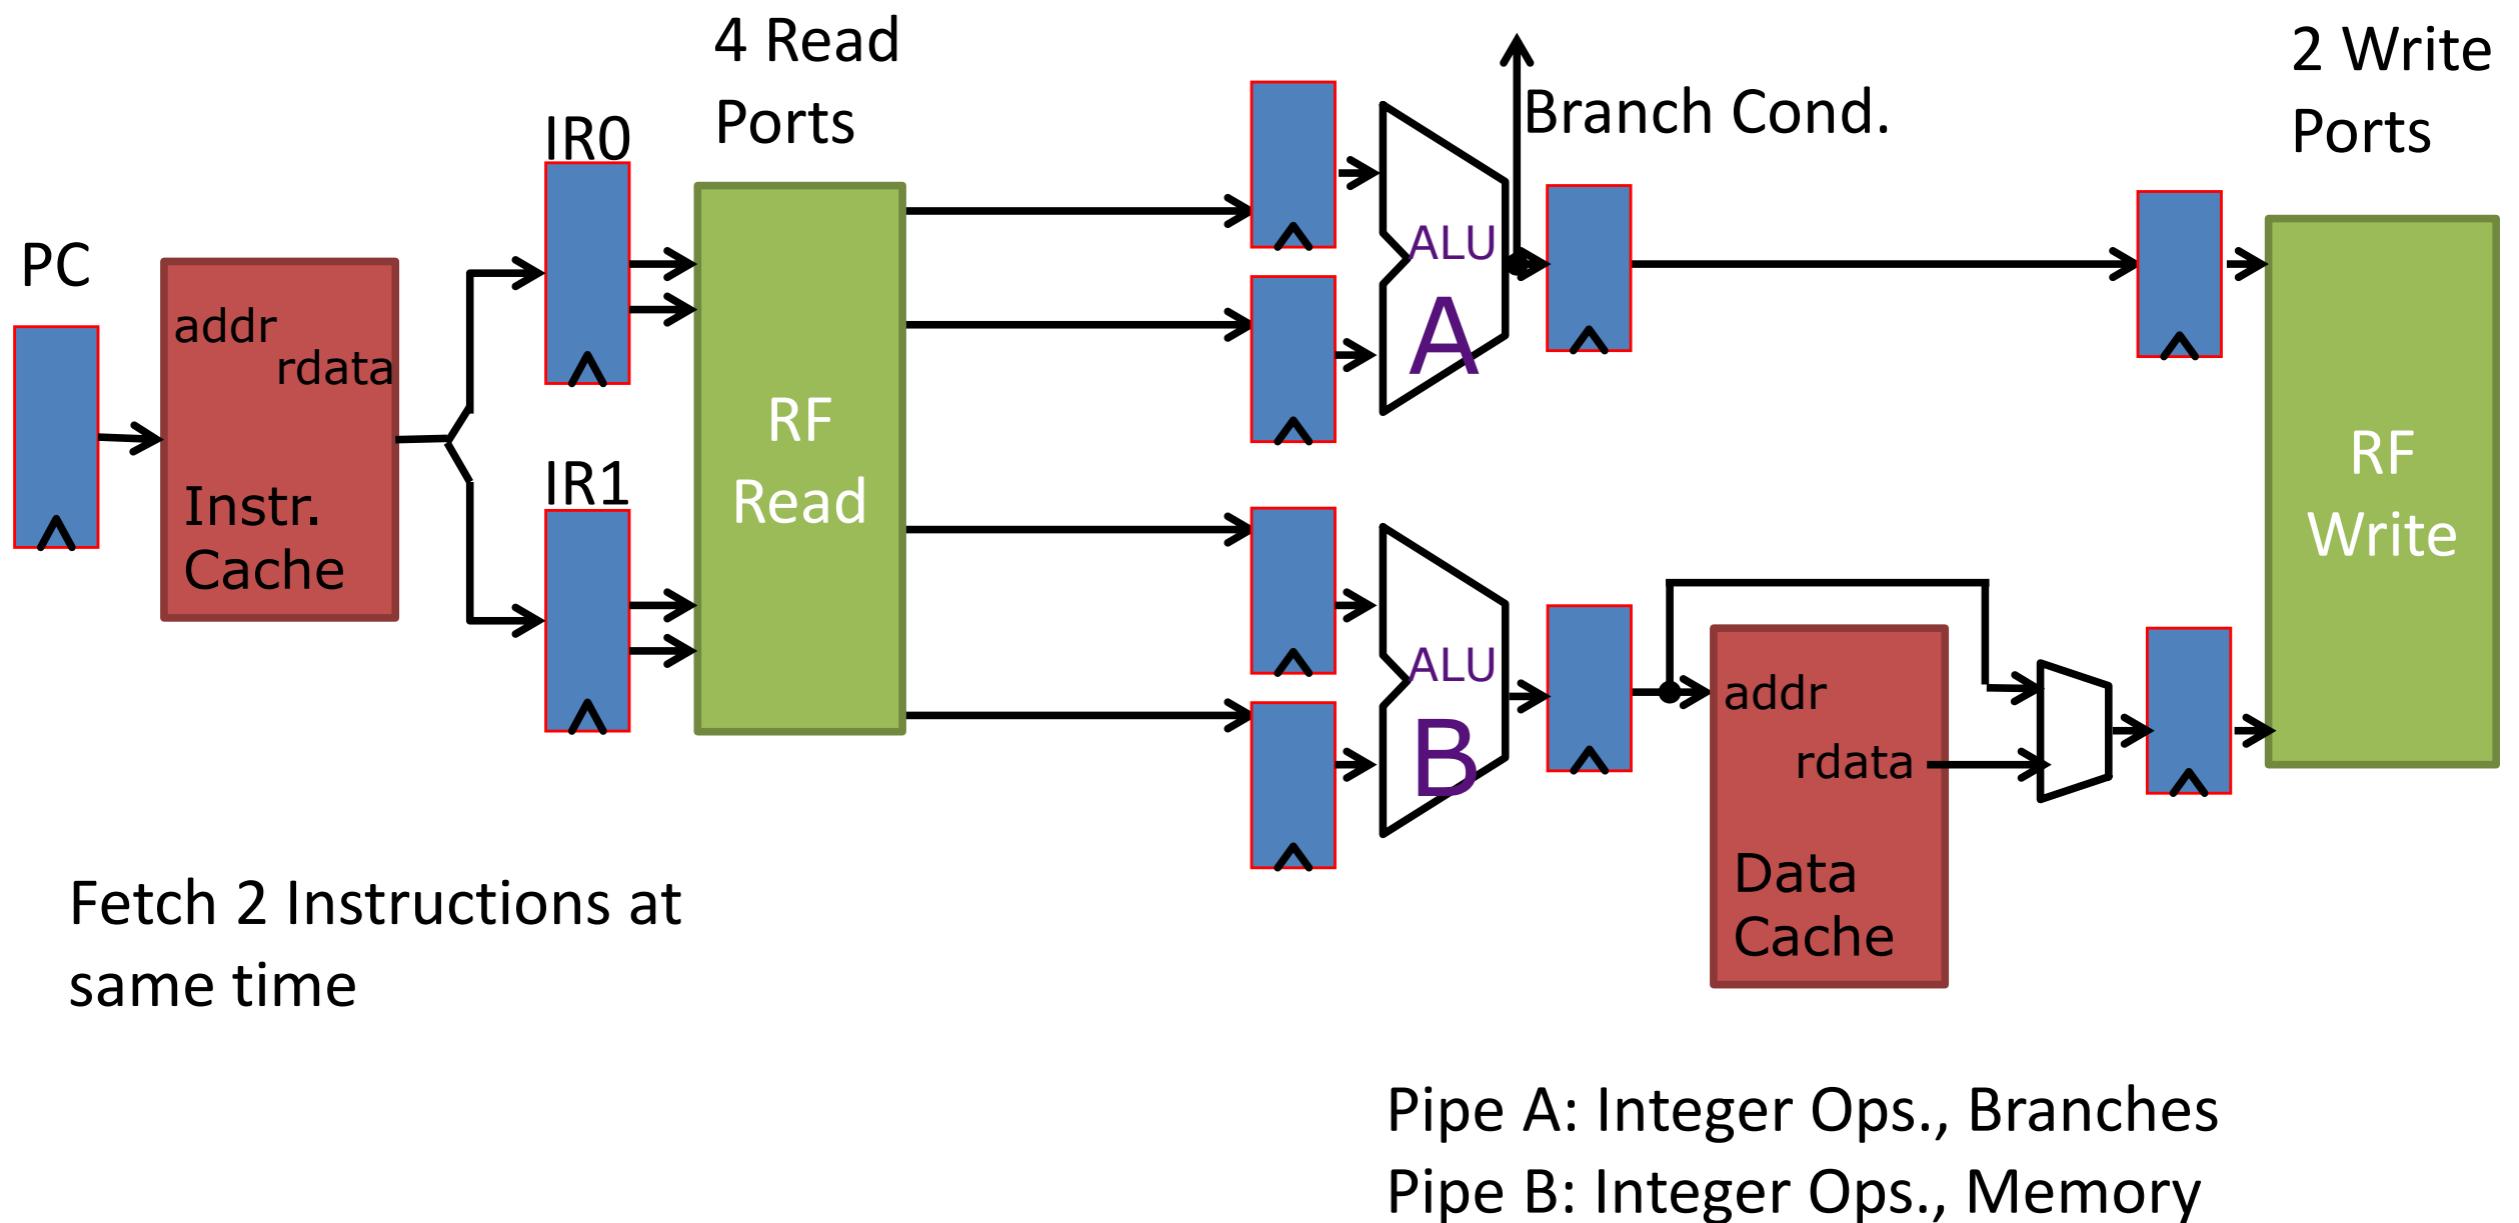
- Processors studied so far are fundamentally limited to $CPI \geq 1$
- Superscalar processors enable $CPI < 1$ ($IPC > 1$) by executing multiple instructions in parallel
- Can have both in-order and out-of-order superscalar processors. We will start with in-order.

Baseline 2-Way In-Order Superscalar Processor

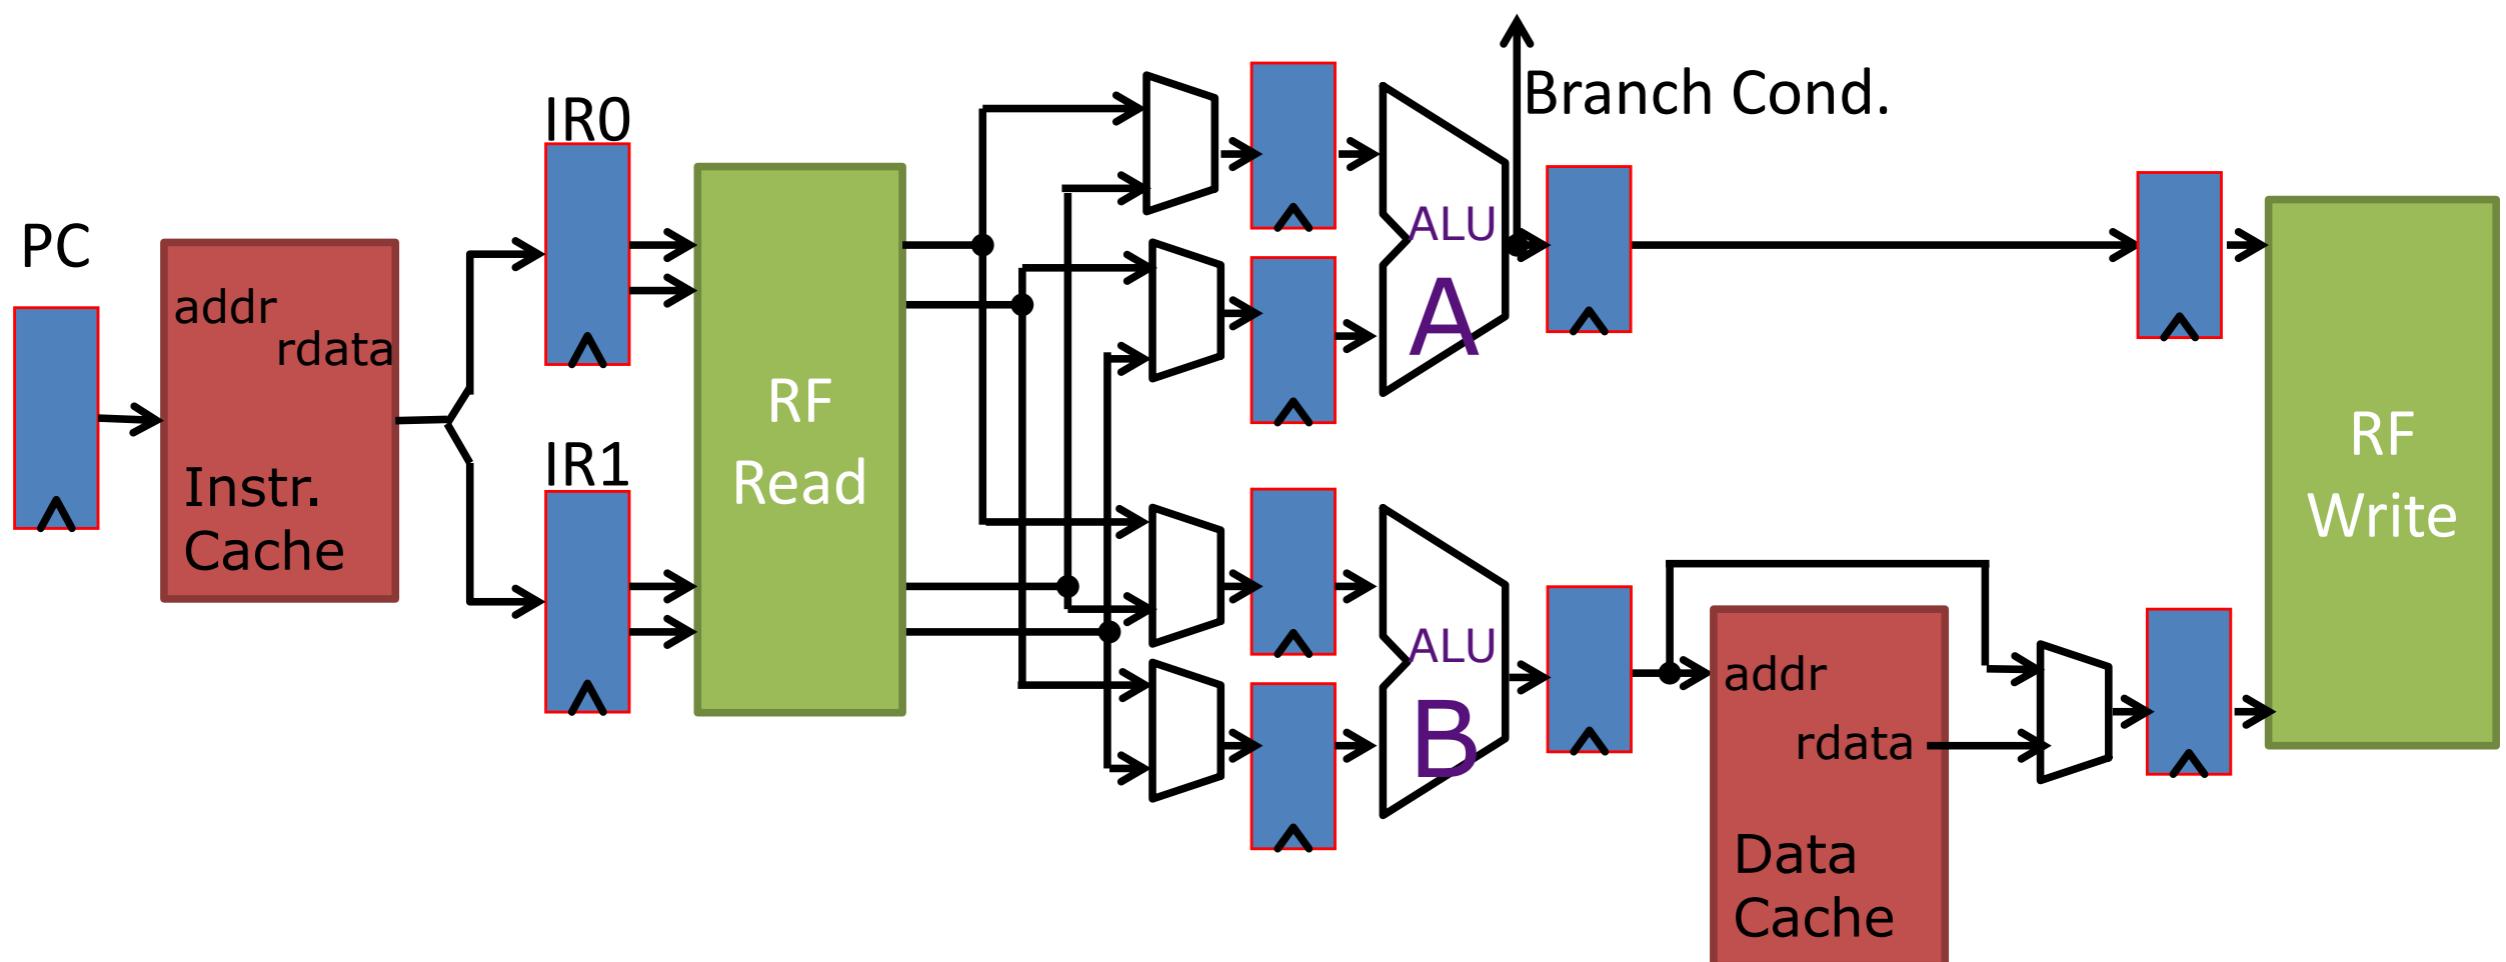


Pipe A: Integer Ops., Branches
Pipe B: Integer Ops., Memory

Baseline 2-Way In-Order Superscalar Processor

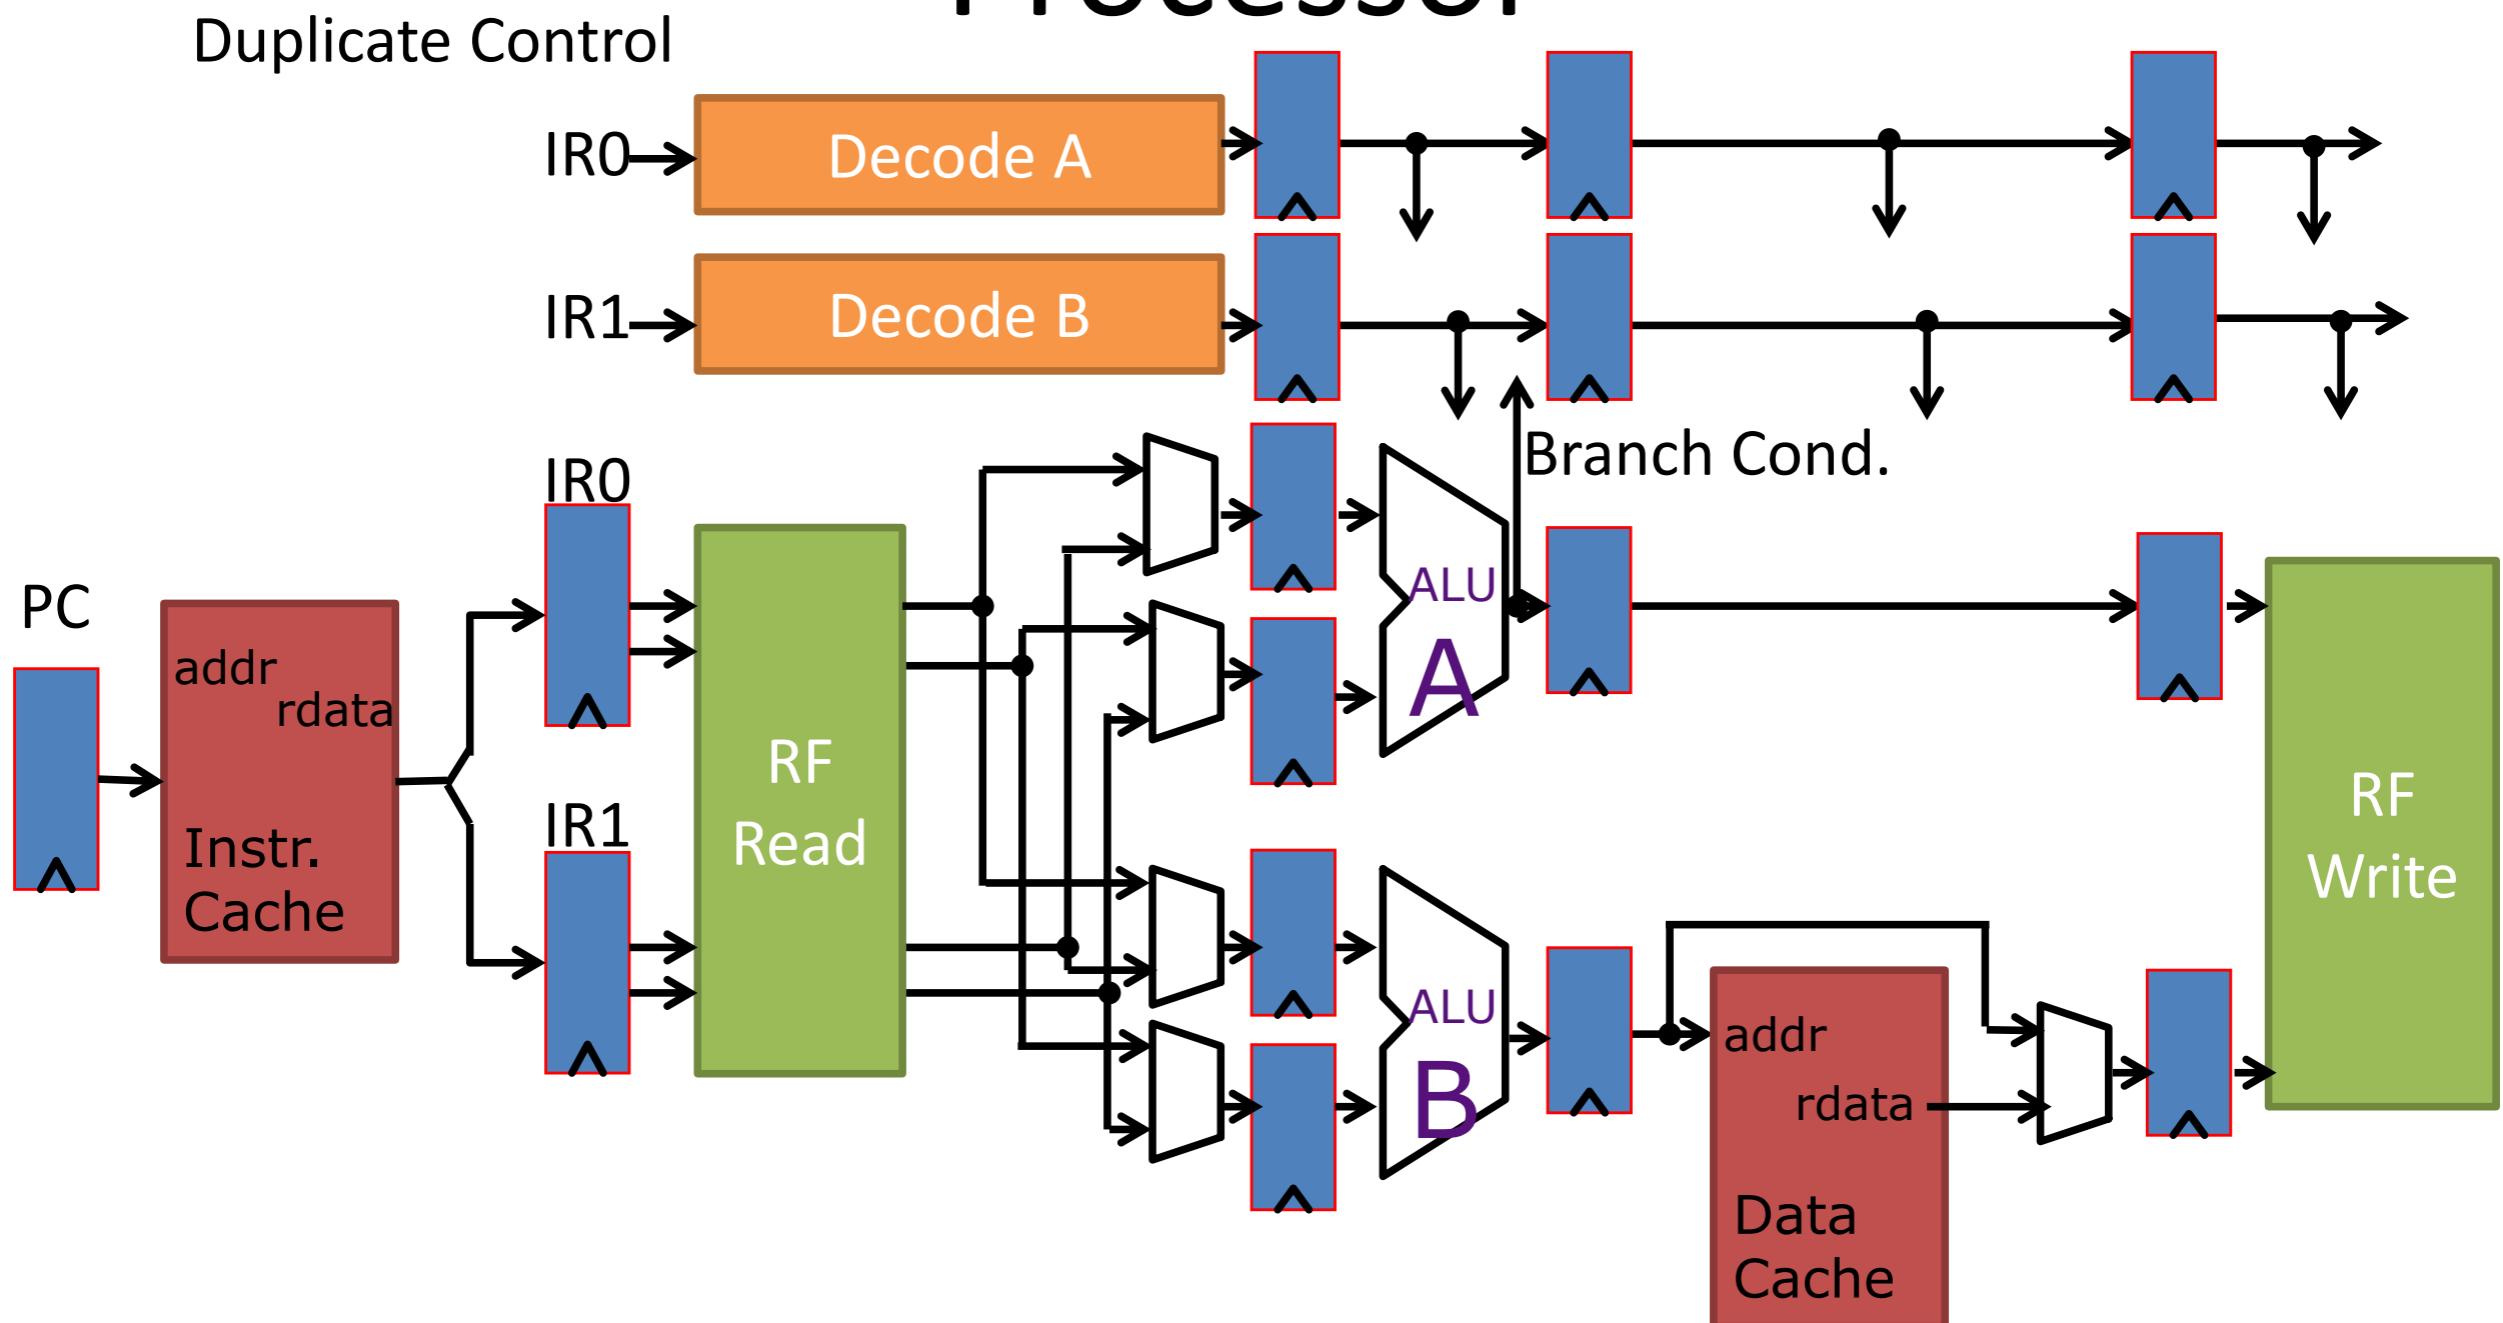


Baseline 2-Way In-Order Superscalar Processor



Pipe A: Integer Ops., Branches
Pipe B: Integer Ops., Memory

Baseline 2-Way In-Order Superscalar Processor



Pipe A: Integer Ops., Branches
Pipe B: Integer Ops., Memory

Issue Logic Pipeline Diagrams

OpA	F	D	A0	A1	W		
OpB	F	D	B0	B1	W		
OpC		F	D	A0	A1	W	
OpD		F	D	B0	B1	W	
OpE			F	D	A0	A1	W
OpF			F	D	B0	B1	W

CPI = 0.5 (IPC = 2)

Double Issue Pipeline
Can have two instructions in
same stage at same time

Dual Issue Data Hazards

No Bypassing:

ADDIU R1,R1,1 F D A0 A1 W

ADDIU R3,R4,1 F D B0 B1 W

ADDIU R5,R6,1 F D A0 A1 W

ADDIU R7,R5,1 F D D D D B0 B1 W

Dual Issue Data Hazards

No Bypassing:

ADDIU R1,R1,1 F D A0 A1 W

ADDIU R3,R4,1 F D B0 B1 W

ADDIU R5,R6,1 F D A0 A1 W

ADDIU R7,R5,1 F D D D D B0 B1 W

Full Bypassing:

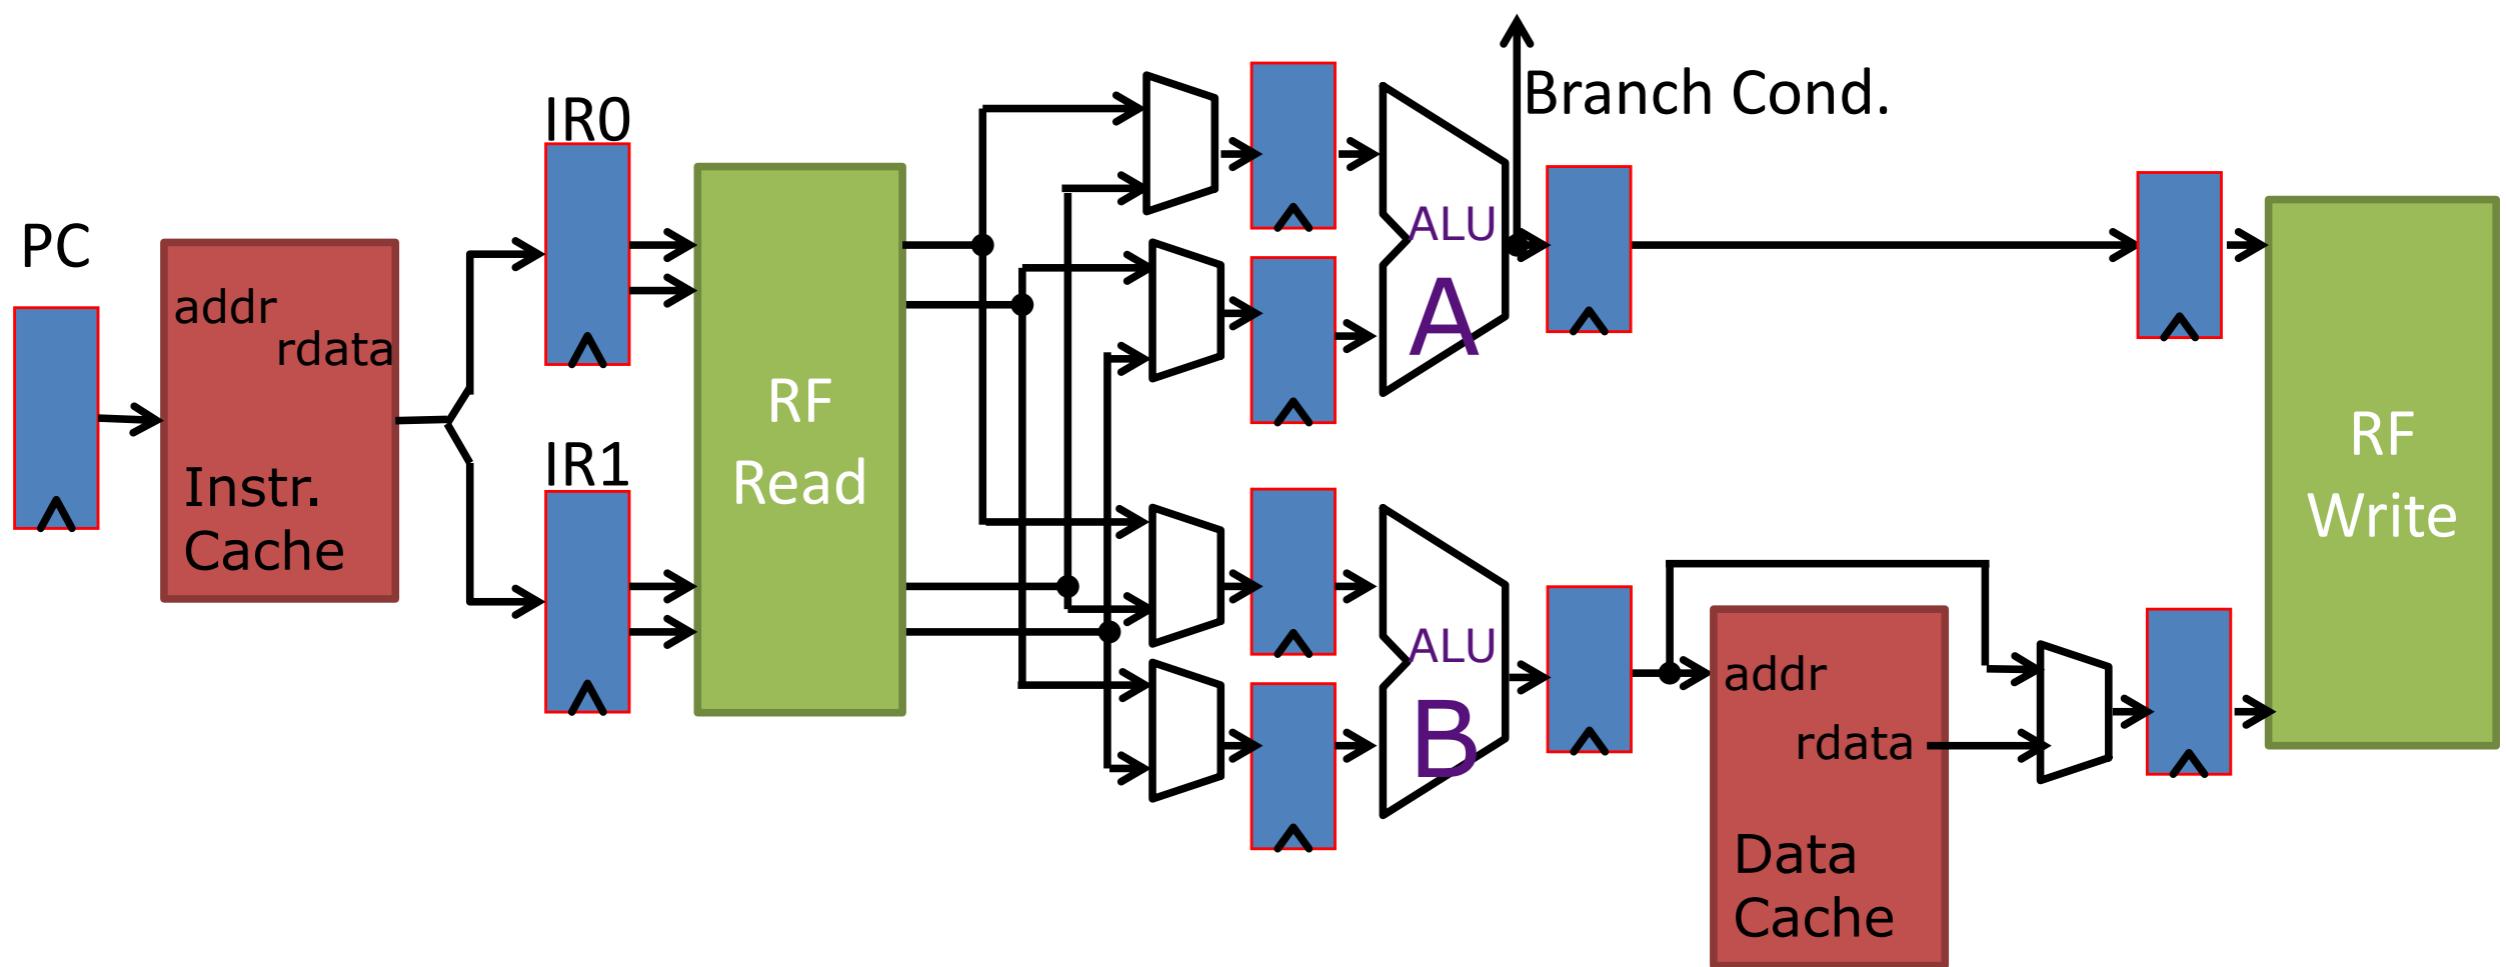
ADDIU R1,R1,1 F D A0 A1 W

ADDIU R3,R4,1 F D B0 B1 W

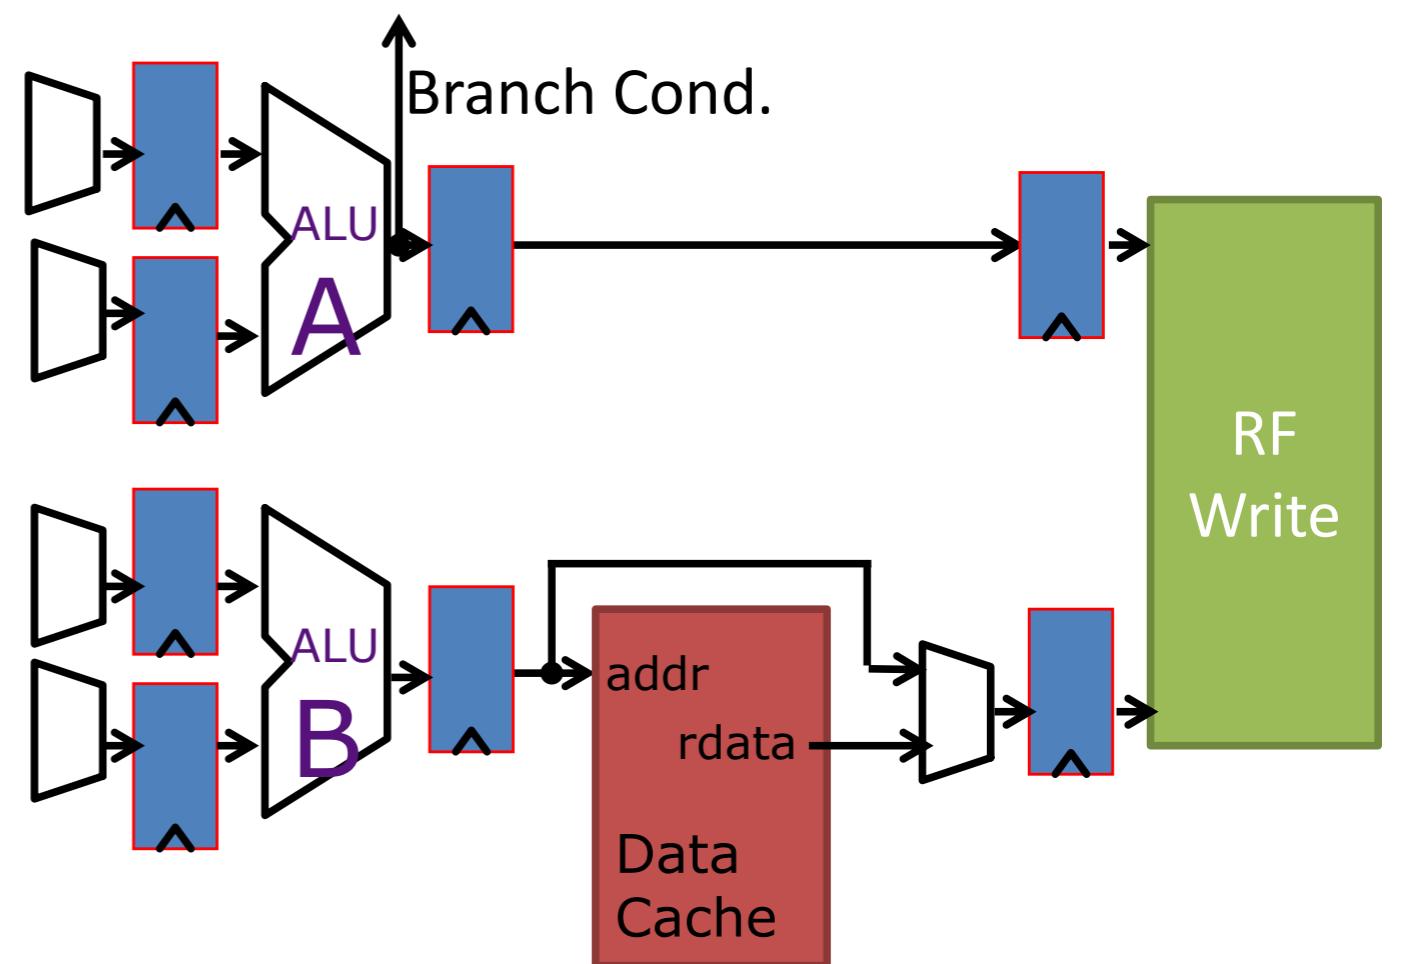
ADDIU R5,R6,1 F D A0 A1 W

ADDIU R7,R5,1 F D D B0 B1 W

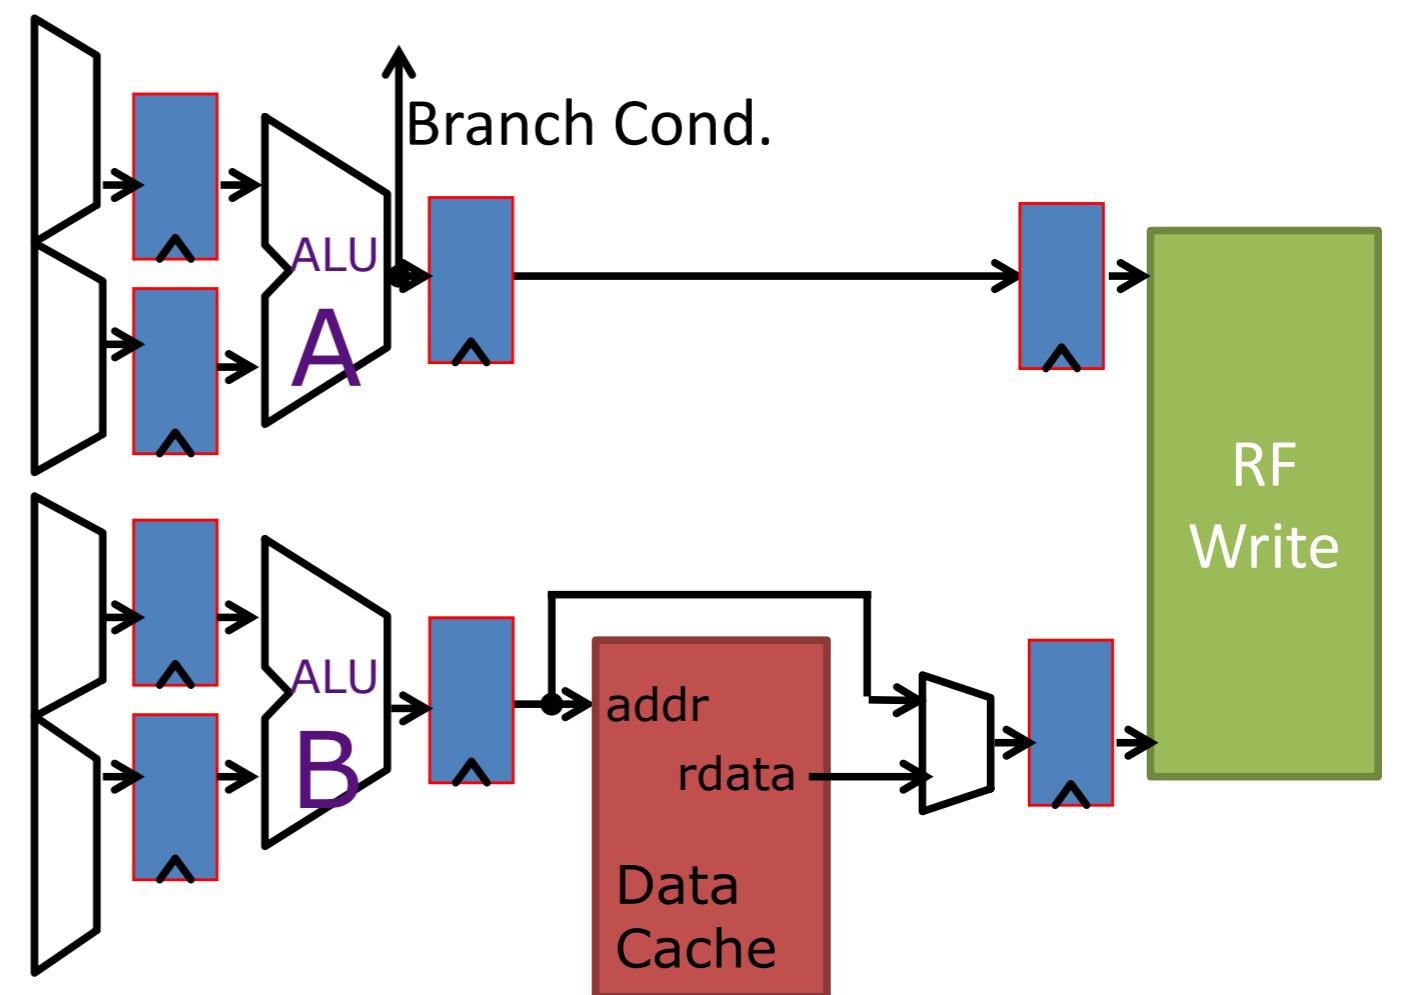
Bypassing in Superscalar Pipelines



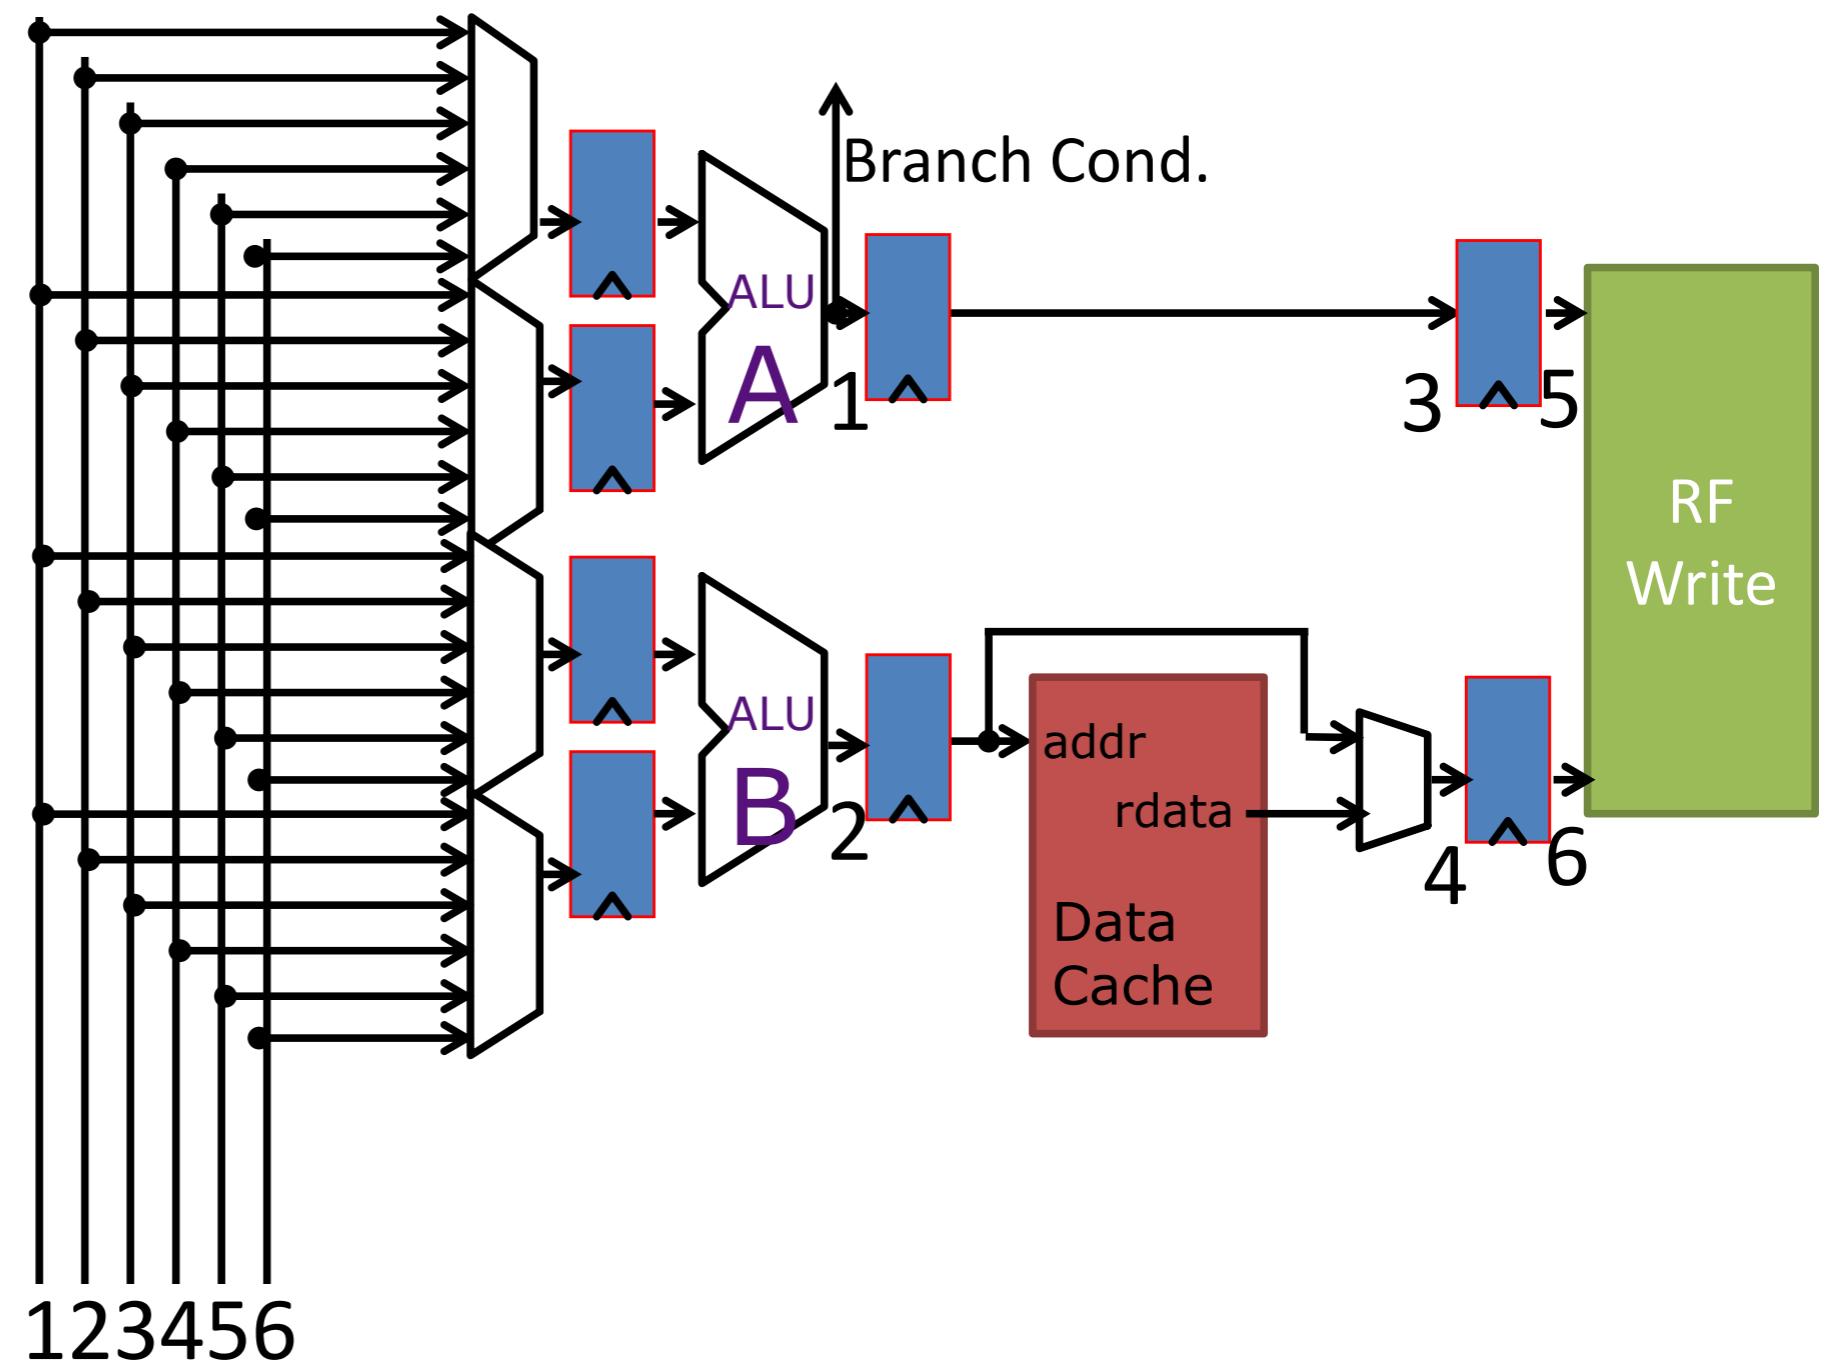
Bypassing in Superscalar Pipelines



Bypassing in Superscalar Pipelines



Bypassing in Superscalar Pipelines



Breaking Decode and Issue Stage

- Bypass Network can become very complex
- Can motivate breaking Decode and Issue Stage

D = Decode, Possibly resolve structural Hazards

I = Register file read, Bypassing, Issue/Steer
Instructions to proper unit

OpA F D I A0 A1 W

OpB F D I B0 B1 W

OpC F D I A0 A1 W

OpD F D I B0 B1 W

Dual Issue Data Hazards

Order Matters:

ADDIU R1,R1,1 F D A0 A1 W

ADDIU R3,R4,1 F D B0 B1 W

ADDIU R7,R5,1 F D A0 A1 W

ADDIU R5,R6,1 F D B0 B1 W

WAR Hazard Possible?

Fetch Logic and Alignment

Cyc	Addr	Instr
-----	------	-------

0	0x000	OpA
0	0x004	OpB
1	0x008	OpC
1	0x00C	J 0x100
...		
2	0x100	OpD
2	0x104	J 0x204
...		
3	0x204	OpE
3	0x208	J 0x30C
...		
4	0x30C	OpF
4	0x310	OpG
5	0x314	OpH

0x000	0	0	1	1
...				
0x100	2	2		
...				
0x200		3	3	
...				
0x300				4
0x310	4	5		

Fetching across cache Lines is
very hard. May need extra ports

Fetch Logic and Alignment

Cyc	Addr	Instr	Ideal, No Alignment Constraints					
0	0x000	OpA						
0	0x004	OpB						
1	0x008	OpC	OpA	F	D	A0	A1	W
1	0x00C	J 0x100	OpB	F	D	B0	B1	W
...			OpC		F	D	B0	B1 W
2	0x100	OpD	J		F	D	A0	A1 W
2	0x104	J 0x204	OpD		F	D	B0	B1 W
...			J		F	D	A0	A1 W
3	0x204	OpE	OpE		F	D	B0	B1 W
3	0x208	J 0x30C	J		F	D	A0	A1 W
...			OpF		F	D	A0	A1 W
4	0x30C	OpF	OpG		F	D	B0	B1 W
4	0x310	OpG	OpH		F	D	A0	A1 W
5	0x314	OpH						

With Alignment Constraints

Cyc	Addr	Instr	F	D	A0	A1	W
1	0x000	OpA	F	D	A0	A1	W
1	0x004	OpB	F	D	B0	B1	W
2	0x008	OpC	F	D	B0	B1	W
2	0x00C	J 0x100	F	D	A0	A1	W
3	0x100	OpD	F	D	B0	B1	W
3	0x104	J 0x204	F	D	A0	A1	W
4	0x200	?	F	-	-	-	-
4	0x204	OpE	F	D	A0	A1	W
5	0x208	J 0x30C	F	D	A0	A1	W
5	0x20C	?	F	-	-	-	-
6	0x308	?	F	-	-	-	-
6	0x30C	OpF	F	D	A0	A1	W
7	0x310	OpG	F	D	A0	A1	W
7	0x314	OpH	F	D	B0	B1	W

Superscalars Multiply Branch Cost

BEQZ	F	D	I	A0	A1	W
OpA	F	D	I	B0	-	-
OpB		F	D	I	-	-
OpC		F	D	I	-	-
OpD		F	D	-	-	-
OpE		F	D	-	-	-
OpF		F	-	-	-	-
OpG		F	-	-	-	-
OpH			F	D	I	A0 A1 W
OpI			F	D	I	B0 B1 W