

# CSCI 564 Advanced Computer Architecture

## Lecture 10: Out-of-Order Execution

---

Akshit Sharma (with modifications by Dr. Bo Wu and Ismet Dagli)


March 25, 2024

Colorado School of Mines

# Review of In-Order Pipelines

---

# An In-Order Pipeline



- **Problem:** A true data dependency stalls *dispatch* of younger instructions into functional (that is, execution) units.
- **Dispatching** an instruction is the act of sending that instruction to a functional unit.

# Know the Enemy I

What do the following two pieces of code have in common (with respect to execution in the previous design)?

```
IMUL R3 <- R1, R2
ADD  R3 <- R3, R1
ADD  R1 <- R6, R7
IMUL R5 <- R6, R8
ADD  R7 <- R9, R9
```

```
LD   R3 <- R1(0)
ADD  R3 <- R3, R1
ADD  R1 <- R6, R7
IMUL R5 <- R6, R8
ADD  R7 <- R9, R9
```

Answer: the first ADD stalls the whole pipeline!

- ADD cannot dispatch because its source registers are unavailable.
- Later *independent* instructions cannot get executed.

# Know the Enemy I

What do the following two pieces of code have in common (with respect to execution in the previous design)?

```
IMUL R3 <- R1, R2
```

```
ADD  R3 <- R3, R1
```

```
ADD  R1 <- R6, R7
```

```
IMUL R5 <- R6, R8
```

```
ADD  R7 <- R9, R9
```

```
LD   R3 <- R1(0)
```

```
ADD  R3 <- R3, R1
```

```
ADD  R1 <- R6, R7
```

```
IMUL R5 <- R6, R8
```

```
ADD  R7 <- R9, R9
```

**Answer: the first ADD stalls the whole pipeline!**

- ADD cannot dispatch because its source registers are unavailable.
- Later *independent* instructions cannot get executed.

# Know the Enemy II

What is *different* between the following two pieces of code (with respect to execution in the previous design)?

```
IMUL R3 <- R1, R2
ADD  R3 <- R3, R1
ADD  R1 <- R6, R7
IMUL R5 <- R6, R8
ADD  R7 <- R9, R9
```

```
LD   R3 <- R1(0)
ADD  R3 <- R3, R1
ADD  R1 <- R6, R7
IMUL R5 <- R6, R8
ADD  R7 <- R9, R9
```

**Answer: load latency is variable** (we don't know if it will be a cache hit or miss until runtime)

- What does this affect? One can be solved by the compiler, the other requires microarchitectural support.

# Know the Enemy II

What is *different* between the following two pieces of code (with respect to execution in the previous design)?

```
IMUL R3 <- R1, R2
```

```
ADD  R3 <- R3, R1
```

```
ADD  R1 <- R6, R7
```

```
IMUL R5 <- R6, R8
```

```
ADD  R7 <- R9, R9
```

```
LD   R3 <- R1(0)
```

```
ADD  R3 <- R3, R1
```

```
ADD  R1 <- R6, R7
```

```
IMUL R5 <- R6, R8
```

```
ADD  R7 <- R9, R9
```

**Answer: load latency is variable** (we don't know if it will be a cache hit or miss until runtime)

- What does this affect? One can be solved by the compiler, the other requires microarchitectural support.

# Preventing Dispatch Stalls

There are many ways to prevent dispatch stalls. We have already seen two:



# Preventing Dispatch Stalls

There are many ways to prevent dispatch stalls. We have already seen two:

1. Fine-grained multithreading (execute a different thread each cycle)
2. Compile-time instruction scheduling/reordering

# Preventing Dispatch Stalls

There are many ways to prevent dispatch stalls. We have already seen two:

1. Fine-grained multithreading (execute a different thread each cycle)
2. Compile-time instruction scheduling/reordering

What are the disadvantages of the above two?

# Preventing Dispatch Stalls

There are many ways to prevent dispatch stalls. We have already seen two:

1. Fine-grained multithreading (execute a different thread each cycle)
2. Compile-time instruction scheduling/reordering

What are the disadvantages of the above two?

- Requires thread saturation
- Not dynamic

Thoughts for how to improve?

# Introducing Out-of-Order Execution

---

# Out-of-Order Execution (Dynamic Scheduling)

**Idea:** Move the dependent instructions out of the way of independent ones.

- Add *rest areas* for dependent instructions (*reservation stations*)

Then, we can monitor the source *values* of each instruction in the rest area.

When all source *values* of an instruction are available, dispatch the instruction.

**Benefit: latency tolerance:** allows independent instructions to execute and complete in the presence of a long-latency operation.

# Out-of-Order Execution (Dynamic Scheduling)

**Idea:** Move the dependent instructions out of the way of independent ones.

- Add *rest areas* for dependent instructions (*reservation stations*)

Then, we can monitor the source *values* of each instruction in the rest area.

When all source *values* of an instruction are available, dispatch the instruction.

**Benefit: latency tolerance:** allows independent instructions to execute and complete in the presence of a long-latency operation.

# Out-of-Order Execution (Dynamic Scheduling)

**Idea:** Move the dependent instructions out of the way of independent ones.

- Add *rest areas* for dependent instructions (*reservation stations*)


Then, we can monitor the source *values* of each instruction in the rest area.

When all source *values* of an instruction are available, dispatch the instruction.

**Benefit: latency tolerance:** allows independent instructions to execute and complete in the presence of a long-latency operation.

# In-order vs. Out-of-order Dispatch

- In order dispatch + precise exceptions:



IMUL R3  $\leftarrow$  R1, R2


ADD R3  $\leftarrow$  R3, R1

ADD R1  $\leftarrow$  R6, R7

IMUL R5  $\leftarrow$  R6, R8

ADD R7  $\leftarrow$  R3, R5

- Out-of-order dispatch + precise exceptions:




- 15 vs. 12 cycles



# Data Dependence

Flow dependence


$r_3 \leftarrow r_1 \text{ op } r_2$   
 $r_5 \leftarrow r_3 \text{ op } r_4$



Read-after-Write  
(RAW)

Anti dependence


$r_3 \leftarrow r_1 \text{ op } r_2$   
 $r_1 \leftarrow r_4 \text{ op } r_5$



Write-after-Read  
(WAR)

Output-dependence

$r_3 \leftarrow r_1 \text{ op } r_2$   
 $r_5 \leftarrow r_3 \text{ op } r_4$   
 $r_3 \leftarrow r_6 \text{ op } r_7$



Write-after-Write  
(WAW)

# Register Renaming

```
ADD R1 <- R2, R3
```

```
ADD R1 <- R4, R5
```

```
ADD R1 <- R2, R3
```

```
ADD R6 <- R4, R5
```

# Register Renaming

```
ADD R1 <- R2, R3
```

```
ADD R1 <- R4, R5
```

```
ADD R8 <- R1, R5
```

```
ADD R1 <- R2, R3
```

```
ADD R6 <- R4, R5
```

# Register Renaming

```
ADD R1 <- R2, R3
```

```
ADD R1 <- R4, R5
```

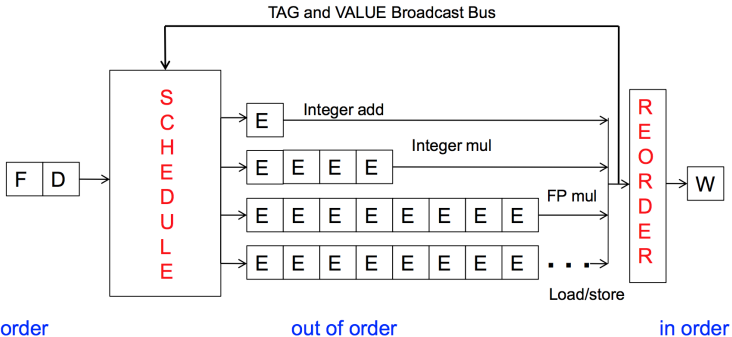
```
ADD R8 <- R1, R5
```

```
ADD R1 <- R2, R3
```

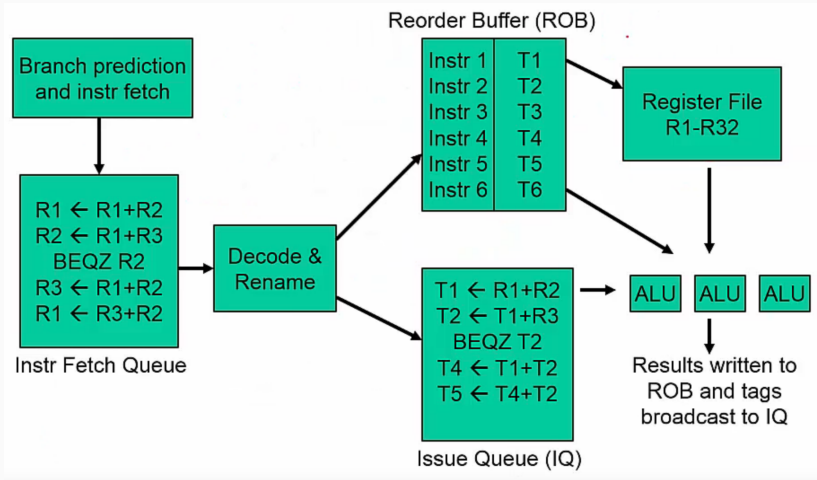
```
ADD R6 <- R4, R5
```

```
ADD R8 <- R6, R5
```

# Overview of Out-of-Order Execution



# Out-of-Order Implementation



## Practice: Rename Registers: Worksheet Problem 1

Show the renamed versions of the following code. Assume you have 4 rename registers, T1-T4.

```
ADD    R1, R2, R3
ADD    R3, R4, R5
BEQZ   R1, 1000
ADD    R1, R1, R3
ADD    R1, R1, R3
ADD    R3, R1, R3
```

## Practice: Rename Registers: Worksheet Problem 2

Show the renamed versions of the following code. Assume you have 36 physical registers and 32 architected registers.

```
ADD    R1, R2, R3
ADD    R3, R4, R5
BEQZ   R1, 1000
ADD    R1, R1, R3
ADD    R1, R1, R3
ADD    R3, R1, R3
ADD    R4, R3, R1
```



# The Dataflow Model

---

# The Dataflow Model


- Von Neumann model: An instruction is fetched and executed in **control flow order**
  - As specified by the **instruction pointer**
  - Sequential unless explicit control flow instruction
- Dataflow model: An instruction is fetched and executed in **data flow order**
  - i.e., when its operands are ready
  - i.e., there is **no instruction pointer**
  - Instruction ordering specified by data flow dependence
    - Each instruction specifies “who” should receive the result
    - An instruction can “fire” whenever all operands are received
  - Potentially many instructions can execute at the same time
    - Inherently more parallel

# von Neumann vs Dataflow

- Consider a von Neumann program
  - What is the significance of the program order?
  - What is the significance of the storage locations?

```
v <= a + b;  
w <= b * 2;  
x <= v - w  
y <= v + w  
z <= x * y
```

Sequential





Dataflow

- Which model is more natural to you as a programmer?

# More on Data Flow


- In a data flow machine, a program consists of data flow nodes
  - A data flow node fires (fetched and executed) when all its inputs are ready
    - i.e. when all inputs have tokens
- Data flow node and its ISA representation




# Dataflow Nodes

- A small set of dataflow operators can be used to define a general programming language


Fork








Primitive Ops



Switch



Merge



# Dataflow Graphs

$\{x = a + b;$   
 $y = b * 7$   
 $\text{in}$   
 $(x - y) * (x + y)\}$

- Values in dataflow graphs are represented as tokens

token

$\langle \text{ip}, p, v \rangle$


instruction ptr

port

data

$\text{ip} = 3$   
 $p = L$

- An operator executes when all its input tokens are present; copies of the result token are distributed to the destination operators




no separate control flow

# Control Flow vs. Data Flow

$a := x + y$

$b := a \times a$

$c := 4 - a$



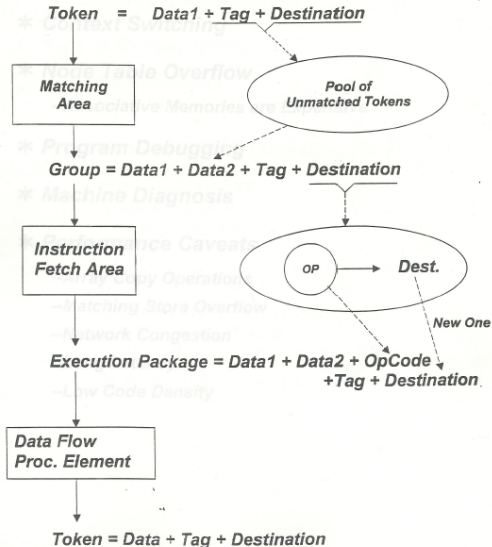
**Figure 2.** A comparison of control flow and dataflow programs. On the left a control flow program for a computer with memory-to-memory instructions. The arcs point to the locations of data that are to be used or created. Control flow arcs are indicated with dashed arrows; usually most of them are implicit. In the equivalent dataflow program on the right only one memory is involved. Each instruction contains pointers to all instructions that consume its results.

# Data Flow Characteristics

- Data-driven execution of instruction-level graphical code
  - Nodes are operators
  - Arcs are data (I/O)
  - As opposed to control-driven execution
- Only real dependencies constrain processing
- No sequential instruction stream
  - No program counter
- Execution triggered by the presence/readiness of data



# A Dataflow Processor



# Data Flow Advantages/Disadvantages

- Advantages

- Very good at exploiting **irregular parallelism**
- Only real dependencies constrain processing

- Disadvantages

- No precise state
  - Debugging very difficult
  - Interrupt/exception handling is difficult
- Bookkeeping overhead (tag matching)
- Too much parallelism? (Parallelism control needed)